

Foundations of Machine Learning CentraleSupélec — Fall 2017

10. Artificial Neural Networks

Chloé-Agathe Azencott

Centre for Computational Biology, Mines ParisTech
chloe-agathe.azencott@mines-paristech.fr




Learning objectives

- Draw a **perceptron** and write out its **decision function**.
- Implement the **learning algorithm** for a perceptron.
- Write out the **decision function** and **weight updates** for any **multiple layer perceptron**.
- Design and train a **multiple layer perceptron**.

The human brain


- Networks of **processing units** (neurons) with **connections** (synapses) between them
- Large number of neurons: 10^{10}
- Large connectivity: 10^4
- Parallel processing
- Distributed computation/memory
- Robust to noise, failures



1950s – 1970s: The perceptron


Perceptron

[Rosenblatt, 1958]




Perceptron

[Rosenblatt, 1958]




How can we do classification?

Classification with the perceptron



What is the shape of the decision boundary?


Classification with the perceptron



The decision boundary is a hyperplane (a line in dim 2).

Which other methods have we seen that yield decision boundaries that are lines/hyperplanes?

Classification with the perceptron




$$f(\mathbf{x}) = s \left(w_0 + \sum_{j=1}^p w_j x_j \right) = s(\mathbf{w}^\top \mathbf{x})$$

$$s(o(\mathbf{x})) = \begin{cases} 1 & \text{if } o(\mathbf{x}) > 0 \\ 0 & \text{otherwise.} \end{cases}$$


What if instead of just a decision (+/-) we want to output the probability of belonging to the positive class?

Classification with the perceptron



$$f(\mathbf{x}) = \sigma \left(w_0 + \sum_{j=1}^p w_j x_j \right) = \sigma(\mathbf{w}^\top \mathbf{x})$$

$$\sigma(u) = \frac{1}{1 + \exp(-u)}$$





Probability of belonging to the positive class:

$$f(\mathbf{x}) = \text{logistic}(\mathbf{w}^\top \mathbf{x}).$$


Perceptron: 1D summary

- **Regression:** $f(x) = w.x + w_0$
- **Classification:** $f(x) = \frac{1}{1 + \exp - (w.x + w_0)}$




Multiclass classification

Use K output units



How do we take a final decision?

Multiclass classification



- Choose C_k if

$$f_k(\mathbf{x}) = \max_{l \in \{1, \dots, K\}} f_l(\mathbf{x})$$

- To get probabilities, use the **softmax**:

$$f_k(\mathbf{x}) = \frac{\exp(o_k)}{\sum_{l=1}^K \exp(o_l)} \quad o_k = \mathbf{w}^k \top \mathbf{x}$$

$$\sigma(u) = \frac{1}{1 + e^{-u}} = \frac{e^u}{1 + e^u}$$

- If the output for one class is sufficiently larger than for the others, its softmax will be close to 1 (0 otherwise)
- Similar to the max, but differentiable.

Training a perceptron

Training a perceptron

- **Online** (instances seen one by one) vs **batch** (whole sample) learning:
 - No need to store the whole sample
 - Problem may change in time
 - Wear and degradation in system components.

Training a perceptron

- **Online** (instances seen one by one) vs **batch** (whole sample) learning:
 - No need to store the whole sample
 - Problem may change in time
 - Wear and degradation in system components.
- **Gradient descent:**
 - Start from random weights
 - After each data point, adjust the weights to minimize the error.

Training a perceptron

- Generic update rule:


$$\Delta w_j = -\eta \frac{\partial \text{Error}(f(\mathbf{x}^i), y^i)}{\partial w_j}$$

Learning rate

- After each training instance, for each weight:

$$w_j^{(t+1)} = w_j^{(t)} + \Delta w_j^{(t)}$$

$$w_j \leftarrow w_j + \Delta w_j$$



Training a perceptron: regression

- Regression

$$\text{Error}(f(\mathbf{x}^i), y^i) = \frac{1}{2}(y^i - f(\mathbf{x}^i))^2 = \frac{1}{2}(y^i - \mathbf{w}^\top \mathbf{x}^i)^2$$

- What is the update rule?

Remember the generic update rule:

$$w_j^{(t+1)} = w_j^{(t)} + \Delta w_j^{(t)} \quad \Delta w_j = -\eta \frac{\partial \text{Error}(f(\mathbf{x}^i), y^i)}{\partial w_j}$$

Training a perceptron: regression

- Regression

$$\text{Error}(f(\mathbf{x}^i), y^i) = \frac{1}{2}(y^i - f(\mathbf{x}^i))^2 = \frac{1}{2}(y^i - \mathbf{w}^\top \mathbf{x}^i)^2$$

- The update rule for the regression is:

$$\Delta w_j = \eta(y^i - f(\mathbf{x}^i))x_j^i$$

Training a perceptron: classification



- Sigmoid output:

$$f(\mathbf{x}^i) = \sigma(\mathbf{w}^\top \mathbf{x}^i)$$

$$\sigma(u) = \frac{1}{1 + \exp(-u)}$$

- Cross-entropy error:

$$\text{Error}(f(\mathbf{x}^i), y^i) = -y^i \log f(\mathbf{x}^i) - (1 - y^i) \log(1 - f(\mathbf{x}^i))$$



- What is the update rule now?

Training a perceptron: classification

- **Sigmoid output:**



$$f(\mathbf{x}^i) = \sigma(\mathbf{w}^\top \mathbf{x}^i)$$

$$\sigma(u) = \frac{1}{1 + \exp(-u)}$$

$$\sigma'(u) = \sigma(u)(1 - \sigma(u))$$

- **Cross-entropy error:**

$$\text{Error}(f(\mathbf{x}^i), y^i) = -y^i \log f(\mathbf{x}^i) - (1 - y^i) \log(1 - f(\mathbf{x}^i))$$



- **Update rule for binary classification:**

$$\Delta w_j = -\eta \frac{\partial \text{Error}(f(\mathbf{x}^i), y^i)}{\partial w_j} \quad \Delta w_j = \eta (y^i - f(\mathbf{x}^i)) x_j^i$$

Training a perceptron: K classes

- **K > 2 softmax outputs:**

$$f_k(\mathbf{x}^i) = \frac{\exp(\mathbf{w}^k \top \mathbf{x}^i)}{\sum_{l=1}^K \exp(\mathbf{w}^l \top \mathbf{x}^i)}$$

- **Cross-entropy error:**

$$\text{Error}(f(\mathbf{x}^i), y^i) = - \sum_{k=1}^K y_k^i \log f_k(\mathbf{x}^i)$$

- **Update rule for K-way classification:**

$$\Delta w_j^k = \eta(y^i - f_k(\mathbf{x}^i))x_j^i$$

Training a perceptron

- **Generic update rule:**

$$\Delta w_j = \eta(y^i - f(\mathbf{x}^i))x_j^i$$

Update = Learning rate.(Desired output – Actual output).Input

- After each training instance, for each weight:

$$w_i^{(t+1)} = w_i^{(t)} + \Delta w_i^{(t)}$$

$$w_j \leftarrow w_j + \Delta w_j$$

- **What happens if**

- desired output = actual output?
- desired output < actual output?

Training a perceptron: regression

- Generic update rule:

$$\Delta w_j = \eta(y^i - f(\mathbf{x}^i))x_j^i$$

Update = Learning rate.(Desired output – Actual output).Input

- After each training instance, for each weight:

$$w_j^{(t+1)} = w_j^{(t)} + \Delta w_j^{(t)} \qquad w_j \leftarrow w_j + \Delta w_j$$

- If desired output = actual output: no change
- If desired output < actual output:
 - input > 0 \rightarrow update < 0 \rightarrow w_j smaller \rightarrow prediction \searrow
 - input < 0 \rightarrow update > 0 \rightarrow w_j bigger \rightarrow prediction \searrow

Learning boolean functions

Learning AND

x1	x2	y
0	0	0
0	1	0
1	0	0
1	1	1


Design a perceptron that learns AND.

- What is its architecture?

Learning AND

x1	x2	y
0	0	0
0	1	0
1	0	0
1	1	1

Design a perceptron that learns AND.




$$f(x) = s(w_0 + w_1x_1 + w_2x_2)$$

Draw the 4 possible points (x1, x2) and a desirable separating line. What is its equation?


Learning AND

x1	x2	y
0	0	0
0	1	0
1	0	0
1	1	1

Design a perceptron that learns AND.




$$f(x) = s(w_0 + w_1x_1 + w_2x_2)$$



Learning AND


x1	x2	y
0	0	0
0	1	0
1	0	0
1	1	1

Design a perceptron that learns AND.



$$f(x) = s(w_0 + w_1x_1 + w_2x_2)$$

x1	x2	f(x)
0	0	$s(-1.5 + 0 + 0) = s(-1.5) = 0$
0	1	$s(-1.5 + 0 + 1) = s(-0.5) = 0$
1	0	$s(-1.5 + 1 + 0) = s(-0.5) = 0$
1	1	$s(-1.5 + 1 + 1) = s(0.5) = 1$



Learning XOR


x1	x2	y
0	0	0
0	1	1
1	0	1
1	1	0

Design a perceptron that learns XOR

Learning XOR

x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0

Design a perceptron that learns XOR



Learning XOR

x1	x2	y
0	0	0
0	1	1
1	0	1
1	1	0

[Minsky and Papert, 1969]


No w_0, w_1, w_2 satisfy:

$$w_0 \leq 0$$

$$w_0 + w_2 > 0$$

$$w_0 + w_1 > 0$$

$$w_0 + w_1 + w_2 \leq 0$$



Perceptrons


M. Minsky & S. Papert, 1969

The perceptron has shown itself worthy of study despite (and even because of!) its severe limitations. It has many features to attract attention: its linearity; its intriguing learning theorem; its clear paradigmatic simplicity as a kind of parallel computation. There is no reason to suppose that any of these virtues carry over to the many-layered version. Nevertheless, we consider it to be an important research problem to elucidate (or reject) our intuitive judgement that the extension to multilayer systems is sterile.


1980s – early 1990s

Multilayer perceptrons

What is the output of the hidden unit h ?



Multilayer perceptrons




- Output of hidden unit h :

$$z_h = \frac{1}{1 + e^{-w_h^\top x}}$$

- What is the output of the network?

Multilayer perceptrons



- **Output of hidden unit h:**

$$z_h = \frac{1}{1 + e^{-w_h^\top x}}$$



- **Output of the network:**

$$\begin{aligned} f(x) &= \mathbf{v}^\top \mathbf{z} \\ &= v_0 + \sum_{h=1}^H \frac{v_h}{1 + e^{-w_h^\top x}} \end{aligned}$$



Not linear in \mathbf{x} !

Learning XOR with an MLP

Draw the geometric interpretation of this multiple layer perceptron.



Learning XOR with an MLP




Universal approximation

Any continuous function on a compact subset of \mathbb{R}^n can be approximated to any arbitrary degree of precision by a feed-forward multi-layer perceptron with a single hidden layer containing a finite number of neurons.

Cybenko (1989), Hornik (1991)

Backpropagation

Backwards propagation of errors.



$$z_h = \frac{1}{1 + e^{-\mathbf{w}_h^\top \mathbf{x}}}$$


$$\begin{aligned} f(\mathbf{x}) &= \mathbf{v}^\top \mathbf{z} \\ &= v_0 + \sum_{h=1}^H \frac{v_h}{1 + e^{-\mathbf{w}_h^\top \mathbf{x}}} \end{aligned}$$

$$\frac{\partial E}{\partial w_{hj}} = \frac{\partial E}{\partial f(\mathbf{x})} \boxed{\frac{\partial f(\mathbf{x})}{\partial z_h}} \boxed{\frac{\partial z_h}{\partial w_{hj}}}$$

The diagram shows the chain rule for the derivative of the error with respect to the weight w_{hj} . The first term is the derivative of the error with respect to the output function $f(\mathbf{x})$. The second term is the derivative of the output function with respect to the hidden node output z_h , highlighted with an orange box. The third term is the derivative of the hidden node output with respect to the weight w_{hj} , highlighted with a purple box. Arrows point from the orange box to the z_h node in the diagram and from the purple box to the w_{hj} connection in the diagram.

Backpropagation

Backwards propagation of errors.



$$z_h = \frac{1}{1 + e^{-\mathbf{w}_h^\top \mathbf{x}}}$$

$$\begin{aligned} f(\mathbf{x}) &= \mathbf{v}^\top \mathbf{z} \\ &= v_0 + \sum_{h=1}^H \frac{v_h}{1 + e^{-\mathbf{w}_h^\top \mathbf{x}}} \end{aligned}$$

$$\frac{\partial E}{\partial w_{hj}} = \frac{\partial E}{\partial f(\mathbf{x})} \boxed{\frac{\partial f(\mathbf{x})}{\partial z_h}} \boxed{\frac{\partial z_h}{\partial w_{hj}}}$$

v_h (orange arrow pointing to $\frac{\partial f(\mathbf{x})}{\partial z_h}$)
 $z_h^i (1 - z_h^i) x_j^i$ (purple arrow pointing to $\frac{\partial z_h}{\partial w_{hj}}$)

Backprop: Regression

$$\text{Error}(f(\mathbf{x}^i), y^i) = \frac{1}{2}(y^i - f(\mathbf{x}^i))^2$$

$$f(\mathbf{x}) = v_0 + \sum_{h=1}^H v_h z_h$$

$$\Delta v_h = \eta_1 (y^i - f(\mathbf{x}^i)) z_h^i$$

Forward


$$z_h = \sigma(w_h^\top \mathbf{x})$$

\mathbf{x}

Backward

$$\begin{aligned} \Delta w_{hj} &= -\eta \frac{\partial E^i}{\partial w_{hj}} \\ &= -\eta \frac{\partial E}{\partial f(\mathbf{x}^i)} \frac{\partial f(\mathbf{x}^i)}{\partial z_h^i} \frac{\partial z_h^i}{\partial w_{hj}} \\ &= -\eta (y^i - f(\mathbf{x}^i)) v_h z_h^i (1 - z_h^i) x_j^i \end{aligned}$$

Backprop: Regression



Initialize all v_h, w_{hj} to $\text{rand}(-0.01, 0.01)$

Repeat until convergence

For $i = 1, \dots, n$

For $h = 1, \dots, H$

$$z_h^i = \sigma(w_h^\top x^i)$$

$$f(x^i) = v^\top z^i$$

For $h = 1, \dots, H$

$$\Delta v_h = \eta_1(y^i - f(x^i))z_h^i$$

For $h = 1, \dots, H$

For $j = 1, \dots, p$

$$\Delta w_{hj} = \eta((y_k^i - f_k(x^i))v_h)z_h^i(1 - z_h^i)x_j^i$$


For $h = 1, \dots, H$

$$v_h \leftarrow v_h + \Delta v_h$$

For $j = 1, \dots, p$

$$w_{hj} \leftarrow w_{hj} + \Delta w_{hj}$$

Backprop: Regression



Initialize all v_h, w_{hj} to $\text{rand}(-0.01, 0.01)$

Repeat until convergence

For $i = 1, \dots, n$

For $h = 1, \dots, H$

$$z_h^i = \sigma(w_h^\top x^i)$$

$$f(x^i) = v^\top z^i$$

For $h = 1, \dots, H$

$$\Delta v_h = \eta_1(y^i - f(x^i))z_h^i$$

For $h = 1, \dots, H$

For $j = 1, \dots, p$

$$\Delta w_{hj} = \eta((y_k^i - f_k(x^i))v_h)z_h^i(1 - z_h^i)x_j^i$$

For $h = 1, \dots, H$

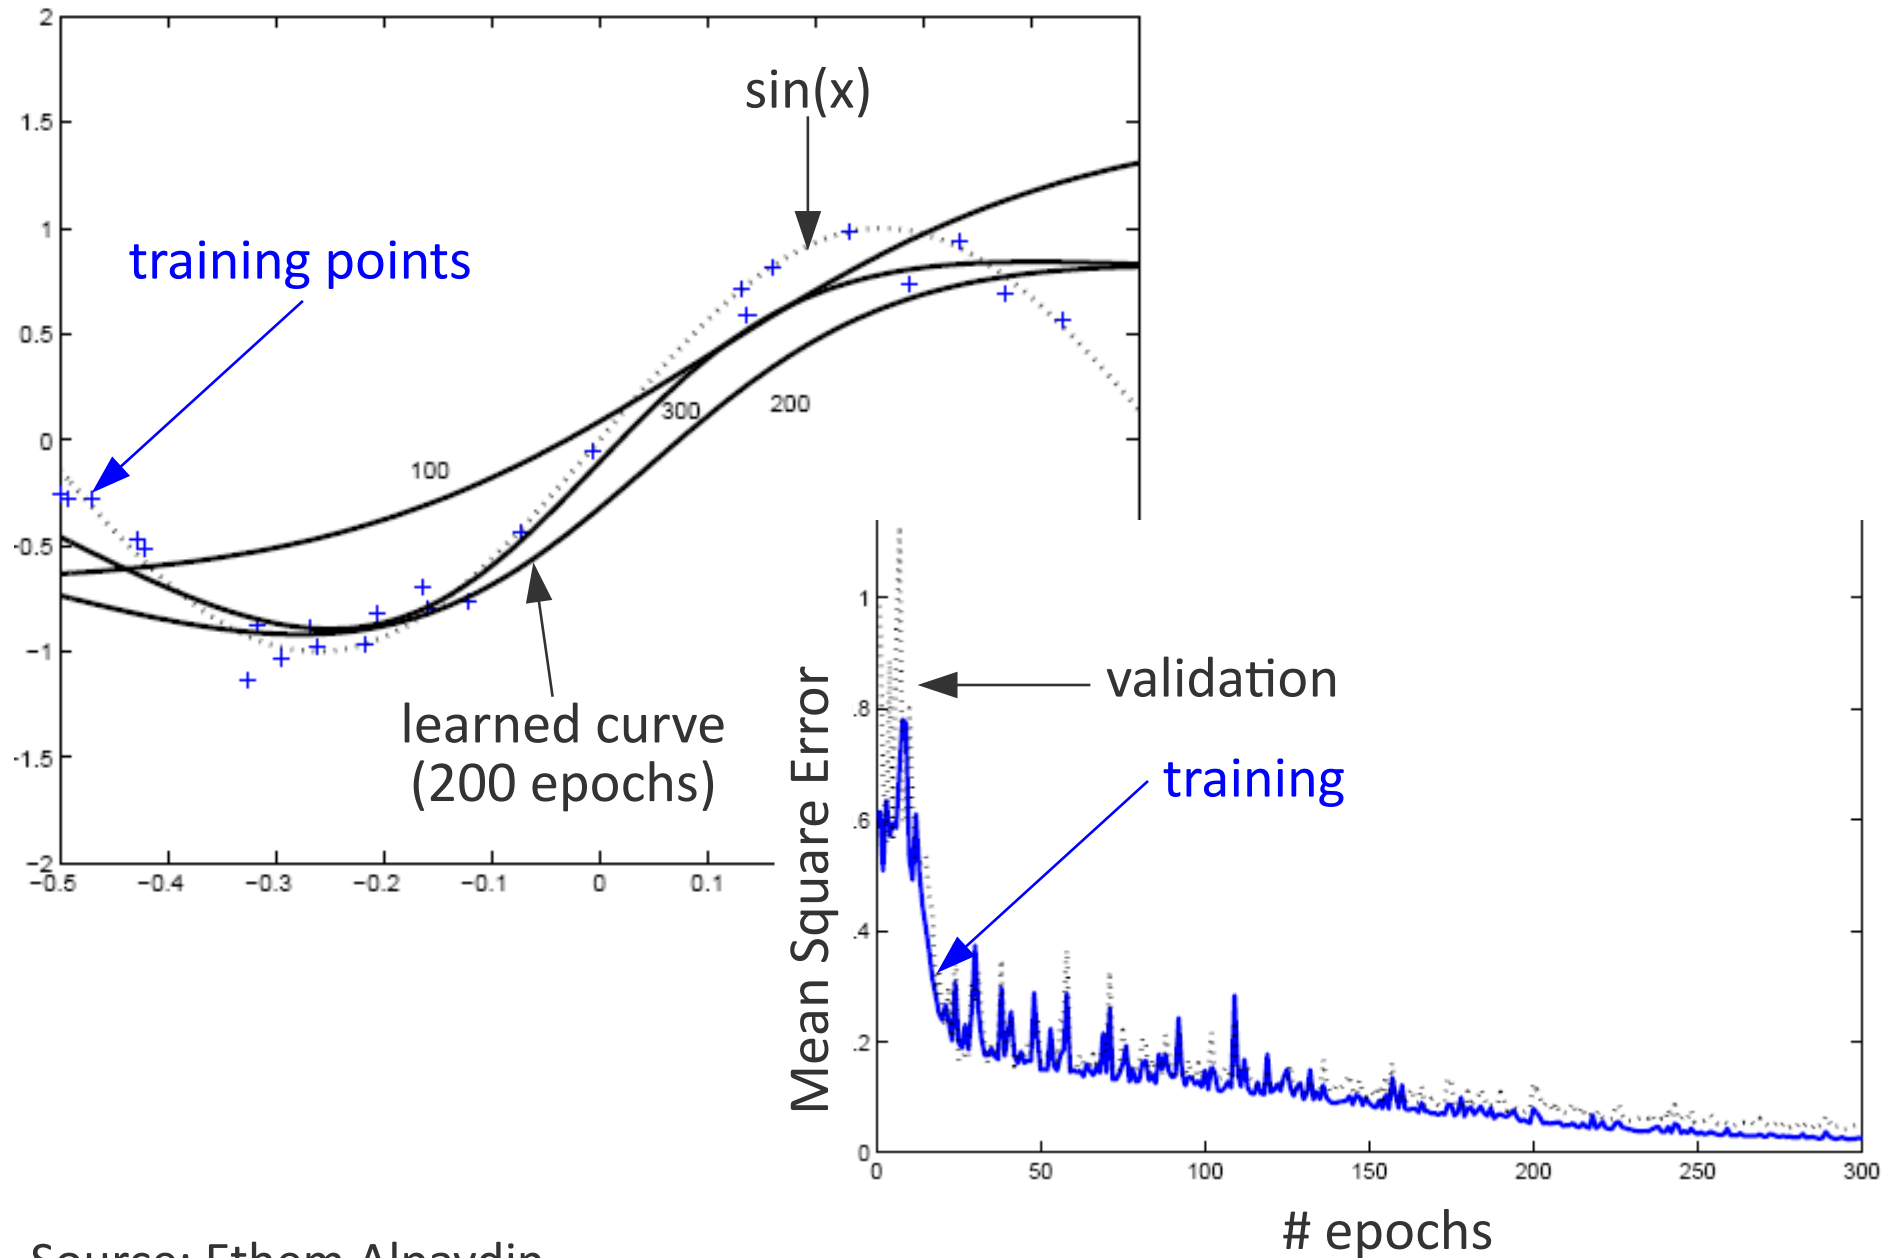
$$v_h \leftarrow v_h + \Delta v_h$$

For $j = 1, \dots, p$


$$w_{hj} \leftarrow w_{hj} + \Delta w_{hj}$$

Epoch: when all the training points have been seen once

E.g.: Learning $\sin(x)$



Backprop: Classification




- Forward:
 - $z_h = ?$
 - $f(x) = ?$

Backprop: Classification


$$z_h = \sigma(w_h^\top \mathbf{x})$$

$$f(\mathbf{x}) = \sigma \left(v_0 + \sum_{h=1}^H v_h z_h \right)$$

- Error (cross-entropy)?



Backprop: Classification



$$z_h = \sigma(w_h^\top \mathbf{x})$$

$$f(\mathbf{x}) = \sigma \left(v_0 + \sum_{h=1}^H v_h z_h \right)$$


$$\text{Error} = - \sum_{i=1}^n y^i \log(f(\mathbf{x}^i)) + (1 - y^i) \log(1 - f(\mathbf{x}^i))$$

• Backward:

- $\Delta \mathbf{v}_h$?

- $\Delta \mathbf{w}_{hj}$?

Backprop: Classification



$$z_h = \sigma(w_h^\top \mathbf{x})$$

$$f(\mathbf{x}) = \sigma \left(v_0 + \sum_{h=1}^H v_h z_h \right)$$

$$\text{Error} = - \sum_{i=1}^n y^i \log(f(x^i)) + (1 - y^i) \log(1 - f(x^i))$$

$$\Delta v_h = \eta_1 \sum_{i=1}^n (y^i - f(\mathbf{x}^i)) z_h^i$$

$$\Delta w_{hj} = \eta \sum_{i=1}^n (y^i - f(\mathbf{x}^i)) v_h z_h^i (1 - z_h^i) x_j^i$$

Backprop: K classes


$$o_k^i = v_{k0} + \sum_{h=1}^H v_{kh} z_h^i$$

$$f_k(\mathbf{x}^i) = \frac{\exp(o_k^i)}{\sum_{l=1}^K \exp(o_l^i)}$$

$$\text{Error} = - \sum_{i=1}^n \sum_{k=1}^K y_k^i \log f_k(\mathbf{x}^i)$$


$$\Delta v_{kh} = \eta_1 \sum_{i=1}^n (y_k^i - f_k(\mathbf{x}^i)) z_h^i$$

$$\Delta w_{hj} = \eta \sum_{i=1}^n \left(\sum_{k=1}^K (y_k^i - f_k(\mathbf{x}^i)) v_{kh} \right) z_h^i (1 - z_h^i) x_j^i$$



Multiple hidden layers

- The MLP with one hidden layer is a **universal approximator**
- But using **multiple layers** may lead to **simpler networks**.



$$f(\mathbf{x}) = v_0 + \sum_{l=1}^{H_2} v_l z_{2l}$$

$$z_{2l} = \sigma(\mathbf{w}_{2l}^\top \mathbf{z}_1) = \sigma \left(w_{2l0} + \sum_{h=1}^{H_1} w_{2lh} z_{1h} \right)$$

$$z_{1h} = \sigma(\mathbf{w}_{1h}^\top \mathbf{x}) = \sigma \left(w_{1h0} + \sum_{j=1}^p w_{1hj} x_j \right)$$

$$f(\mathbf{x}) = v_0 + \sum_{l=1}^{H_2} v_l \sigma \left(w_{2l0} + \sum_{h=1}^{H_1} w_{2lh} \cdot \sigma \left(w_{1h0} + \sum_{j=1}^p w_{1hj} x_j \right) \right)$$

Deep learning


- Multi-layer perceptrons with “enough” layers are **deep feed-forward neural networks**
- Nothing more than a (possibly very complicated) **parametric model**!
- Coefficients are learned by **gradient descent**
 - **local minima**
 - **vanishing/exploding gradient**
- Each layer learns a **new representation** of the data
⇒ **“representation learning”**



What makes deep networks hard to train? by Michael Nielsen
<http://neuralnetworksanddeeplearning.com/chap5.html>

(Deep) neural networks

Internal representation of the digits data



Yann Le Cun et al. (1990)

Puppy or bagel?



Photo credit: Karen Zack @teenybiscuit

Adversarial examples



panda

$+ .007 \times$



$=$



gibbon

Goodfellow et al. ICLR 2015
<https://arxiv.org/pdf/1412.6572v3.pdf>

Types of (deep) neural networks

- **Deep feed-forward** (= multilayer perceptrons)
- **Unsupervised networks**
 - **autoencoders** / **variational autoencoders** (VAE) — learn a new representation of the data
 - **deep belief networks** (DBNs) — model the distribution of the data but can add a supervised layer in the end
 - **generative adversarial networks** (GANs) — learn to separate real data from fake data they generate
- **Convolutional neural networks** (CNNs)
 - for image/audio modeling
- **Recurrent Neural Networks**
 - nodes are fed information from the previous layer and also from themselves (i.e. the past)
 - **long short-term memory networks** (LSTM) for sequence modeling.


Types of (deep) neural networks

- **Deep feed-forward** (= multilayer perceptrons)
- **Unsupervised networks**
 - **autoencoders** / **variational autoencoders** (VAE) — learn a new representation of the data
 - **deep belief networks** (DBNs) — model the distribution of the data but can add a supervised layer in the end
 - **generative adversarial networks** (GANs) — learn to separate real data from fake data they generate
- **Convolutional neural networks** (CNNs)
 - for image/audio modeling
- **Recurrent Neural Networks**
 - nodes are fed information from the previous layer and also from themselves (i.e. the past) \Rightarrow time series modeling
 - **long short-term memory networks** (LSTM) for sequence modeling.

Feature extraction: Autoencoders

Autoencoders

- Dimensionality reduction with neural networks
Rumelhart, Hinton & Williams (1986)
- **Goal:** output matches input




Compact representation of input

$$\min_{f,g} \sum_x \Delta(f \circ g(x), x)$$

$$\Delta(y, x) = ||y - x||_2^2$$

Restricted Boltzmann Machines

- Boltzmann Machines Hinton & Sejnowsky (1985)
- RBM Smolensky (1986)



- **binary units** $x_j \in \{0, 1\}, j = 1, \dots, p$ $z_h \in \{0, 1\}, h = 1, \dots, H$
(e.g. pixels in an image)

- **stochastic activation** $P(x_j = 1 | \mathbf{z}) = \sigma(a_j + \sum_{h=1}^H w_{jh} z_h)$
 $P(z_h = 1 | \mathbf{x}) = \sigma(b_h + \sum_{j=1}^p w_{jh} x_j)$

Restricted Boltzmann Machines

- **Restricted:**

Boltzmann Machines are fully connected, here there are no connections between units of the same layer.

- **Boltzmann: energy-based probabilistic models**

- **Energy** of the network:

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

Ising model (statistical physics):

- nodes = sites
- edges = adjacence
- the network is a lattice
- variables = magnetic spin (-1 or +1)

Restricted Boltzmann Machines

- **Restricted:**

Boltzmann Machines are fully connected, here there are no connections between units of the same layer.

- **Boltzmann: energy-based probabilistic models**

- **Energy** of the network:

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

- **Probability distribution**

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{(-E(\mathbf{x}, \mathbf{z}))}$$

Boltzmann factor

partition function = sum over all \mathbf{x} and \mathbf{z} of $P(\mathbf{x}, \mathbf{z})$

$$Z = \sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}$$

Restricted Boltzmann Machines

- Restricted:

Boltzmann Machines are fully connected, here there are no connections between units of the same layer.

- Boltzmann: energy-based probabilistic models

- Energy of the network:

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

- Probability distribution

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{(-E(\mathbf{x}, \mathbf{z}))}$$

Boltzmann factor

$$P(\mathbf{x}|\mathbf{z}) = \prod_{j=1}^p P(x_j|\mathbf{z})$$

$$P(x_j = 1|\mathbf{z}) = \frac{P(x_j = 1, \mathbf{z})}{P(x_j = 0, \mathbf{z}) + P(x_j = 1, \mathbf{z})} = \sigma\left(a_j + \sum_{h=1}^H w_{jh} z_h\right)$$

partition function = sum over all \mathbf{x} and \mathbf{z} of $P(\mathbf{x}, \mathbf{z})$

Restricted Boltzmann Machines

- **Restricted:**

Boltzmann Machines are fully connected, here there are no connections between units of the same layer.

- **Boltzmann: energy-based probabilistic models**

- **Energy** of the network:

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

- **Probability distribution**

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{(-E(\mathbf{x}, \mathbf{z}))}$$

- Minimizing the energy of the network = minimizing the negative log likelihood of the observed data.
- Connection to **Markov Random Fields**.

Restricted Boltzmann Machines

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{-E(\mathbf{x}, \mathbf{z})} \quad Z = \sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} \quad P(\mathbf{z}|\mathbf{x}) = \frac{P(\mathbf{x}, \mathbf{z})}{P(\mathbf{x})} = \frac{e^{-E(\mathbf{x}, \mathbf{z})}}{\sum_{\mathbf{z}'} e^{-E(\mathbf{x}, \mathbf{z}')}}$$

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p \underbrace{a_j}_{\text{bias}} x_j - \sum_{h=1}^H \underbrace{b_h}_{\text{bias}} z_h - \sum_{j=1}^p \sum_{h=1}^H x_j \underbrace{w_{jh}}_{\text{weight}} z_h$$

Gradient of the negative log likelihood:

$$P(\mathbf{x}) = \sum_{\mathbf{z}} P(\mathbf{x}, \mathbf{z})$$

$$-\log P(\mathbf{x}) = \log Z - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} = \log \sum_{\mathbf{x}', \mathbf{z}} e^{-E(\mathbf{x}', \mathbf{z})} - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$\frac{\partial -\log P(\mathbf{x})}{\partial \theta} = - \frac{1}{\sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{x}', \mathbf{z}} \frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}', \mathbf{z})} + \frac{1}{\sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{z}} \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$= - \underbrace{\sum_{\mathbf{x}', \mathbf{z}} P(\mathbf{x}', \mathbf{z}) \frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta}}_{\text{negative gradient}} + \underbrace{\sum_{\mathbf{z}} P(\mathbf{z}|\mathbf{x}) \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta}}_{\text{positive gradient}}$$

negative gradient

positive gradient

Restricted Boltzmann Machines

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{-E(\mathbf{x}, \mathbf{z})} \quad Z = \sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} \quad P(\mathbf{z}|\mathbf{x}) = \frac{P(\mathbf{x}, \mathbf{z})}{P(\mathbf{x})} = \frac{e^{-E(\mathbf{x}, \mathbf{z})}}{\sum_{\mathbf{z}'} e^{-E(\mathbf{x}, \mathbf{z}')}}$$

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

Gradient of the negative log likelihood:

$$P(\mathbf{x}) = \sum_{\mathbf{z}} P(\mathbf{x}, \mathbf{z})$$

$$-\log P(\mathbf{x}) = \log Z - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} = \log \sum_{\mathbf{x}', \mathbf{z}} e^{-E(\mathbf{x}', \mathbf{z})} - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$\frac{\partial -\log P(\mathbf{x})}{\partial \theta} = - \frac{1}{\sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{x}', \mathbf{z}} \frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}', \mathbf{z})} + \frac{1}{\sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{z}} \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$= - \sum_{\mathbf{x}', \mathbf{z}} P(\mathbf{x}', \mathbf{z}) \boxed{\frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta}} + \sum_{\mathbf{z}} P(\mathbf{z}|\mathbf{x}) \boxed{\frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta}} \quad \text{easy to compute}$$

Restricted Boltzmann Machines

$$P(\mathbf{x}, \mathbf{z}) = \frac{1}{Z} e^{-E(\mathbf{x}, \mathbf{z})} \quad Z = \sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} \quad P(\mathbf{z}|\mathbf{x}) = \frac{P(\mathbf{x}, \mathbf{z})}{P(\mathbf{x})} = \frac{e^{-E(\mathbf{x}, \mathbf{z})}}{\sum_{\mathbf{z}'} e^{-E(\mathbf{x}, \mathbf{z}')}}$$

$$E(\mathbf{x}, \mathbf{z}) = - \sum_{j=1}^p a_j x_j - \sum_{h=1}^H b_h z_h - \sum_{j=1}^p \sum_{h=1}^H x_j w_{jh} z_h$$

Gradient of the negative log likelihood:

$$P(\mathbf{x}) = \sum_{\mathbf{z}} P(\mathbf{x}, \mathbf{z})$$

$$-\log P(\mathbf{x}) = \log Z - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})} = \log \sum_{\mathbf{x}', \mathbf{z}} e^{-E(\mathbf{x}', \mathbf{z})} - \log \sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$\frac{\partial -\log P(\mathbf{x})}{\partial \theta} = - \frac{1}{\sum_{\mathbf{x}, \mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{x}', \mathbf{z}} \frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}', \mathbf{z})} + \frac{1}{\sum_{\mathbf{z}} e^{-E(\mathbf{x}, \mathbf{z})}} \sum_{\mathbf{z}} \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta} e^{-E(\mathbf{x}, \mathbf{z})}$$

$$= - \sum_{\mathbf{x}', \mathbf{z}} P(\mathbf{x}', \mathbf{z}) \frac{\partial E(\mathbf{x}', \mathbf{z})}{\partial \theta} + \sum_{\mathbf{z}} P(\mathbf{z}|\mathbf{x}) \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta}$$

approximation: replace

expectation with a single sample!

$$\approx \frac{\partial E(\mathbf{x}, \mathbf{z})}{\partial \theta} - \frac{\partial E(\mathbf{x}', \mathbf{z}')}{\partial \theta}$$

Gibbs sampling

Restricted Boltzmann Machines

- Training procedure: **Contrastive Divergence**

For a training sample x^i

- Compute $P(z|x^i)$
 - Sample a hidden activation vector z^i
 - **positive gradient** = $(x_j^i z_h^i)_{j,h}$
 - Compute $P(x|z^i)$
 - Sample a reconstruction vector x'
 - Compute $P(z|x')$ and sample a hidden activation vector z'
 - **negative gradient** = $(x'_j z'_h)_{j,h}$
 - update weights: $w_{jh} \leftarrow w_{jh} + \eta(x_j^i z_h^i - x'_j z'_h)$
- $$a_j \leftarrow a_j - \eta(x_j^i - x'_j) \qquad b_h \leftarrow b_h - \eta(z_h^i - z'_h)$$

Deep Belief Networks

- Stack multiple layers of RBM

G. E. Hinton & R. R. Salakhutdinov. *Reducing the dimensionality of data with neural networks*. (2006).

Neural network magic: How to train your (feed-forward) neural network



Architecture

- Start with one hidden layer.
- Stop adding layers when you **overfit**.
- Never use more weights than training samples.
- **Weight sharing:**

Different units have connections to different inputs but sharing the same weights.

E.g. Image analyses, looking for edges in different regions of space


Tuning the network size

- **Destructive:**


weight decay $\Delta w_j = -\eta \frac{\partial E}{\partial w_j} - \lambda w_j$ $E' = E + \frac{\lambda}{2} \sum_j w_j^2$

- **Constructive:**

Growing networks until satisfactory error rate is reached.



Dynamic node creation [Ash 1989]:
Add new hidden units





Cascade correlation
[Fahlman & Lebiere 1989]:
Add new hidden layers with one unit.

Overtraining

Number of weights: $H(p+1) + (H+1).K$

hidden units



Optimization algorithm

- **Batch learning:**

Update the weights after a complete pass over the training set.

- **Mini-batch learning:**

- Update the weights after a pass over a set of training points of fixed size.


- Use **(quasi-)Newton methods** for small number of weights.

Prefer Levenberg-Marquardt.

- Use **conjugate gradient descent** for large number of weights.

Preconditioning

- An **ill-conditioned** network cannot learn.
- The best **learning rate** is typically **different for each weight**.
- Hence
 - **Standardize** inputs and targets
 - **Initialize** weights carefully
 - **Local learning rates**
- Use **tanh** rather than a logistic sigmoid for hidden layers so as to avoid low coefficients of variation (stdev/mean).



Standardization

- Remove outliers
- Features
 - Mean 0, standard deviation 1
 - Midrange 0, range 2
 - Orthonormalize (SVD, PCs...)
- Targets
 - Mean 0, standard deviation 1
 - Midrange 0, range 2

Use lower/upper bounds rather than min/max

Escaping saturation

- Large weights \Rightarrow saturation
- **Weight initialization**

$$w_{ij} = [-r, r] \quad r = \frac{1}{\sqrt{|\mathcal{N}(i)|}}$$

- **Weight decay** \equiv regularization

$E \rightarrow E + \text{weight decay}$

$$E' = E + \frac{\lambda}{2} \sum_j w_j^2 \quad \Delta w_j = -\eta \frac{\partial E}{\partial w_j} - \lambda w_j$$

$$E' = E + \sum_j \frac{w_j^2}{w_j^2 + \text{Cte}}$$

Escaping local minima

- Online learning or mini-batch
- Momentum

$$w_j \leftarrow w_j + \eta \Delta w_j + m \Delta^{(t+1)} w_j$$

- Adaptive learning rate

$$w_j \leftarrow w_j + \mu_j \Delta w_j$$

- $\mu_j \nearrow$ while the gradient keeps pointing in the same direction



$$\mu_j \leftarrow \mu_j + q \Delta^{(t)} w_j \Delta^{(t-1)} w_j$$

- Prevent $\mu_j < 0$: apply to $\log(\mu_j)$ instead
- Approximate to avoid computing the exp and avoid too small values for μ_j

$$\mu_j \leftarrow \mu_j \times \max(0.5, 1 + q \Delta^{(t)} w_j \Delta^{(t-1)} w_j)$$

Dropout

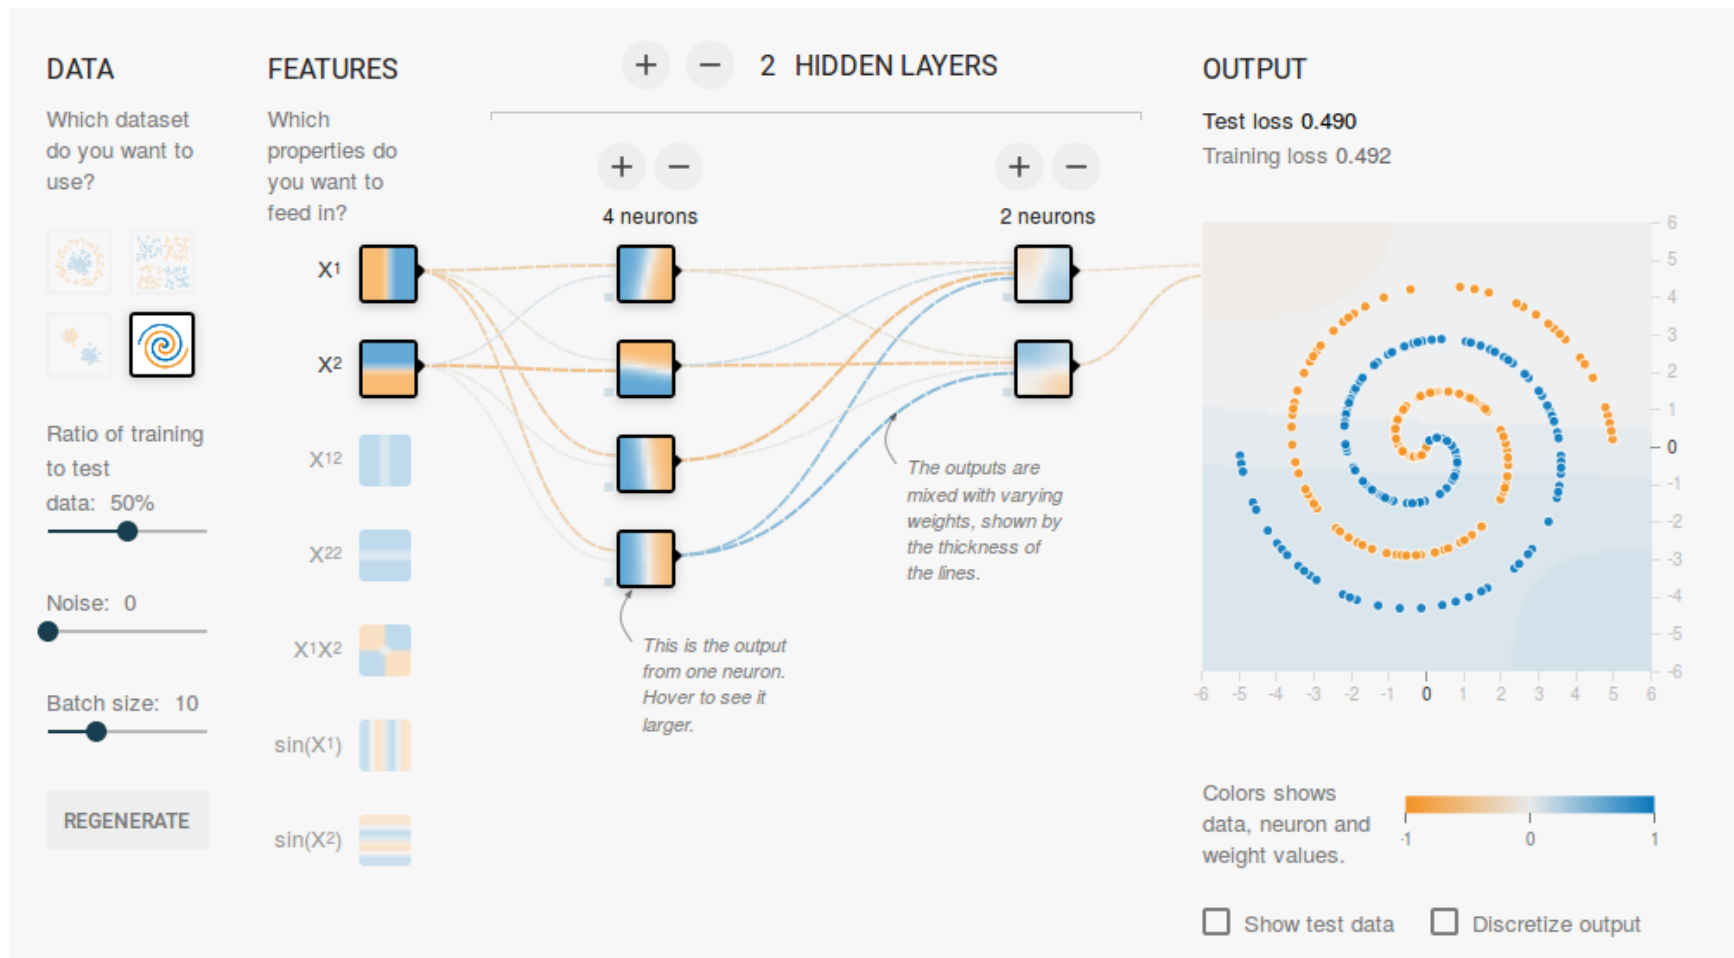
- At each iteration, **set half the units (randomly) to 0.**
- Avoid **overfitting**
- Helps focusing on **informative features**



(Srivastava, Hinton, Krizhevsky, Sutskever & Salakhutdinov 2012)

Playing with a neural network

<http://playground.tensorflow.org/>



Neural networks packages

http://deeplearning.net/software_links/

- Python

Theano, **TensorFlow**, Caffe, Keras...

- Java

Deeplearning4j, TensorFlow for Java

- Matlab

NeuralNetwork toolbox

- R

deepnet, H2O, MXNetR

References

- *A Course in Machine Learning*.
http://ciml.info/dl/v0_99/ciml-v0_99-all.pdf
 - **Perceptron**: Chap 4
 - **Multi-layer perceptron**: Chap 10.1 – 10.4
- Deep learning references
 - Le Cun, Y., Bengio, Y. and Hinton, G. (2015). **Deep learning**. *Nature* 521, 436-444.
 - <http://neuralnetworksanddeeplearning.com>
 - <http://deeplearning.net/>
- Playing with a (deep) neural network
 - <http://playground.tensorflow.org/>

Summary

- **Perceptrons** learn linear discriminants.
- Learning is done by **weight update**.
- **Multiple layer perceptrons** with one hidden unit are **universal approximators**.
- Learning is done by **backpropagation**.
- Neural networks are **hard to train**, caution must be applied.
- (Deep) neural networks can be very powerful!

Exam: Fri, Dec 22 8:30am–11:30am

- **No documents, no calculators, no computer.**
- **Theoretical, technical, and practical** questions
- **Short answers!**
- **How to study**
 - Homework + previous year's exams
 - Labs
 - Answer the questions on the slides
 - **Review session Dec. 15 (15:30-17:00):** ask your questions!
- **Formulas**
 - To know: **Bayes**, how to compute derivatives.
 - Everything else will be given. **Interpretation** is key.

kaggle challenge project



- Detailed instructions on the course website

http://cazencott.info/dotclear/public/lectures/ma2823_2017/kaggle-project.pdf

- Deadline for submissions & for the report:

Sat, Dec 23 at 23:59 <http://tinyurl.com/ma2823-2017-hw/>

- Report: Practical instructions:

- PDF document

- No more than 2 pages
- Project_<LastName1><Initial>_<Lastname2><Initial>_<Lastname3><Initial>.pdf

- Starts with

- Full names
- Kaggle user names
- Kaggle team name(s)

