



NoSQL databases

Elisabetta Di Nitto
elisabetta.dinitto@polimi.it

13/01/2015

Lecture for the course:
A Multidisciplinary Perspective On Big Data

Credits to Danilo Ardagna, Oscar Locatelli,
Santo Lombardo, Marco Scavuzzo



Databases: basic concepts of the relational model

DBMS



- A DBMS does the following
 - ▶ Allows an **efficient access to data**, with a granularity level lower than the one offered by a file system
 - ▶ Allows a **direct access to data** based on their properties
 - ▶ Allows **concurrency control** at the level of the single piece of data (record)
 - ▶ Offers languages to **describe data and queries** on data in a way that is independent on the way they are physically stored
 - ▶ Offers sophisticated mechanisms for
 - **Reliability**
 - **Privacy**
 - **Atomicity of transactions**

RDBMS: assumptions and benefits



- Assumptions
 - ▶ Well-defined structure for data
 - ▶ Data is dense and uniform
 - ▶ Indexes can be defined a priori and used for queries
 - ▶ Data stays within a few Gigabyte
- Benefits
 - ▶ It uses the lowest amount of disk space
 - ▶ It is a well-understood model and query language
 - ▶ It can support a wide variety of use cases
 - ▶ It has schema-enforced data consistency

RDBMS: disadvantages and new requirements



- Disadvantages
 - ▶ Relatively slow performance
 - ▶ Schemas mean a higher programmer overhead for iterating changes
 - ▶ High degree of complexity with many tuning knobs
- New requirements
 - ▶ In 2009 Google was processing 24 Petabyte per day
 - ▶ In 2009 Facebook declared to store about 60 millions of images
 - ▶ The Internet archive stores 2 Petabyte of data

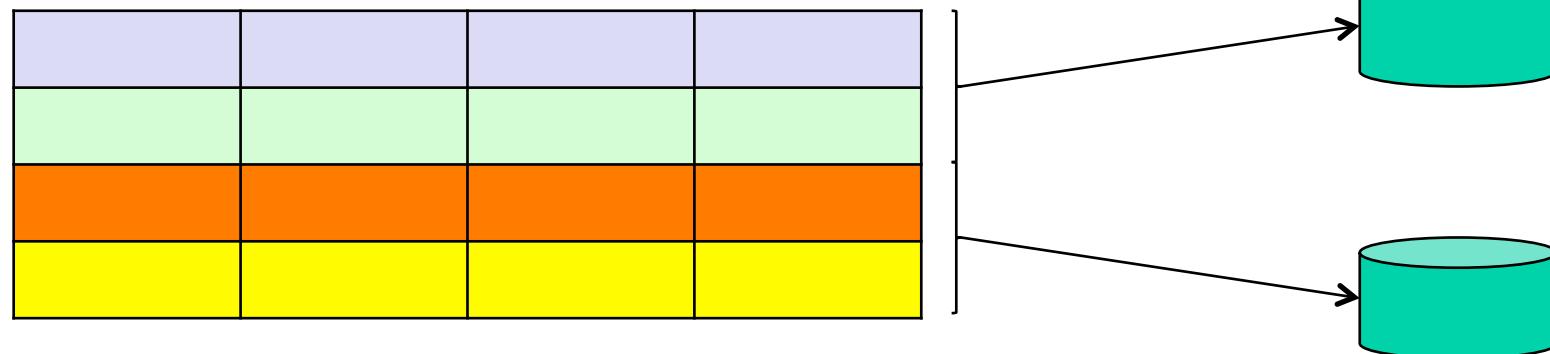


NoSQL: basic concepts

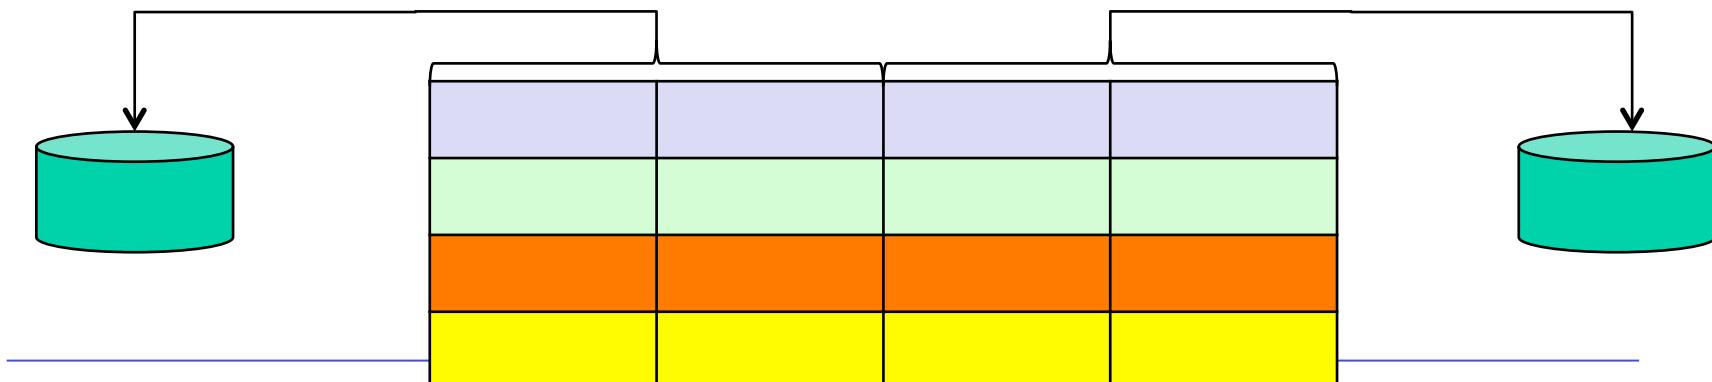
Data partitioning

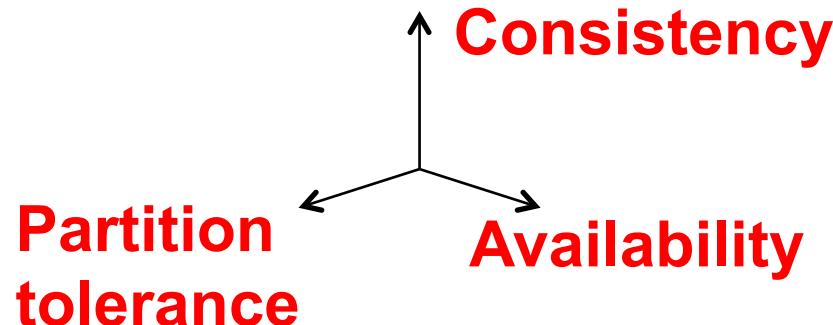


- In case of big data set: partition data storing them on different servers
- Horizontal partitioning (*sharding*)



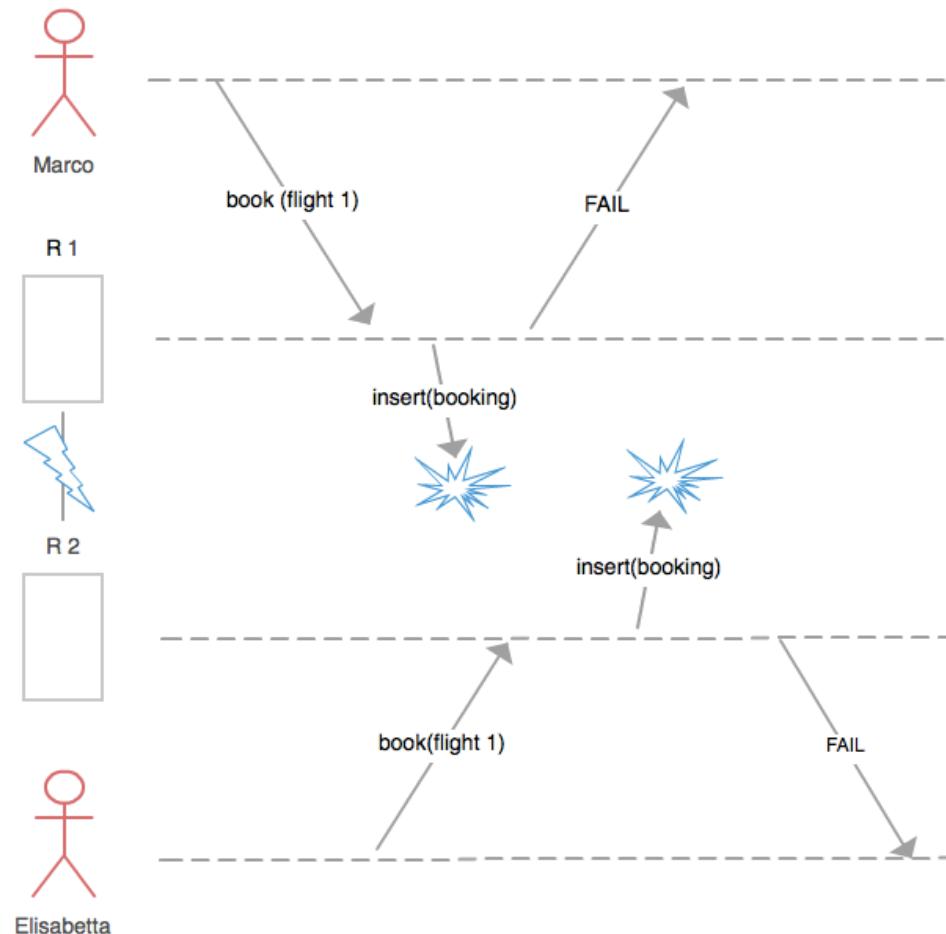
- Vertical partitioning





- Three possible guarantees
 - ▶ *Consistency*: all nodes see the same data at the same time
 - ▶ *Availability*: every request receives a response about whether it was successful or failed
 - ▶ *Partition tolerance*: the system continues to operate despite arbitrary message loss or failure of part of the system
- **CAP theorem**: it is impossible for a distributed computer system to simultaneously provide all three guarantees

CP (Consistency & Partition tolerance) = CAP-Consistency system example

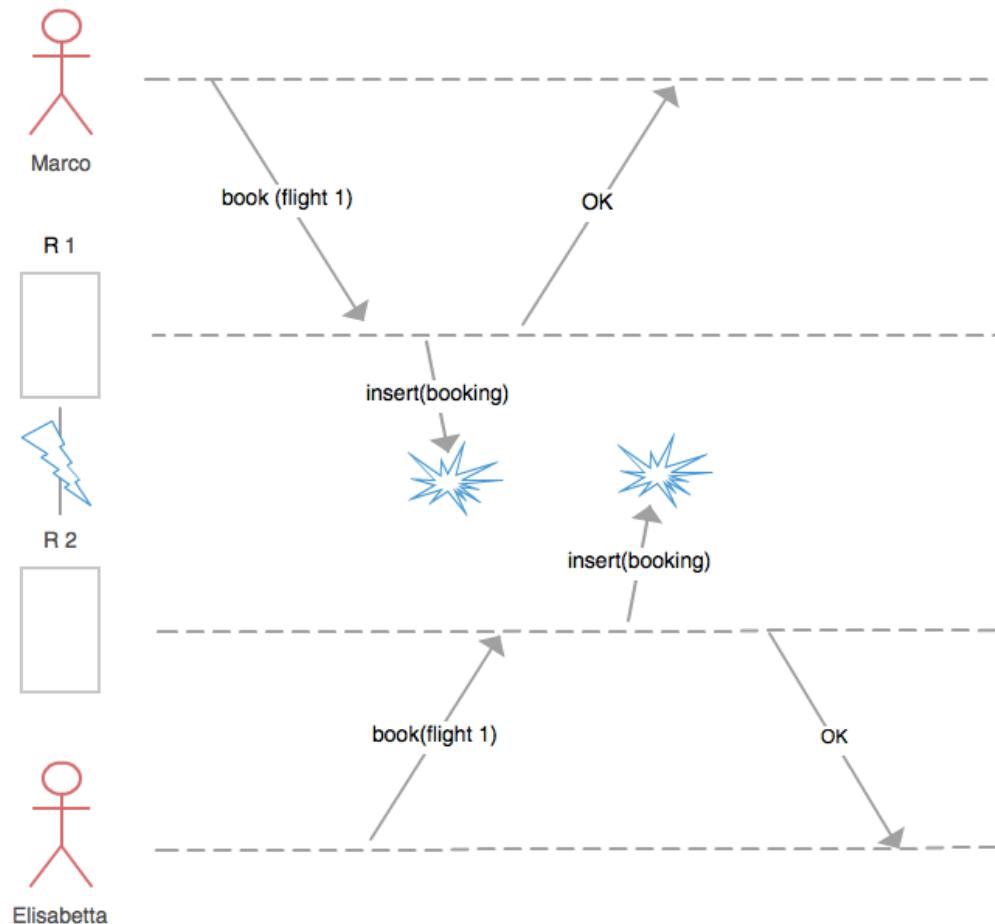


Non ACID Consistency Management



- Eventual Consistency:
 - ▶ Changes made at one replica will be transmitted asynchronously to the others (e.g., DNS)
 - ▶ Discrepancies in data state between replicas, and thus between users and locations, for a temporary period may occur

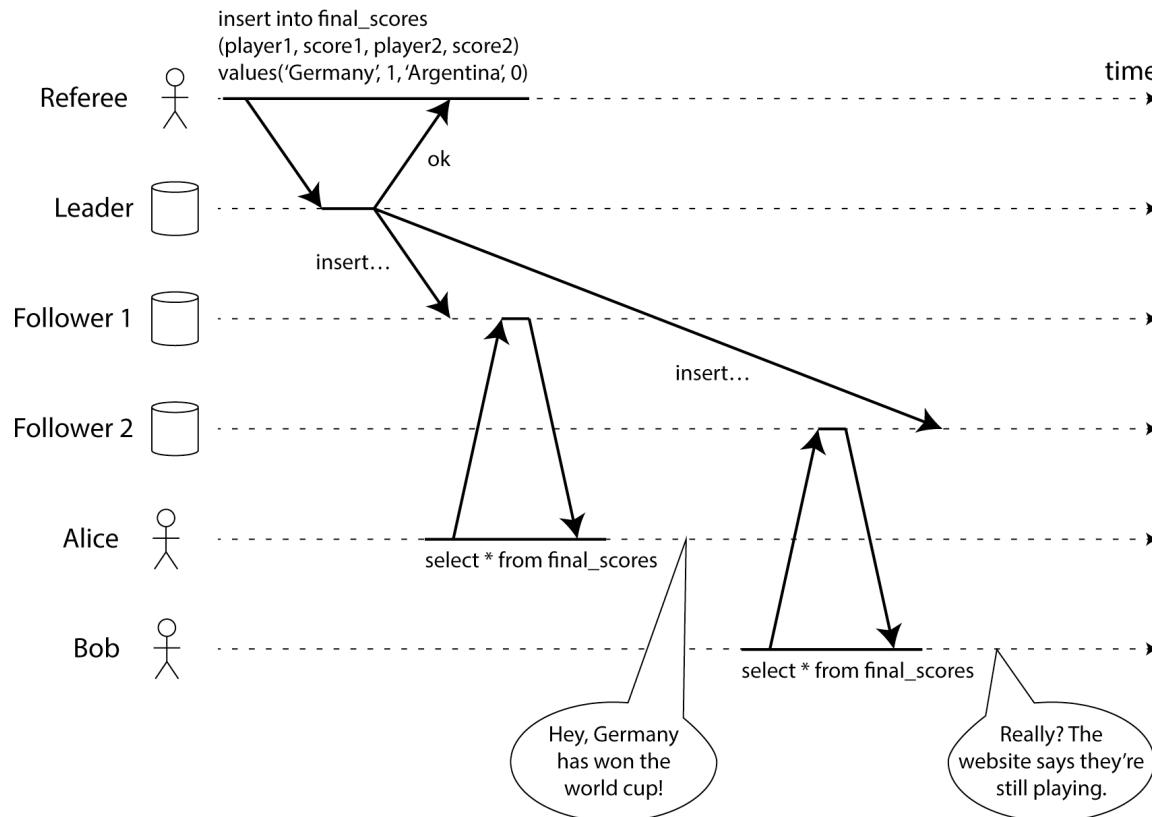
AP (Availability & Partition tolerance) = CAP-Availability system example



Tradeoff Consistency/Latency



- Example of a database executing non linearizable write ops. (without partitions)



Bob gets stale data

Bob realizes it through a separate communication channel (audio)

DBMS and CAP theorem



- Relational databases favor Consistency and Availability
- NoSQL databases favor Availability and Partition tolerance
 - ▶ Don't use them when you need consistency guarantees!

What is NoSQL?



- No use of SQL as query language:
 - ▶ Manage large volumes of data that do not necessarily follow a fixed schema
 - ▶ Data is partitioned among different machines and JOIN operations are not usable
 - ACID guarantees may be relaxed:
 - ▶ E.g., eventual consistency
 - ▶ Transactions limited to single data items
 - Distributed, fault-tolerant architecture
 - ▶ Data held in a redundant manner on several servers
 - ▶ Horizontal scalability
-



Categories of NoSQL

Main categories of NoSQL based on their logical data model



- Key-Value
 - Column-family
 - Document-based
-

Key Value NoSQL Databases



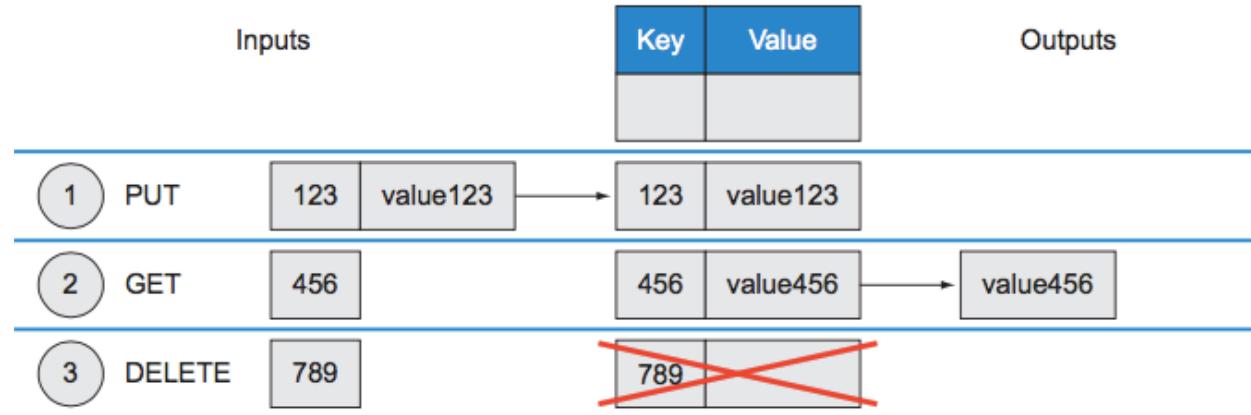
- Common data model: map/dictionary
- Putting and requesting values (BLOBs) is done per key
- Keys are sorted lexicographically
- Favor high scalability over consistency
- No ad-hoc querying or analytics features
- Length of keys is limited, length of values not limited



Key Value query language



- Query operations are limited to
 - ▶ put(key,value)
 - ▶ get(key)
 - ▶ delete(key)

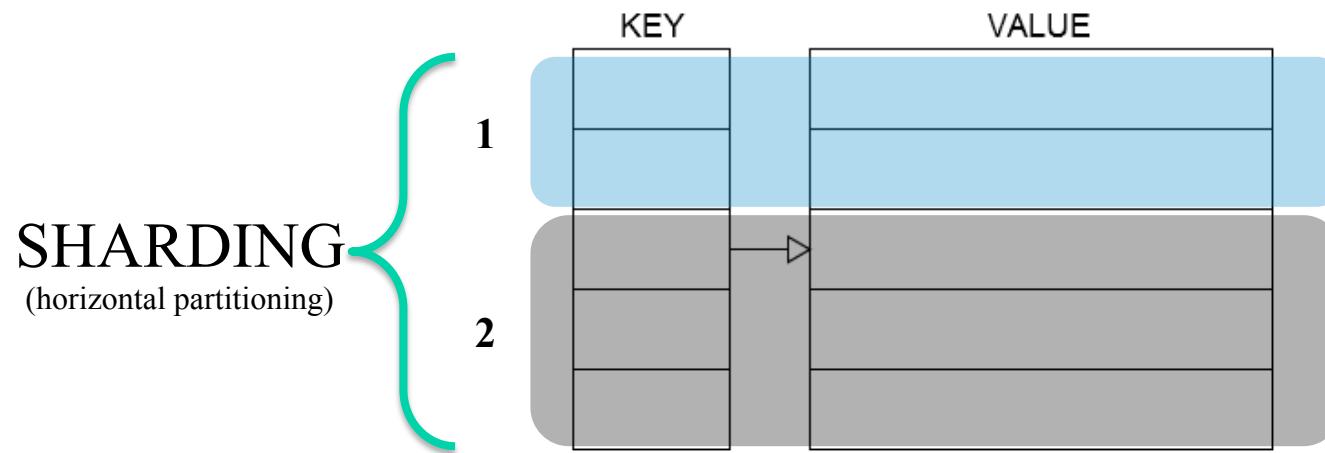


- Redis also provides matching for key-ranges
 - ▶ Matching of numeric ranges
 - ▶ Matching of regular expressions

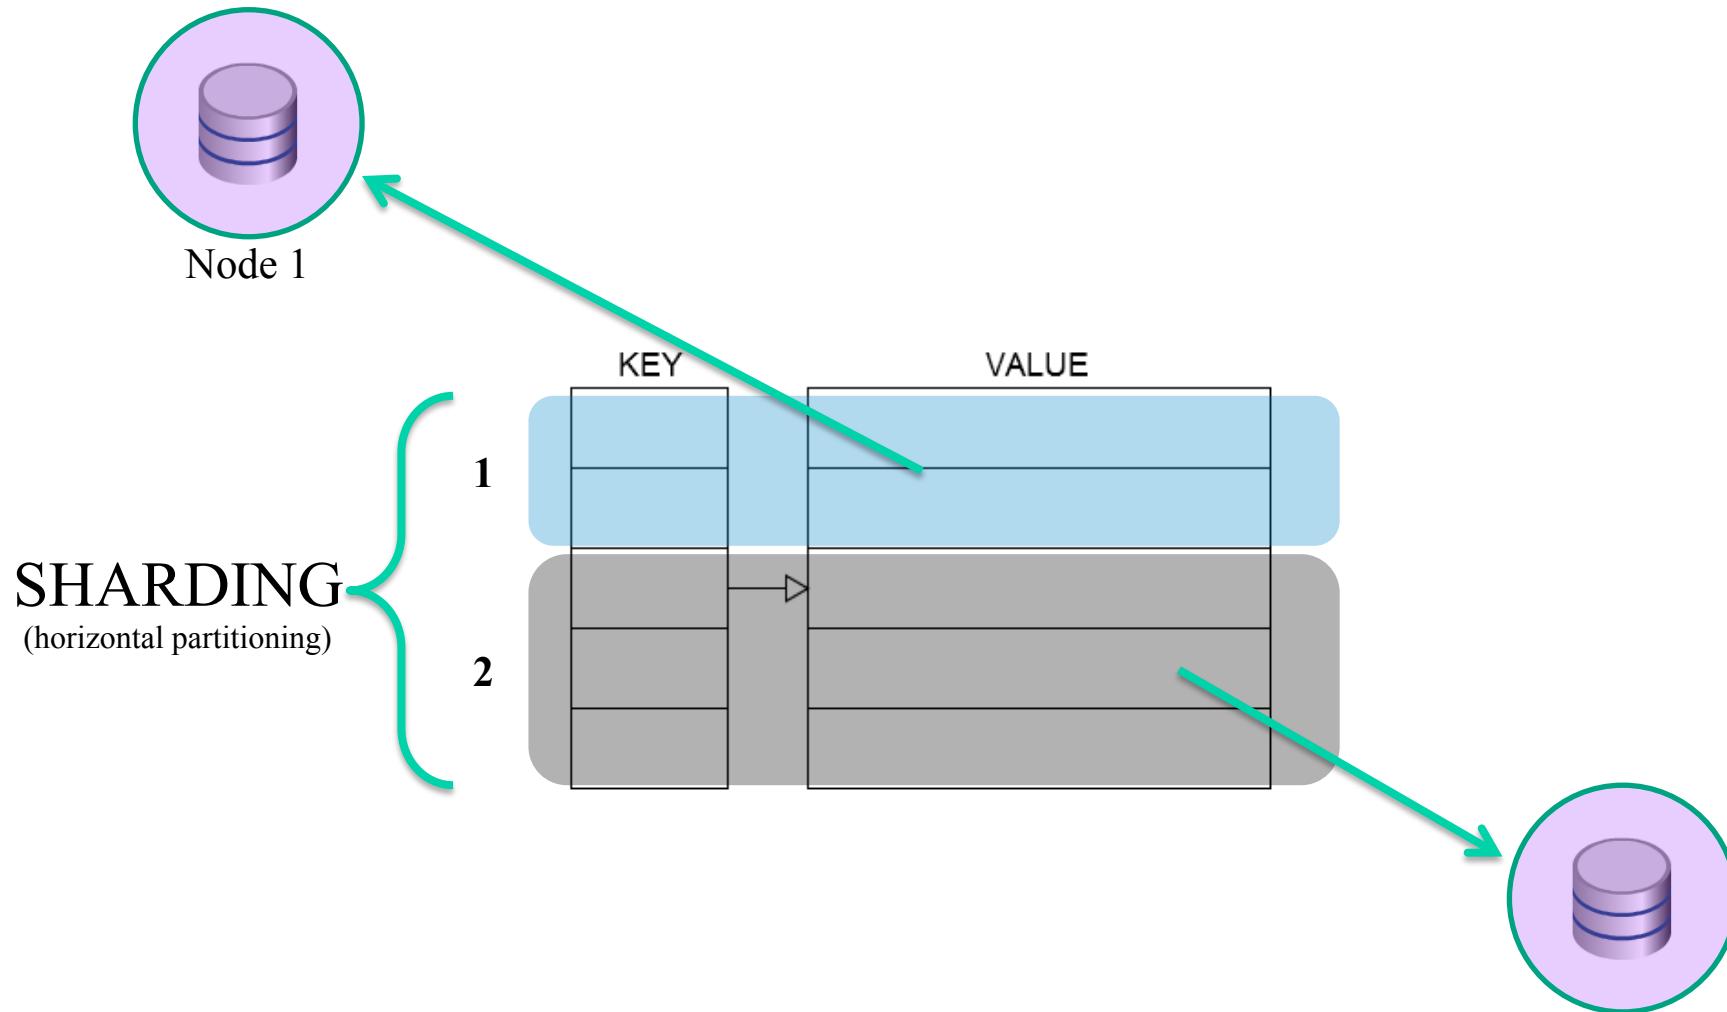
Key-Value and sharding



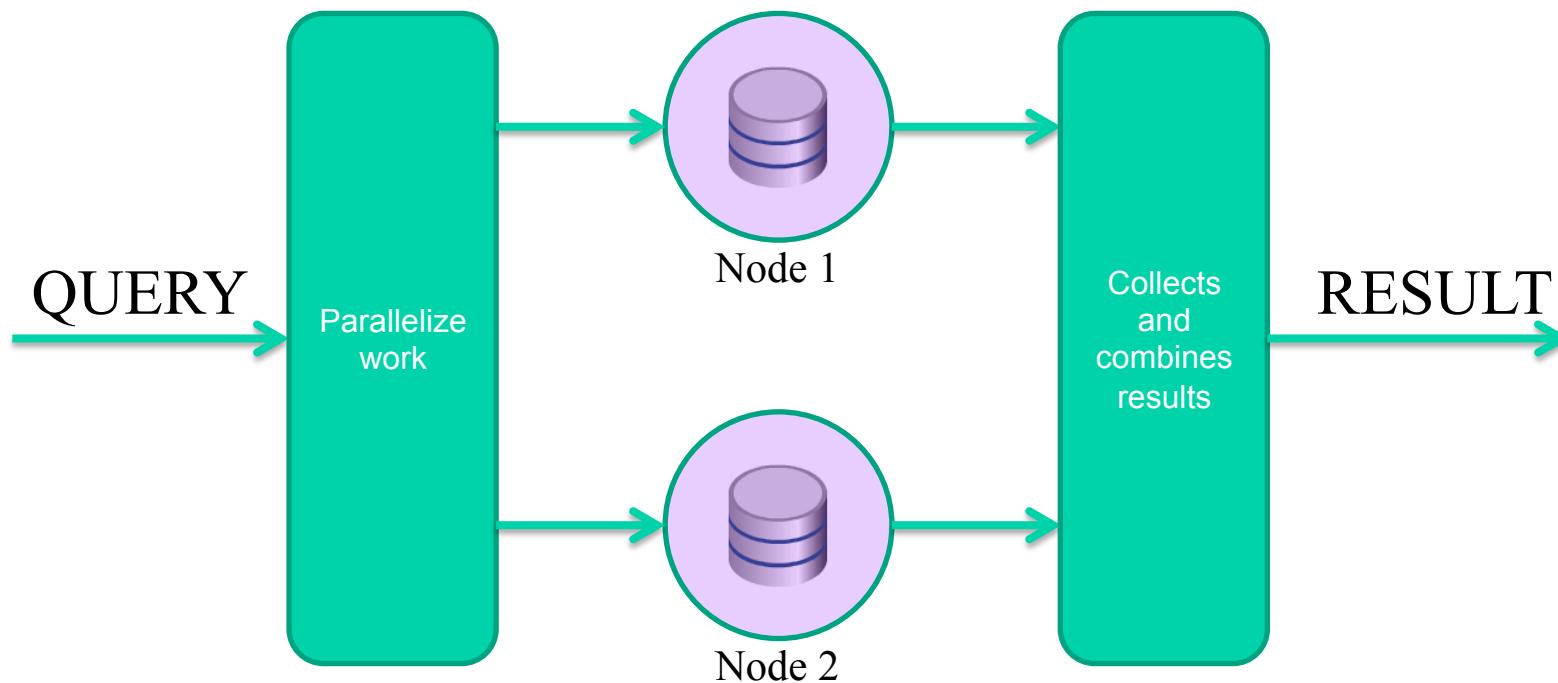
- **Shared Nothing** architecture: each key-value pair is decoupled from the others
- Can be easily moved on different machines



Key-Value and sharding

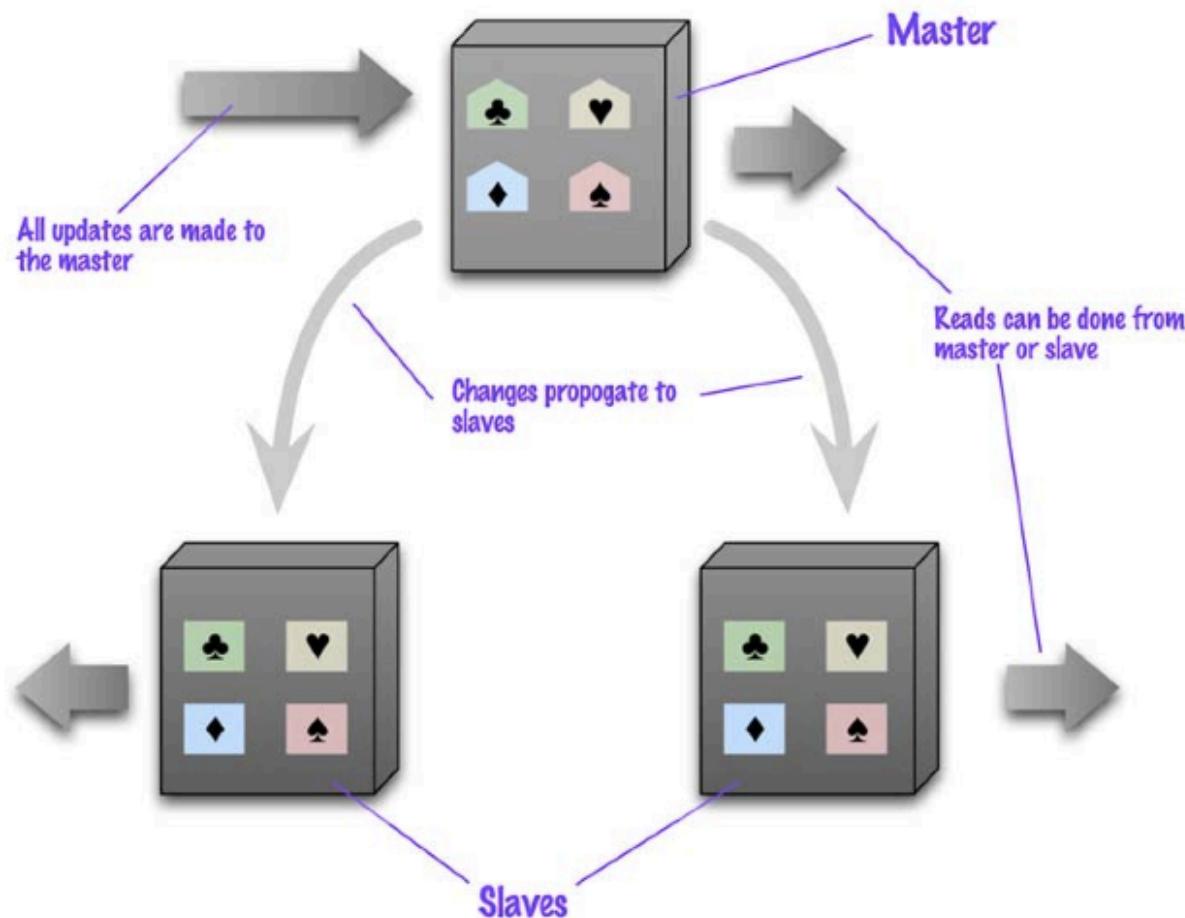


Approach to query execution



PARALLEL EXECUTION

Active-Passive Replication

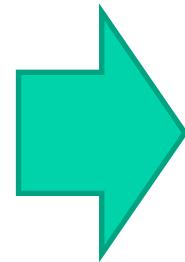
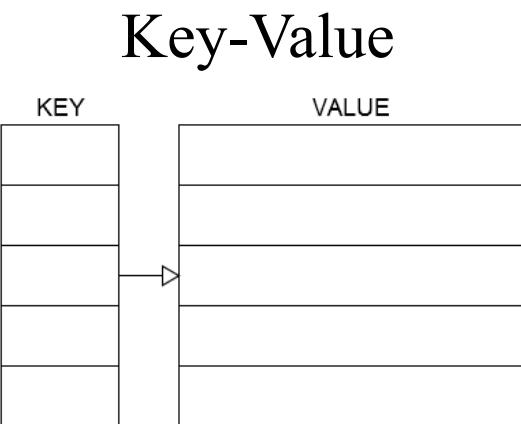


Key Value storage typical usages

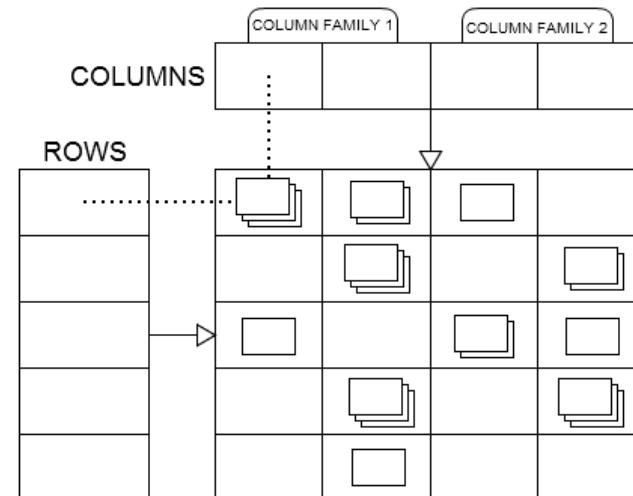


- Session storage
- User Profiles or preferences
- Shopping carts
- Single user analytics

Evolution of key-value



Column oriented



Document based

