

INTRODUCTION TO MEMCACHED

jurriaanpersyn.com
lead web dev at Netlog
since 4 years
php + mysql +
frontend
working on Gatcha

- From Wikipedia: "A cache is a collection of data duplicating original values stored elsewhere or computed earlier, where the original data is expensive to fetch (owing to longer access time) or to compute, compared to the cost of reading the cache."
- Term introduced by IBM in the 60's

- simple key/value storage
- simple operations
 - save
 - get
 - delete

- storage cost
- retrieval cost (network load / algorithm load)
- invalidation (keeping data up to date / removing irrelevant data)
- replacement policy (FIFO/LFU/LRU/MRU/RANDOM vs. Belady's algorithm)
- cold cache / warm cache

- cache hit and cache miss
- typical stats:
 - hit ratio (hits / hits + misses)
 - miss ratio (1 – hit ratio)
- 45 cache hits and 10 cache misses
 - $45 / (45 + 10) = 82\%$ hit ratio
 - 18% miss ratio

- caches are only efficient when the benefits of faster access outweigh the overhead of checking and keeping your cache up to date
- more cache hits than cache misses

Where are caches used?

NETLOG

- at hardware level (cpu, hdd)
- operating systems (ram)
- web stack
- applications
- your own short term vs long term memory

- As PHP backend developer, where to store cache results?
 - in database (computed values, generated html)
 - you'll still need to access your database
 - in static files (generated html or serialized php values)
 - you'll still need to access your file system



in memory!



memcached

- Free & open source, high-performance, distributed memory object caching system
- Generic in nature, intended for use in speeding up dynamic web applications by alleviating database load.
- key/value dictionary

- Developed by Brad Fitzpatrick for LiveJournal in 2003
- Now used by Netlog, Facebook, Flickr, Wikipedia, Twitter, YouTube ...

- It's a server
- Client access over TCP or UDP
- Servers can run in pools
 - eg. 3 servers with 64GB mem each give you a single pool of 192GB storage for caching
 - Servers are independent, clients manage the pool

What to store in memcache?

NETLOG

- high demand (used often)
- expensive (hard to compute)
- common (shared accross users)
- Best? All three

- Typical:
 - user sessions (often)
 - user data (often, shared)
 - homepage data (eg. often, shared, expensive)

- Workflow:
 - monitor application (query logs / profiling)
 - add a caching level
 - compare speed gain

- **Fast network access** (memcached servers close to other application servers)
- **No persistency** (if your server goes down, data in memcached is gone)
- **No redundancy / fail-over**
- **No replication** (single item in cache lives on one server only)
- **No authentication** (not in shared environments)

- 1 key is maximum 1MB
- keys are strings of 250 characters (in application typically MD5 of user readable string)
- No enumeration of keys (thus no list of valid keys in cache at certain moment, list of keys beginnen with “user_”, ...)
- No active clean-up (only clean up when more space needed, LRU)

```
$ telnet localhost 11211
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
get foo
VALUE foo 0 2
hi
END
stats
STAT pid 8861
(etc)
```


- both ASCII as Binary protocol
- in real life:
 - clients available for all major languages
 - C, C++, PHP, Python, Ruby, Java, Perl, Windows, ...

- Pages with high load / expensive to generate
- Very easy
- Very fast
- But: all the dependencies ...
 - language, css, template, logged in user's details, ...

- queries with JOIN and WHERE statements are harder to cache
- often not easy to find the cache key on update/change events
- solution: JOIN in PHP

- queries with JOIN and WHERE statements are harder to cache
 - often not easy to find the cache key on update/change events
 - solution: JOIN in PHP
-
- In following example: what if nickname of user changes?

- Pro's:
 - speed, duh.
 - queries get simpler (better for your db)
 - easier porting to key/value storage solutions
- Cons:
 - You're relying on memcached to be up and have good hit ratios

- We reduced database access
- Memcached is faster, but access to memcache still has it's price
- Solution: multiget
 - fetch multiple keys from memcached in one single call
 - result is array of items

- client is responsible for managing pool
 - hashes a certain key to a certain server
- clients can be naïve: distribute keys on size of pool
- if one server goes down, all keys will now be queried on other servers > cold cache
- use a client with consistent hashing algorithms, so if server goes down, only data on that server gets lost

jurriaan@netlog.com

Resources, a.o.:

- memcached & apc: <http://www.slideshare.net/benramsey/caching-with-memcached-and-apc>
- speed comparison: <http://dealnews.com/developers/memcachedv2.html>
- php client comparison: <http://code.google.com/p/memcached/wiki/PHPClientComparison>
- cakephp-memcached: <http://teknoid.wordpress.com/2009/06/17/send-your-database-on-vacation-by-using-cakephp-memcached/>
- caching basics: <http://www.slideshare.net/soplakanets/caching-basics>
- caching w php: <http://www.slideshare.net/JustinCarmony/effectice-caching-w-php-caching>