




dependable evolvable pervasive software engineering group

Big Data Technologies and Applications



Danilo Ardagna

Politecnico di Milano

danilo.ardagna@polimi.it

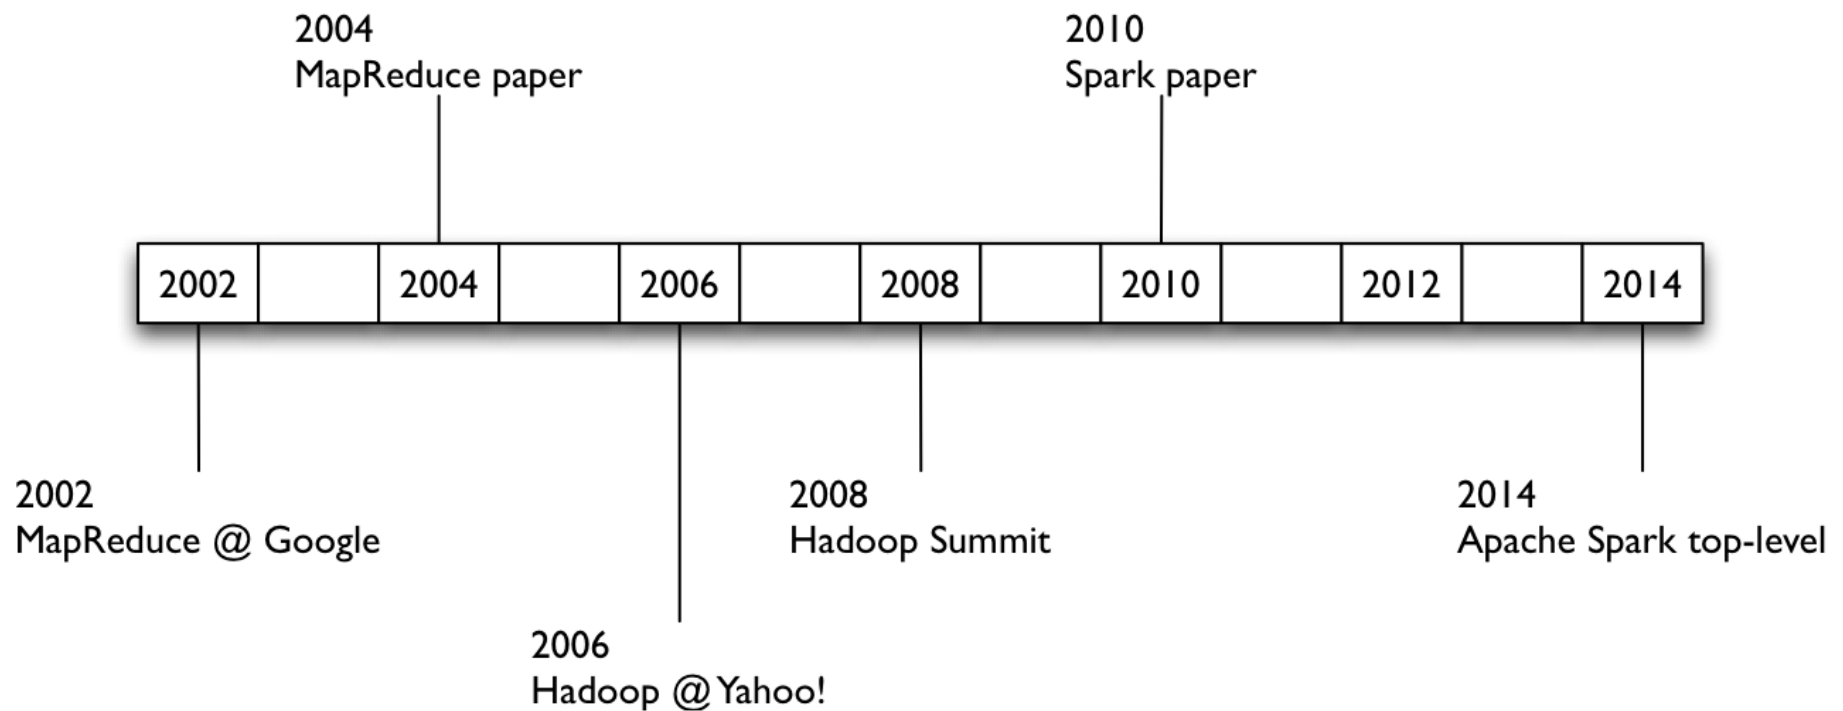


POLITECNICO
DI MILANO


Content

- Big Data Analytics and Data Science
- **Map-reduce and Hadoop fundamentals**
- Map-reduce and Hadoop pros and cons
- Hadoop Success Stories
- Hadoop Evolution and Hadoop eco-system
- Pig and Hive
- Map Reduce Cloud based solutions
- Where is the world going? Spark Apache project

A Brief History




Map-reduce



- Map-reduce: A general **algorithm**, and is prevalent in **functional programming** languages, which supports the notion of **map** and **reduce functions**
- MapReduce: The **patented software framework** from **Google** that the company applies in the realm of managing large datasets over clusters or other distributed topologies
- Hadoop: **Apache project open source** implementation of the MapReduce framework



Functional programming

- Computation as **evaluation of mathematical functions avoiding changing-state and mutable data**
 - functions are **expressions** and running a program means evaluating such expressions to get a **value**
- Mathematical functions have no side effects
 - output value depends only on input arguments
 - calling a ***f*** twice with the same value for an argument ***x*** will produce the same result ***f(x)***
 - eliminating side effects, i.e. changes in state that do not depend on the function inputs, can make much easier to understand and predict the behavior of a program
- Examples: **Scala**, Clojure, Haskell, Lisp



Function view


- Input: a **set** of **key/value** pairs
- User supplies two functions:
 - $\text{map}(k,v) \rightarrow \text{list}(k1,v1)$
 - $\text{reduce}(k1, \text{list}(v1)) \rightarrow v2$
- $(k1,v1)$ is an **intermediate** key/value pair
- **Output** is the **set** of $(k1,v2)$ pairs



MAPPING

SHUFFLING


REDUCING




@manudellavalle

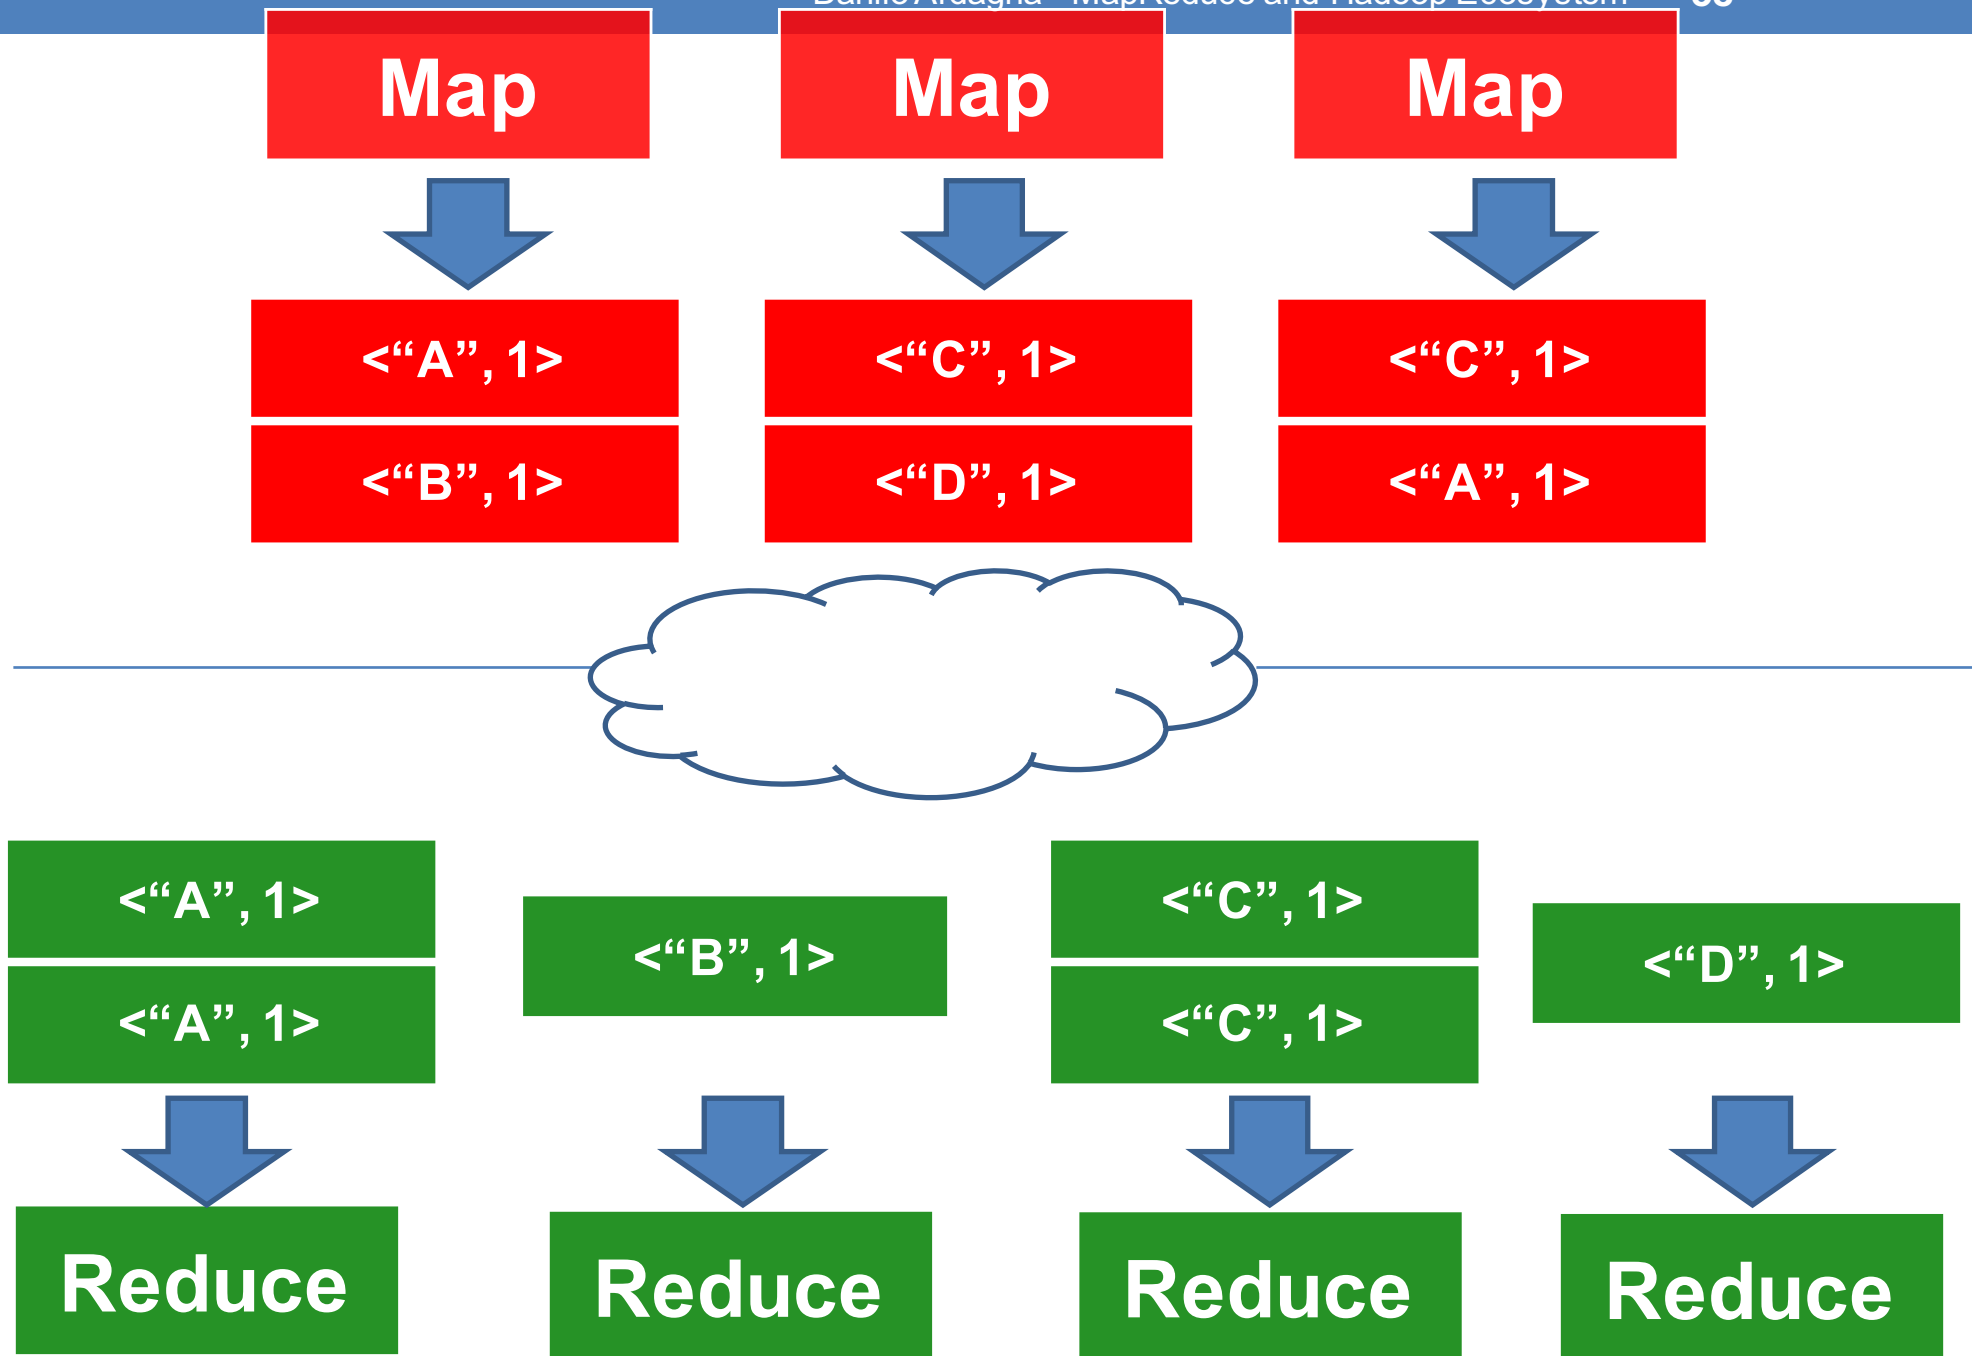
An example: Word Count

- We have many large files of words
- **Count** the number of times each **distinct word** appears in the files
- *Sample application:* analyze web server logs to find **popular URLs**





```
let map(String document_name, String document_content)=  
foreach Word word in document_content :  
emit(word, 1)
```



Intermediate Values



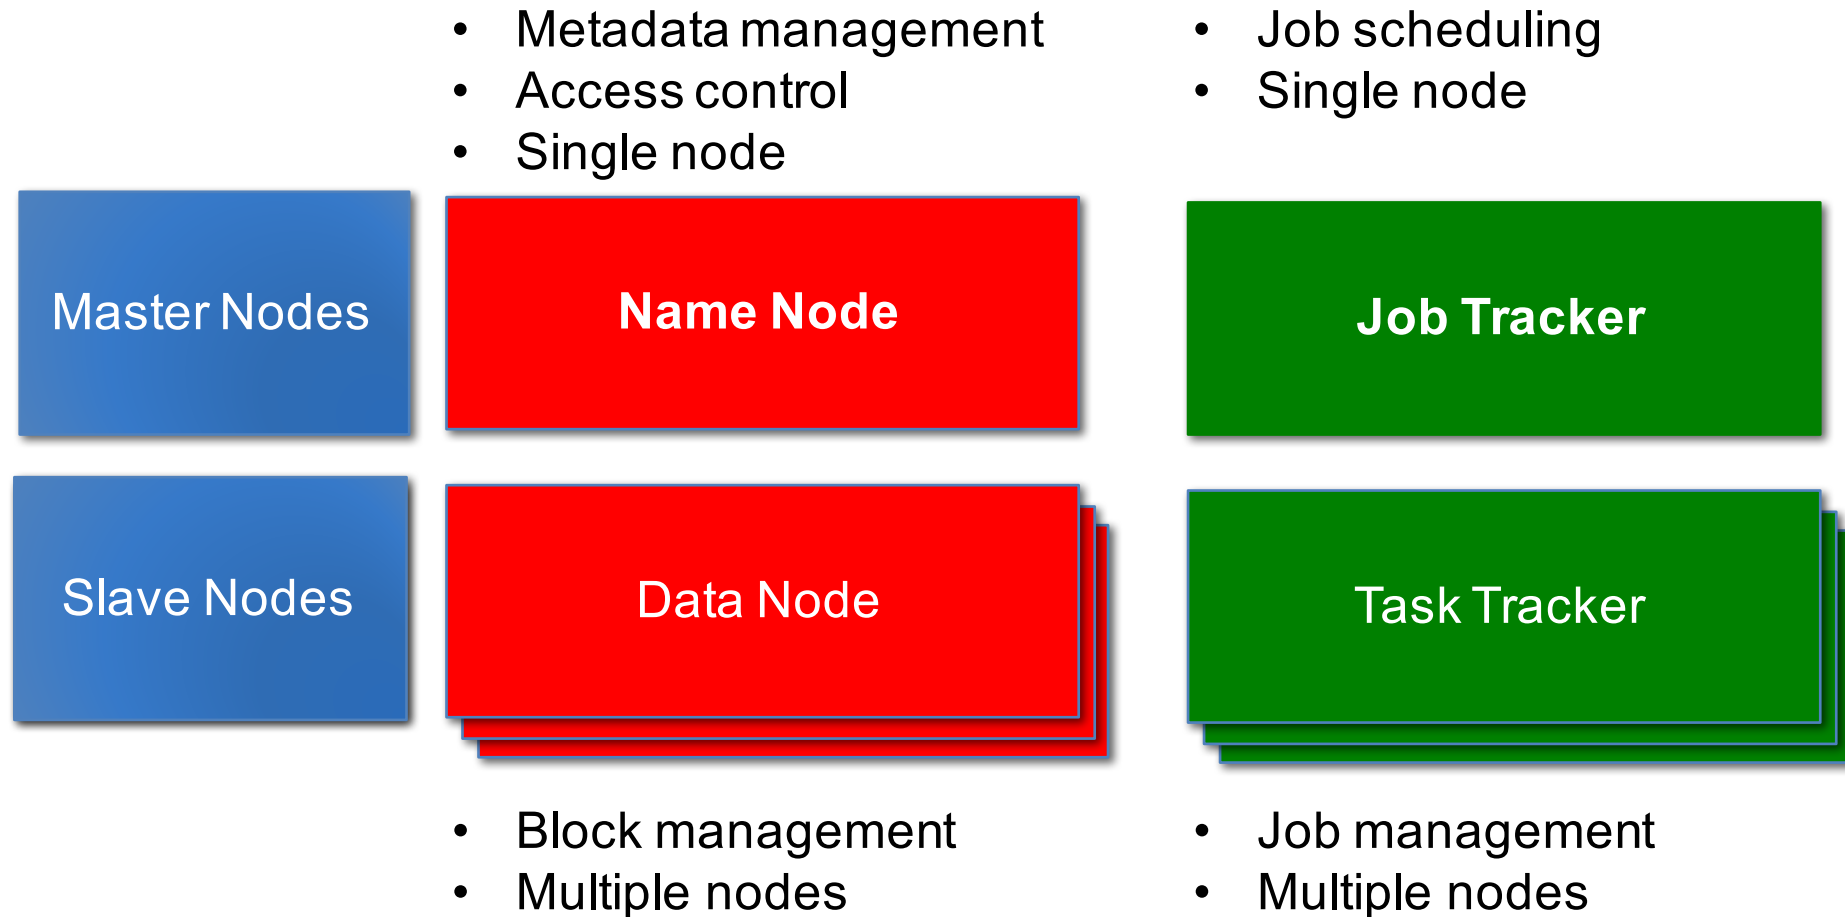
```
let reduce(Word word, Iterator<int> occurrences) =  
  int total_occurrences = 0;  
  foreach int o in occurrences : total_occurrences += o;  
  emit(word, total_occurrences);
```

Motivations for MapReduce

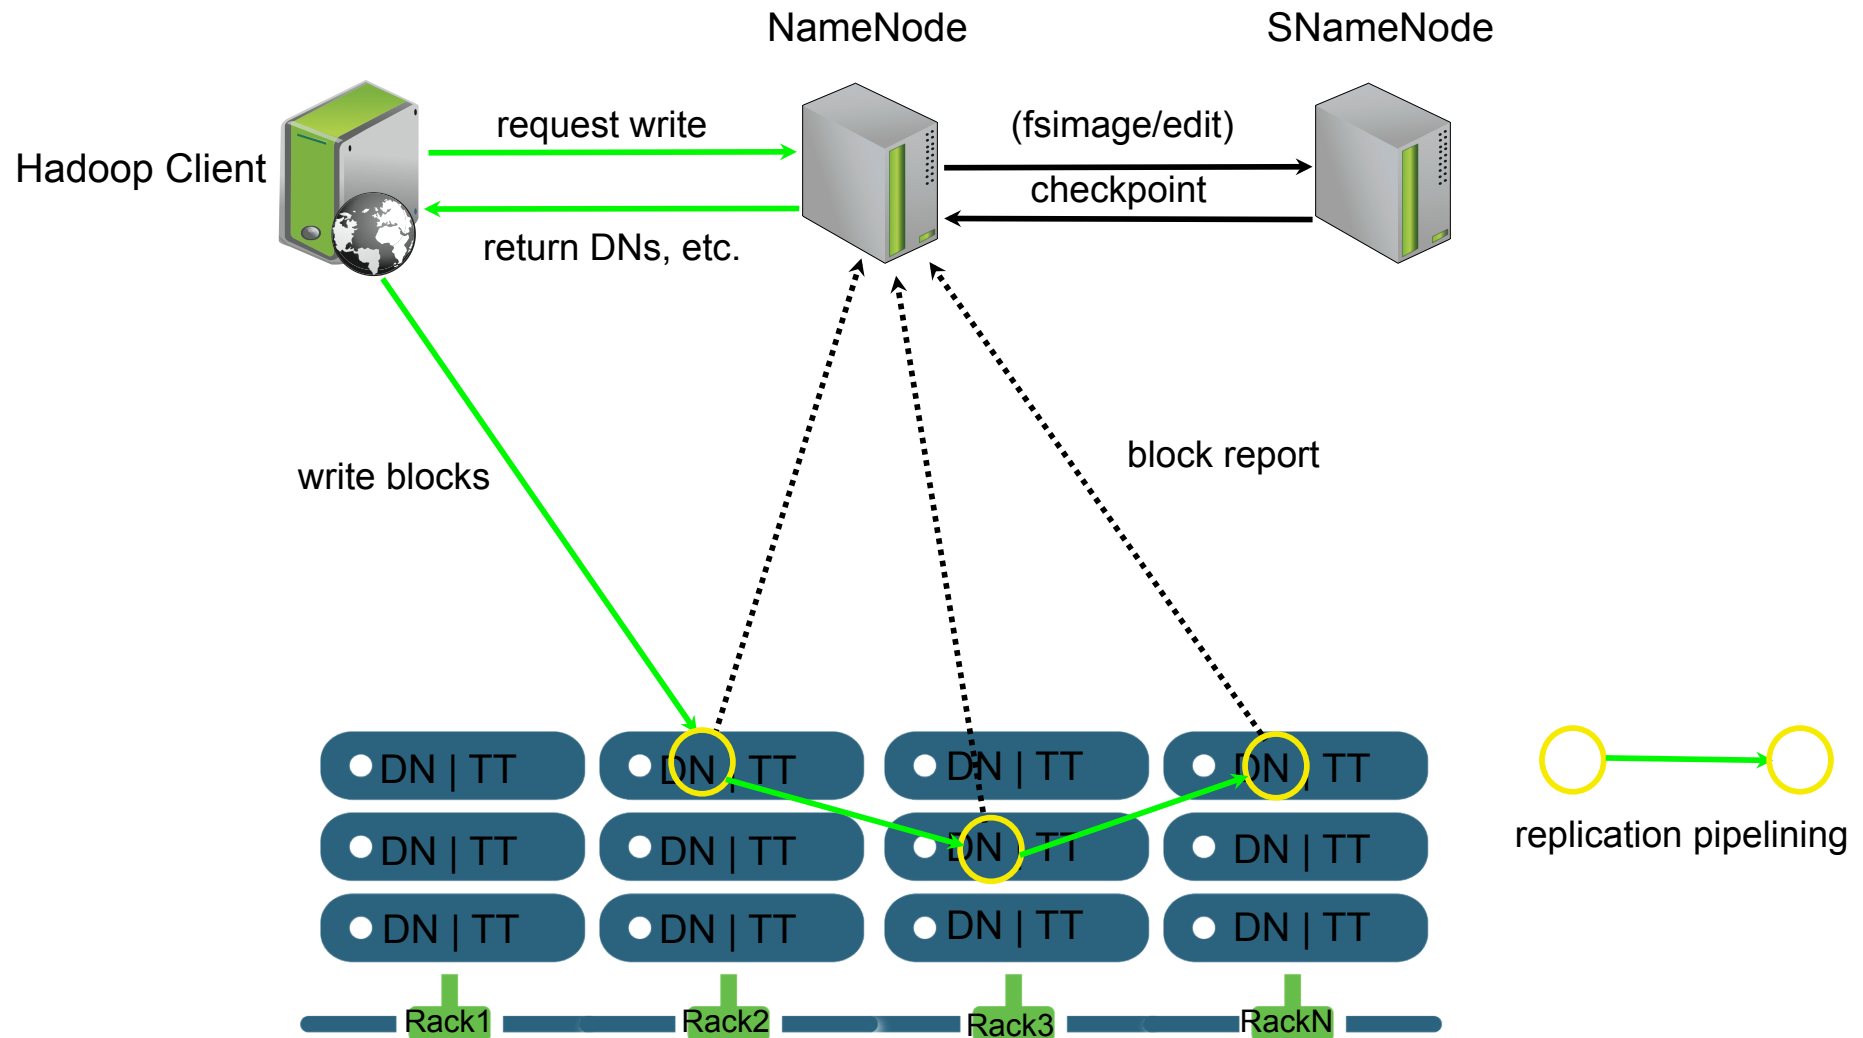
- Large-Scale Data Processing:
 - Want to use 1000s of CPUs
 - But don't want hassle of *managing* things
- Map-reduce Architecture provides:
 - **Automatic** parallelization & distribution
 - **Fault tolerance**
 - **I/O scheduling**
 - **Monitoring** & status updates



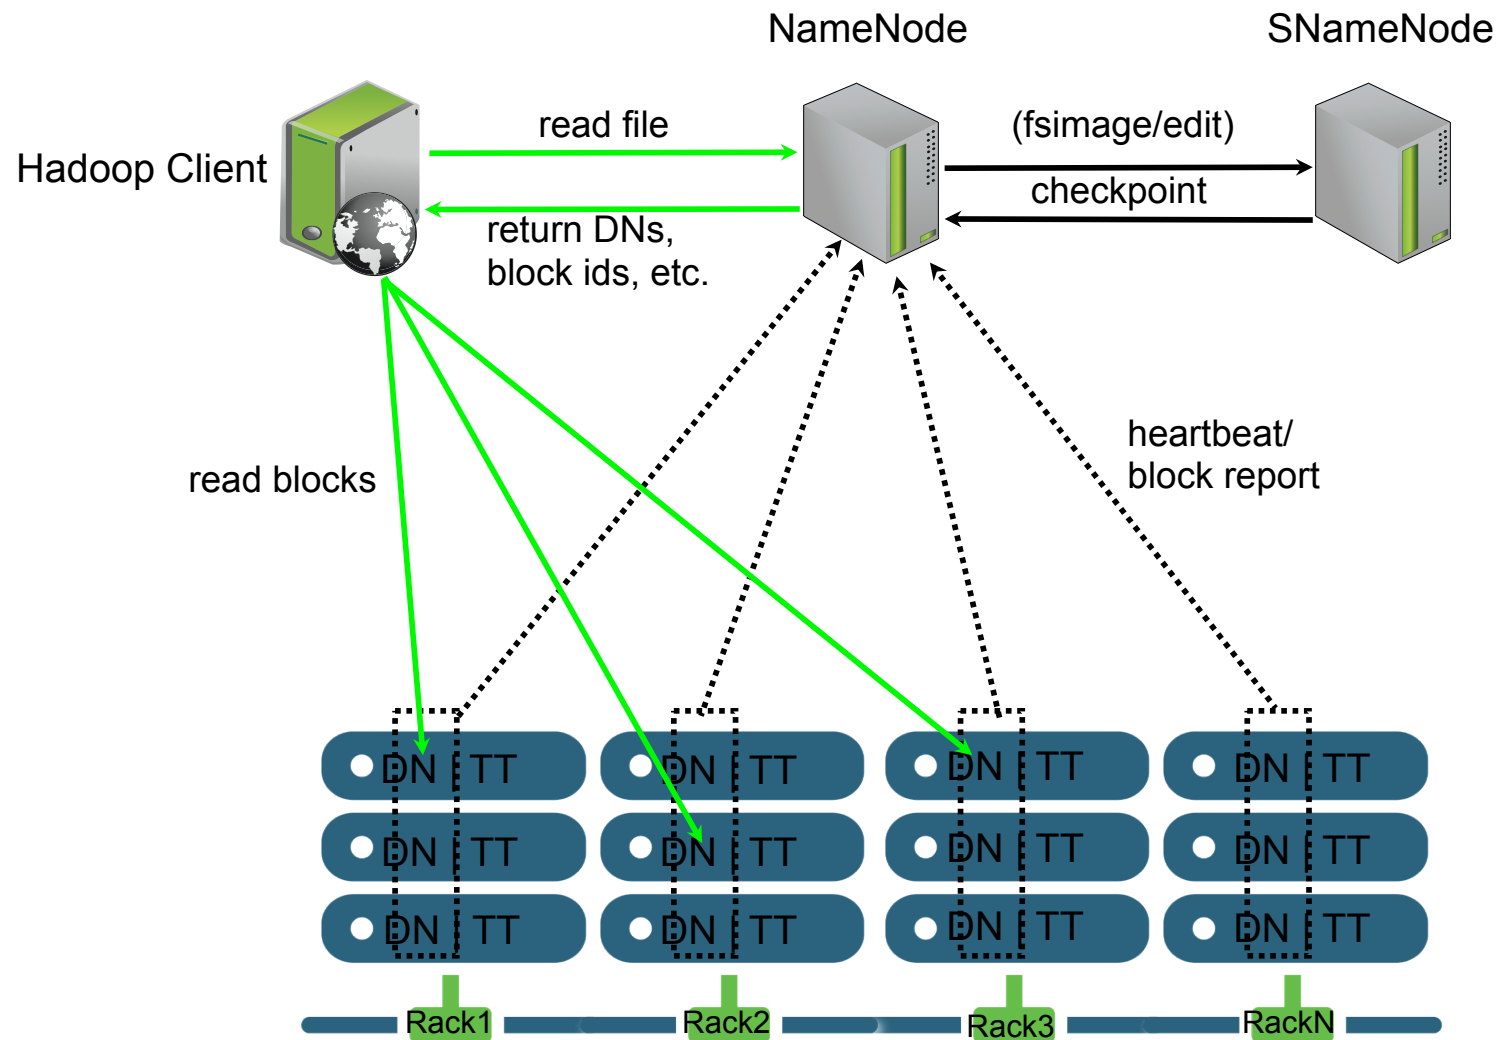
System view: Hadoop 1.0 implementation



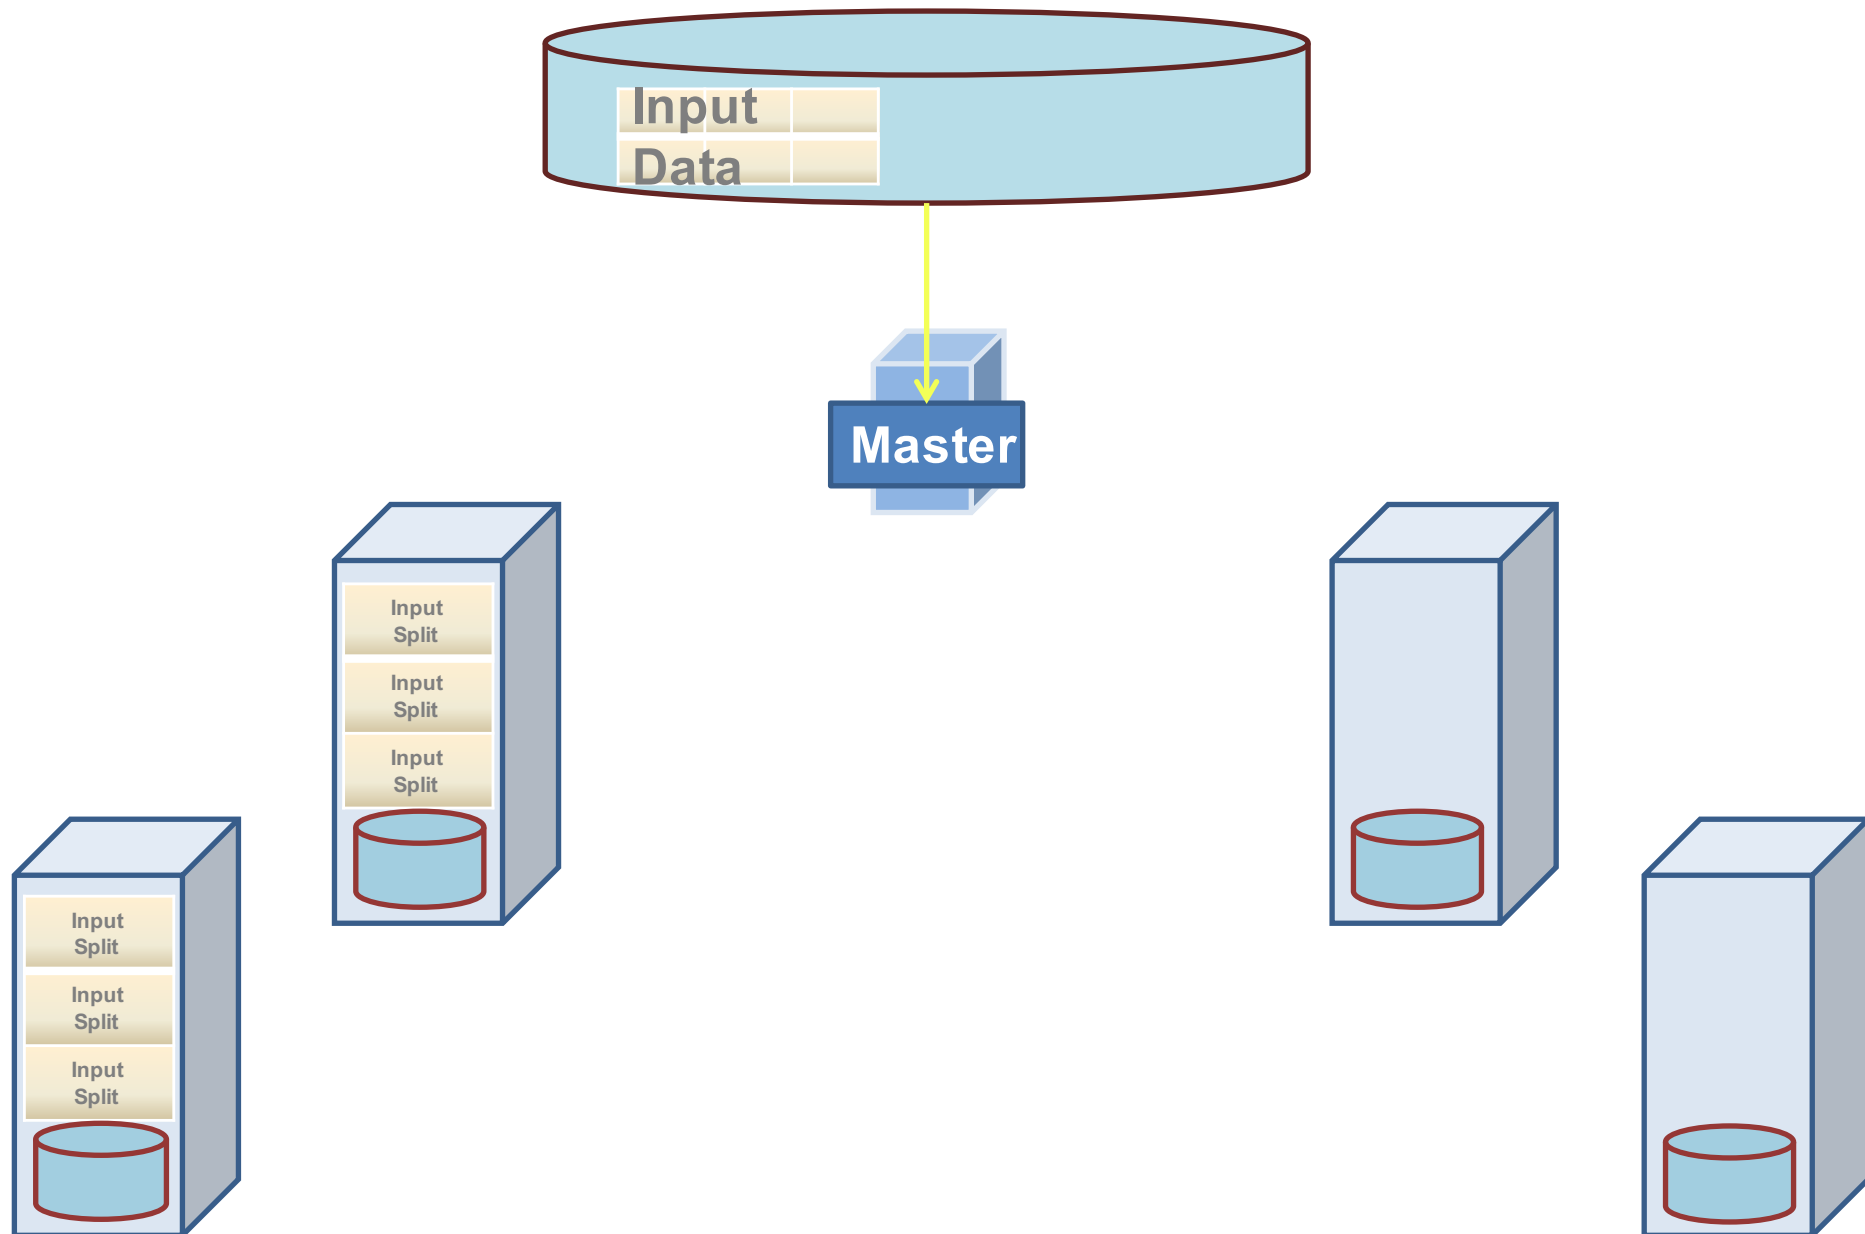
HDFS – Writing files



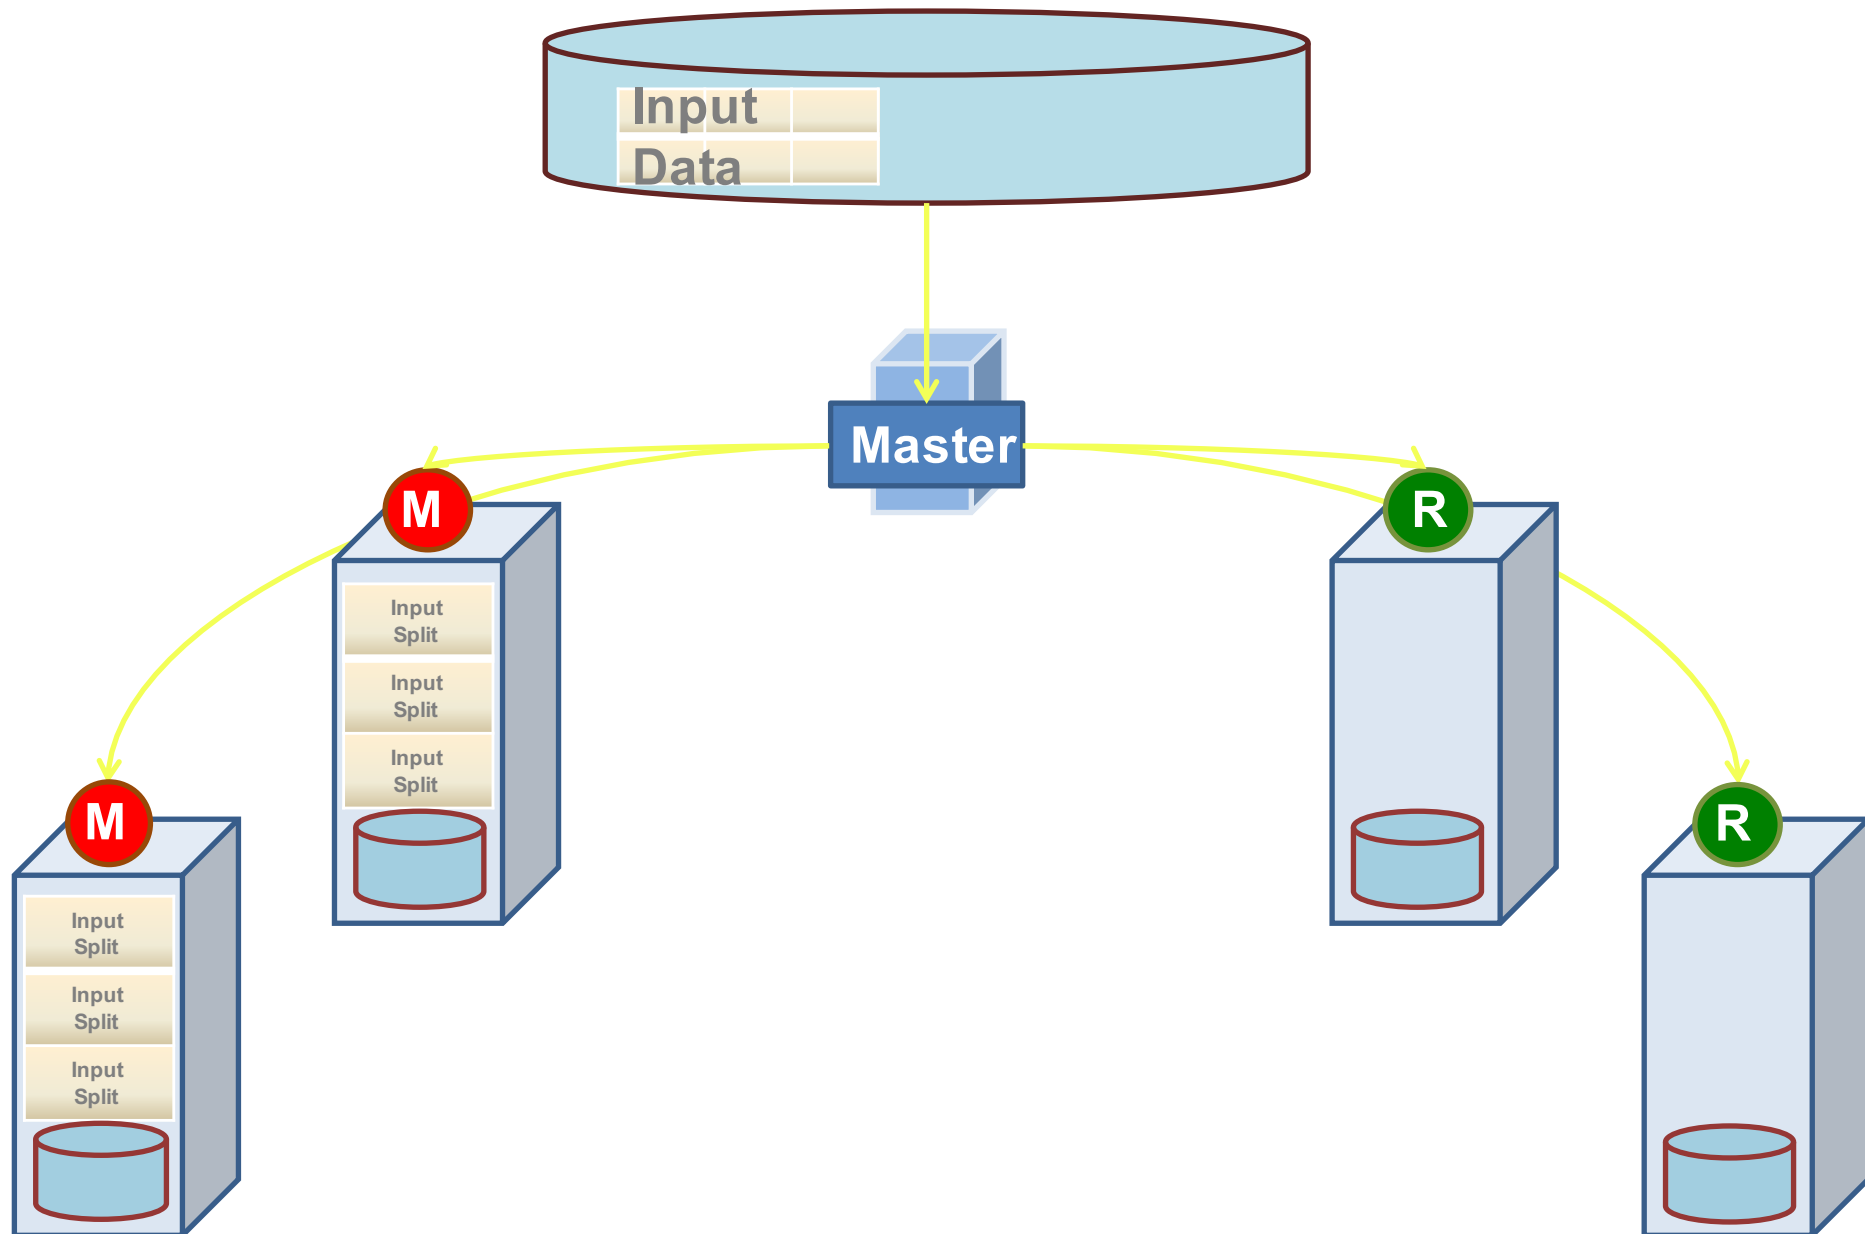
HDFS – Reading files



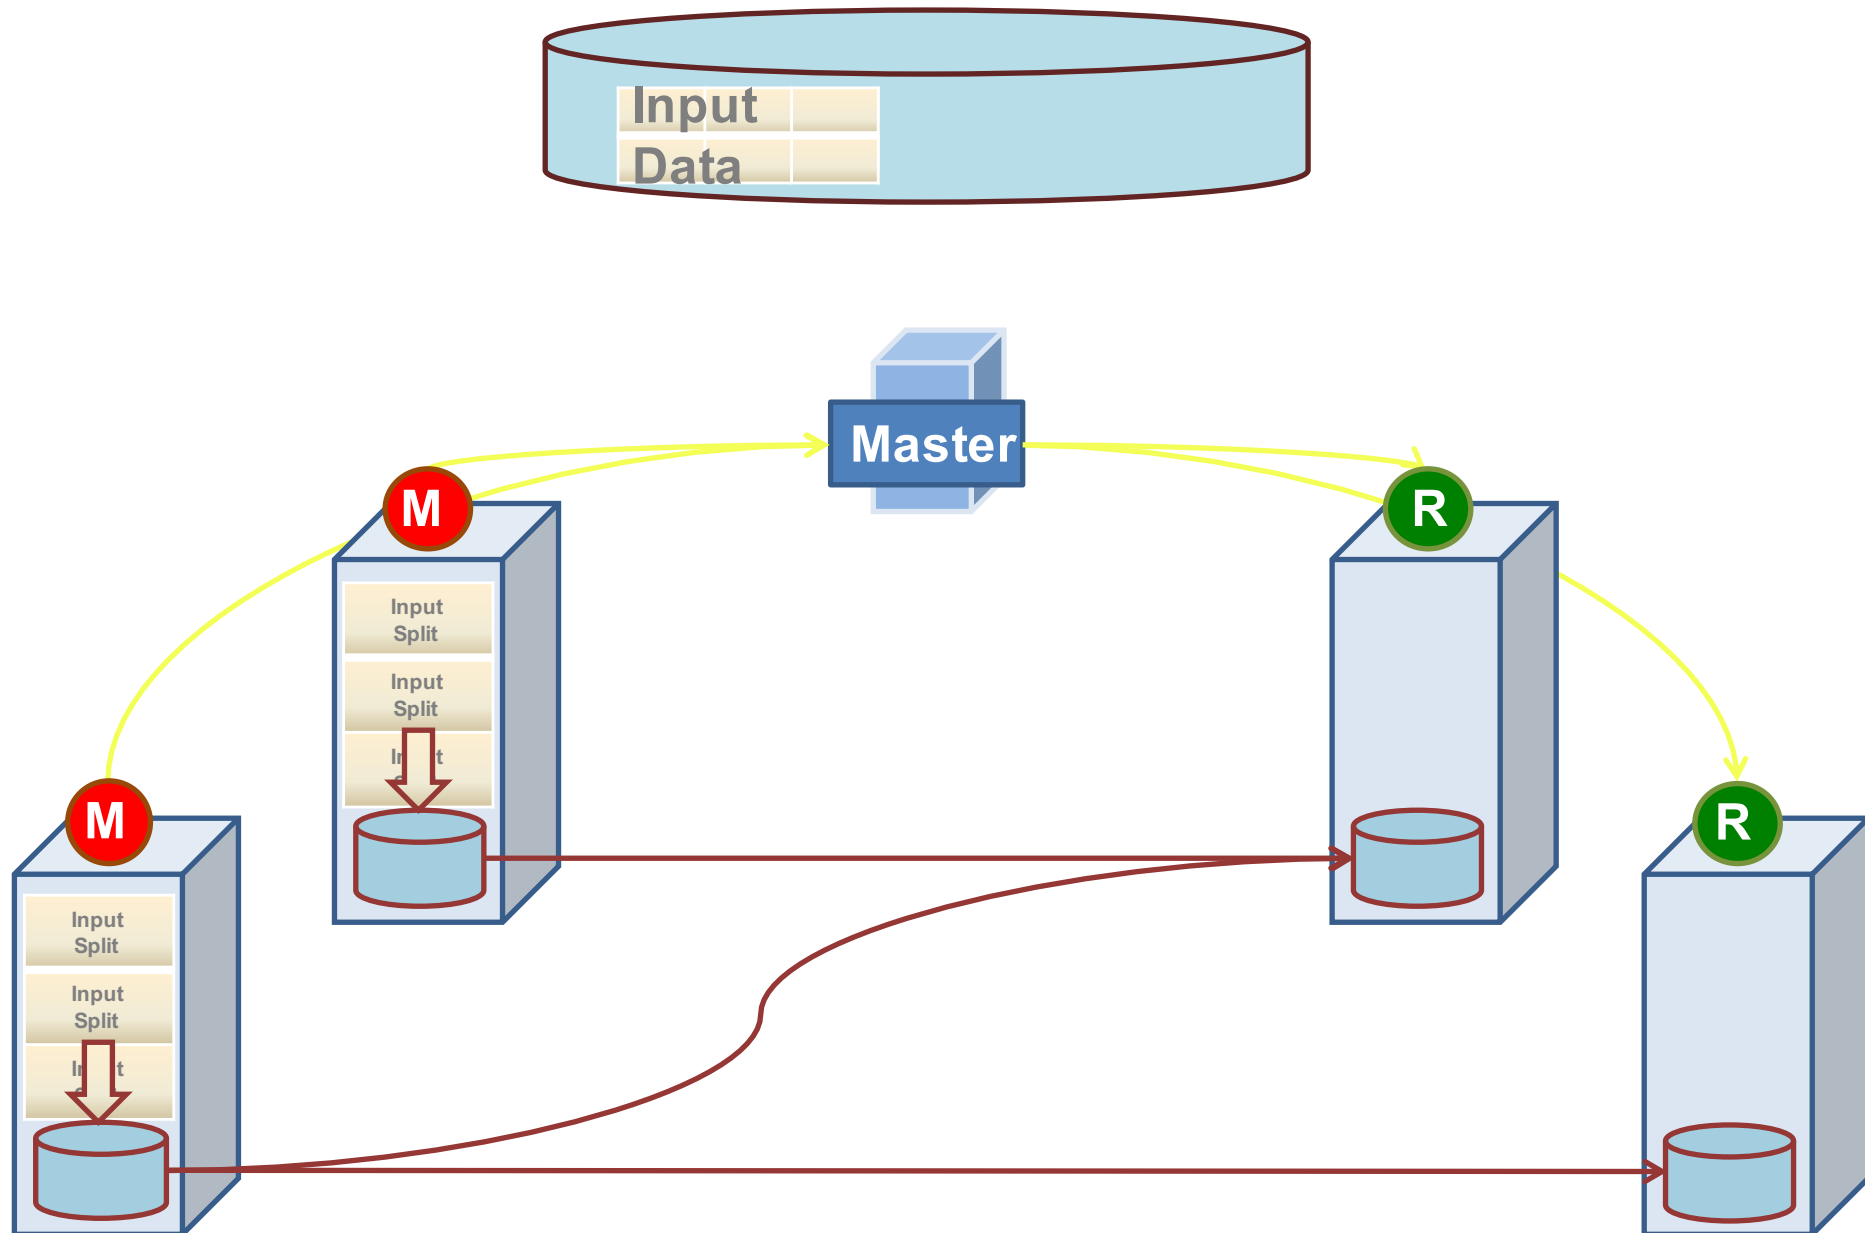
System view – Execution (Hadoop 1.0)



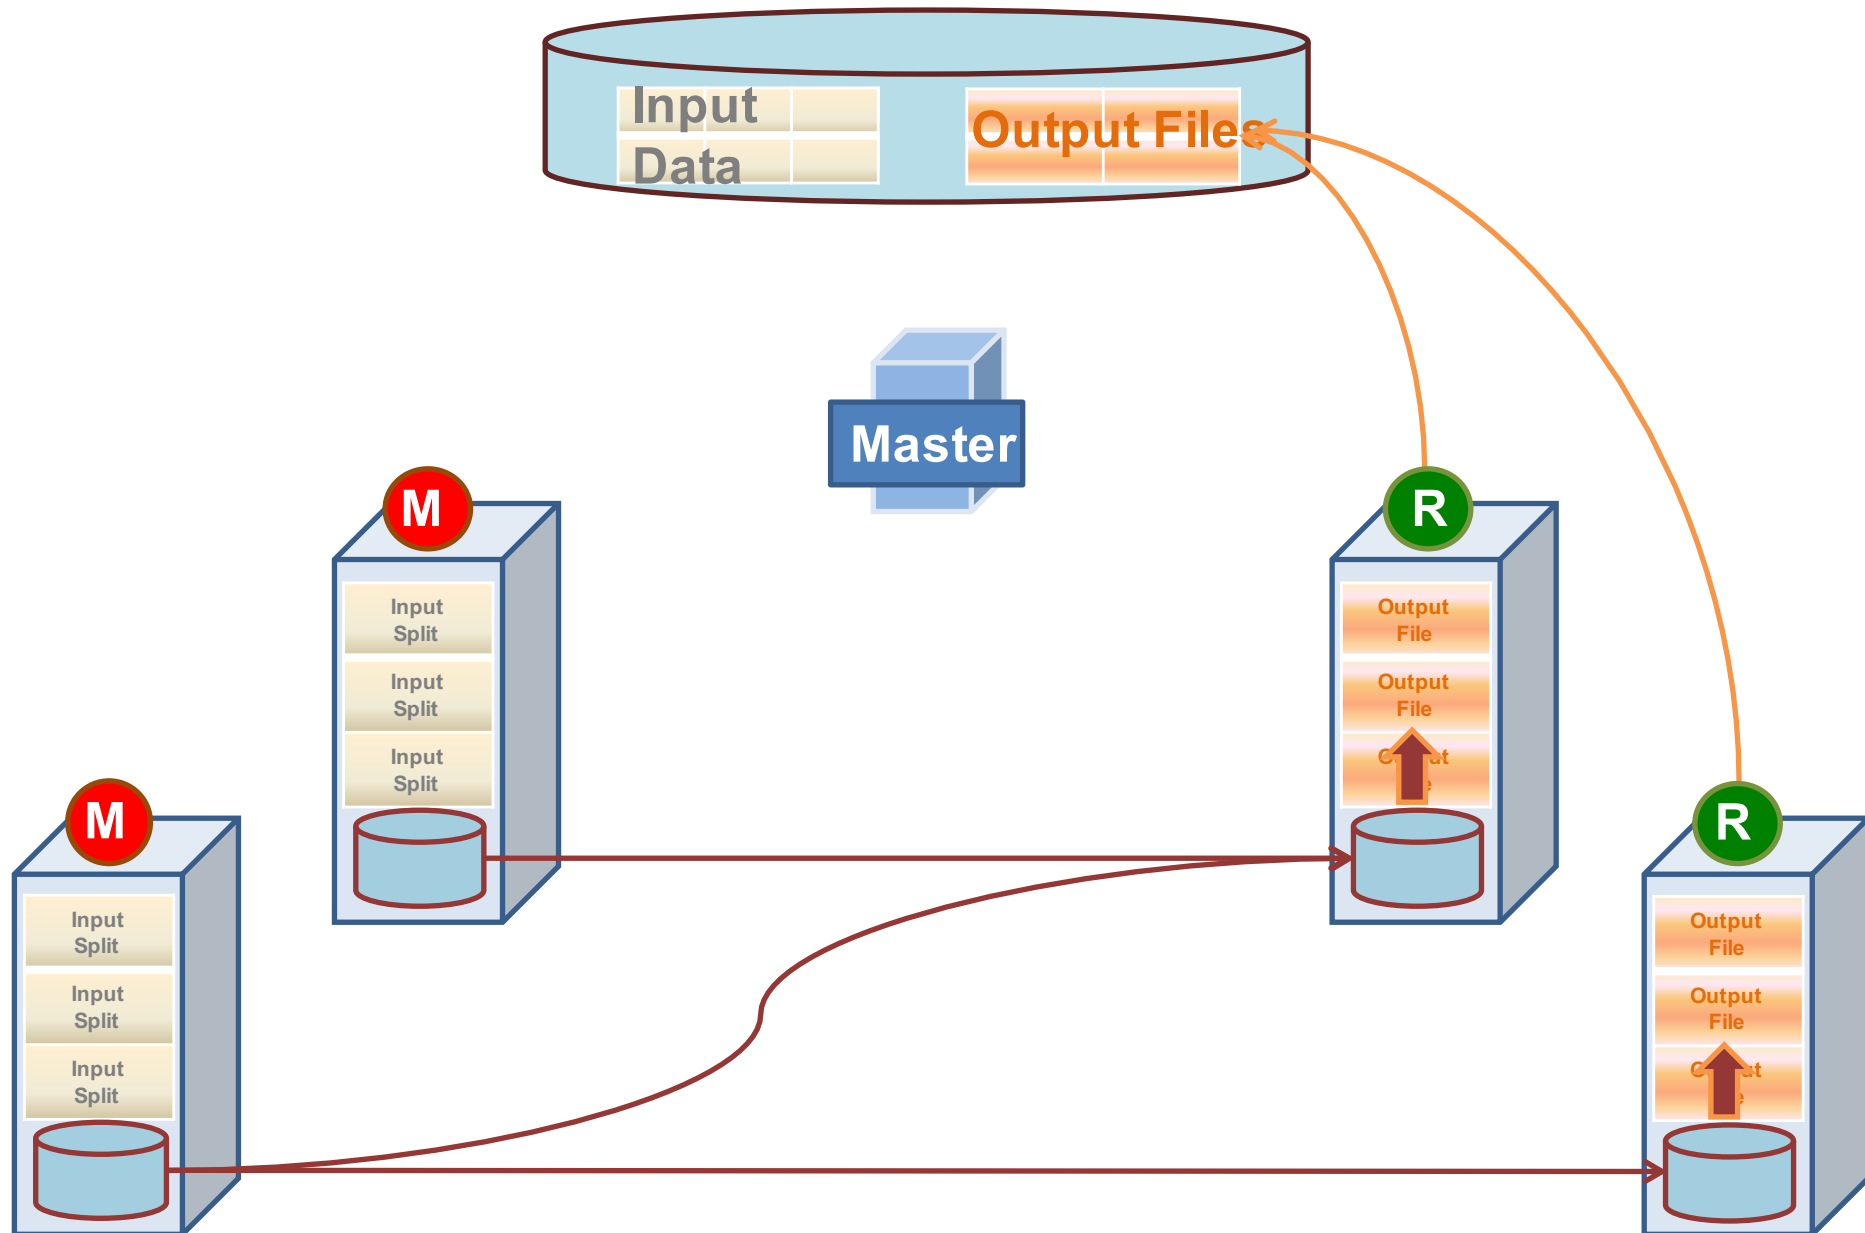
System view – Execution (Hadoop 1.0)



System view – Execution (Hadoop 1.0)



System view – Execution (Hadoop 1.0)



Coordination

- Master data structures
 - **Task status**: idle, in-progress, completed
 - **Idle tasks** get scheduled as workers become available
 - When a **map** task **completes**, it **sends** the **master** the **location** and sizes of its **R intermediate files**, one for each reducer
 - **Master pushes** this **info** to **reducers**
- Master **pings workers** periodically to detect failures



Failures




- Map worker failure
 - **Map** tasks **completed or in-progress** at worker are **reset** to idle
 - Reduce workers are notified when task is rescheduled on another worker
- Reduce worker failure
 - **Only in-progress tasks** are **reset** to idle
- Master failure
 - **MapReduce job** is **aborted** and client is notified

Content


- Big Data Analytics and Data Science
- Big Data Analytics Applications
- Map-reduce and Hadoop fundamentals
- **Map-reduce and Hadoop pros and cons**
- Hadoop Success Stories
- Hadoop Evolution and Hadoop eco-system
- Pig and Hive
- Map Reduce Cloud based solutions
- Where is the world going? Spark Apache project

Hadoop 1.0 pros




- “Simple” and easy to use
 - A programmer defines his job with only map and reduce functions, **without** having to **specify physical distribution** of his **job** across nodes
- Fault tolerance
 - MapReduce is **highly fault-tolerant** (continue to work in spite of an average of 1.2 failures per analysis job at Google)
- Flexible
 - **No dependency on data model** and schema (good for irregular or unstructured data)

Hadoop 1.0 cons



- **A single fixed dataflow**
 - The dataflow is fixed, many **complex algorithms** are **hard** to **implement** in a single job
 - Some algorithms that **require multiple inputs** are **not well supported** since the dataflow of MapReduce is originally designed to read a single input and generate a single output
- **No high-level language**
 - **No declarative language** like SQL in DBMS and any query optimization technique
- **No schema and no index**
 - Each item is parsed at reading input and transform it into data objects for data processing, causing performance degradation

Hadoop 1.0 cons



- Low efficiency
 - With fault-tolerance and scalability as its primary goals, MapReduce operations are **not** always **optimized for I/O efficiency**. In addition, Map and Reduce are **blocking** operations
 - **No** specific **execution** plans and **does not optimize** plans like DBMS does to minimize **data transfer** across nodes

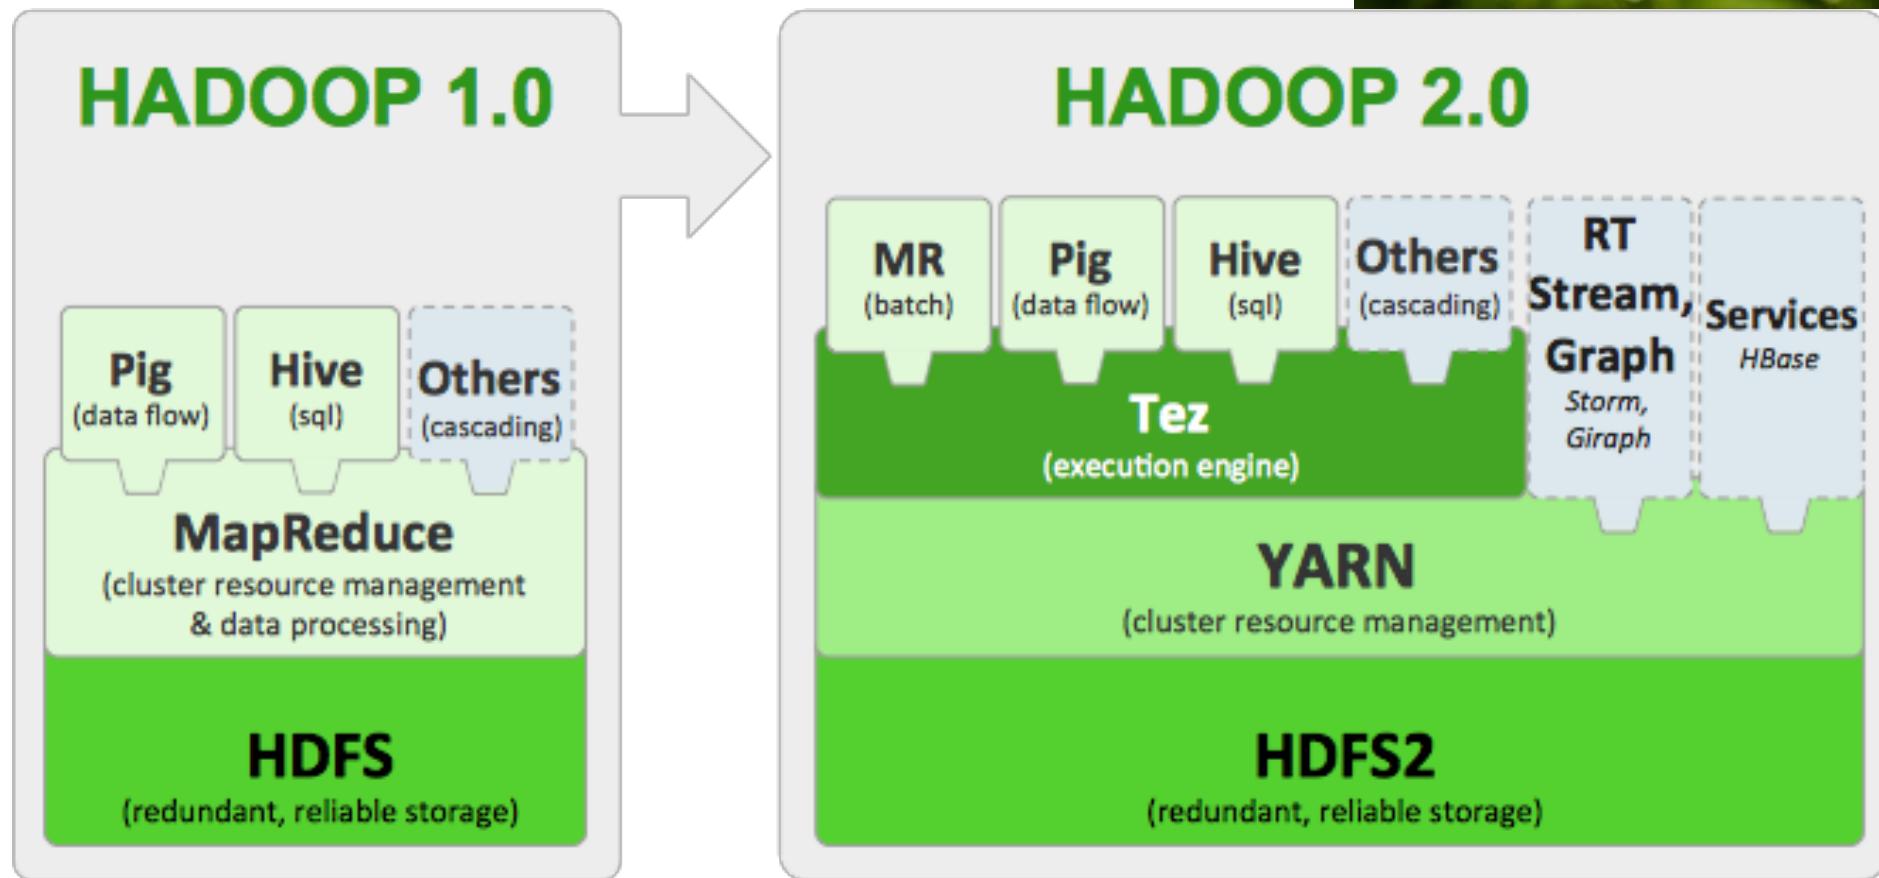
Comparing Hadoop with commercial data warehouses



- Hadoop **2~50 times slower**, except in data loading
- **in 2011** **was** **achieved**
• ~~Current~~ Hadoop system is scalable but ~~achieves~~ **very low efficiency per node**, around 5MB/s processing rate, repeating a mistake that previous studies on HPC did *focusing on scalability but missing efficiency*

Source: K. H. Lee, Y. J. Lee, H. Choi, Y. D. Chung, and B. Moon. Parallel data processing with MapReduce: a survey. *SIGMOD Rec.* 40(4), 2011.

Hadoop Eco-system



<https://www.youtube.com/watch?v=Z5kQR71yJpE>

<http://www.tableausoftware.com>