THE WHY

WE ARE STORING MORE DATA NOW THAN WE EVER HAVE BEFORE

THE WHY

WE ARE STORING MORE DATA NOW THAN WE EVER HAVE BEFORE

CONNECTIONS BETWEEN OUR DATA ARE GROWING ALL THE TIME

THE WHY

WE ARE STORING MORE DATA NOW THAN WE EVER HAVE BEFORE

CONNECTIONS BETWEEN OUR DATA ARE GROWING ALL THE TIME

WE DON'T MAKE THINGS KNOWING THE STRUCTURE FROM DAY 1

THE WHY

WE ARE STORING MORE DATA NOW THAN WE EVER HAVE BEFORE

CONNECTIONS BETWEEN OUR DATA ARE GROWING ALL THE TIME

WE DON'T MAKE THINGS KNOWING THE STRUCTURE FROM DAY 1

SERVER ARCHITECTURE IS NOW AT A STAGE WHERE WE CAN TAKE ADVANTAGE OF IT

NOSQL

USE CASES

LARGE DATA VOLUMES

MASSIVELY DISTRIBUTED ARCHITECTURE
REQUIRED TO STORE THE DATA
GOOGLE, AMAZON, FACEBOOK, 100K SERVERS

NOSQL

USE CASES

LARGE DATA VOLUMES

MASSIVELY DISTRIBUTED ARCHITECTURE
REQUIRED TO STORE THE DATA
GOOGLE, AMAZON, FACEBOOK, 100K SERVERS

EXTREME QUERY WORKLOAD

IMPOSSIBLE TO EFFICIENTLY DO JOINS AT THAT
SCALE WITH AN RDBMS

NOSQL

USE CASES

LARGE DATA VOLUMES

MASSIVELY DISTRIBUTED ARCHITECTURE
REQUIRED TO STORE THE DATA
GOOGLE, AMAZON, FACEBOOK, 100K SERVERS

EXTREME QUERY WORKLOAD

IMPOSSIBLE TO EFFICIENTLY DO JOINS AT THAT
SCALE WITH AN RDBMS

SCHEMA EVOLUTION

SCEMA FLEXIBILITY IS NOT TRIVIAL AT A LARGE
SCALE BUT IT CAN BE WITH NO SQL

NOSQL

PROS AND CONS

PROS

MASSIVE SCALABILITY
HIGH AVAILABILITY
LOWER COST
SCHEMA FLEXIBILITY
SPARCE AND SEMI STRUCTURED DATA

NOSQL

PROS AND CONS

PROS

MASSIVE SCALABILITY
HIGH AVAILABILITY
LOWER COST
SCHEMA FLEXIBILITY
SPARCE AND SEMI STRUCTURED DATA

CONS

LIMITED QUERY CAPABILITIES
NOT STANDARDISED (PORTABILITY MAY BE AN ISSUE)
STILL A DEVELOPING TECHNOLOGY

BIGTABLE

KEY VALUE

FOUR

EMERGING TRENDS IN NOSQL DATABASES

GRAPHDB

DOCUMENT

BUT FIRST...

IMAGINE A LIBRARY

LOTS OF DIFFERENT FLOORS

DIFFERENT SECTIONS ON EACH FLOOR

DIFFERENT BOOKSHELVES IN EACH SECTION

LOTS OF BOOKS ON EACH SHELF

LOTS OF PAGES IN EACH BOOK

LOTS OF WORDS ON EACH PAGE

EVERYTHING IS WELL ORGANISED AND EVERYTHING HAS A SPACE

BUT FIRST...

IMAGINE A LIBRARY

WHAT HAPPENS IF WE BUY TOO MANY BOOKS!?

(THE WORLD EXPLODES AND THE KITTENS WIN)

BUT FIRST...

IMAGINE A LIBRARY

WHAT HAPPENS IF WE WANT TO STORE CDS ALL OF A SUDDEN!?

(THE WORLD EXPLODES AND THE KITTENS WIN)

BUT FIRST...

IMAGINE A LIBRARY

WHAT HAPPENS IF WE WANT TO GET RID OF ALL BOOKS THAT MENTION KITTENS

(KITTENS STILL WIN)

KEY VALUE

AGAIN, DESIGNED TO WORK WITH A LOT OF DATA

EACH BIT OF DATA IS STORED IN A SINGLE COLLECTION

EACH COLLECTION CAN HAVE DIFFERENT TYPES OF DATA

KEY
VALUE

KEY VALUE

OUR VALUES ARE HIDDEN INSIDE THE KEYS

TO FIND OUT WHAT THEY ARE WE NEED TO QUERY THEM

What is in Key B?

The Triangle

KEY VALUE

(VOLDERMORT)

DOCUMENT STORE

DESIGNED TO WORK WITH A LOT OF DATA (BEGINNING TO NOTICE A THEME?)

VERY SIMILAR TO A KEY VALUE DATABASE

MAIN DIFFERENCE IS THAT YOU CAN ACTUALLY SEE THE VALUES

DOCUMENT STORE

DOCUMENT STORE

SIDE

REMEMBER HOW SQL DATABASES ARE LIBRARIES?

NOTE

NO SQL IS MORE LIKE A BAG OF CATS!

SIDE

colour: tabby
name: Gunther

colour: ginger
name: Mylo

WE CAN ADD IN FIELDS AS AND WHEN WE NEED THEM

NOTE

colour: grey
name: Ruffus
age: kitten

colour: ginger(ish)
name: Fred
age: kitten

colour: ginger(ish)
name: Quentin
legs: 3

DOCUMENT STORE

DOCUMENT STORE

BIG TABLE

BEHAVES LIKE A STANDARD RELATIONAL DATABASE BUT WITH A SLIGHT CHANGE

DESIGNED TO WORK WITH A LOT OF DATA...A REALLY BIG CRAP TON

CREATED BY GOOGLE AND NOW USED BY LOTS OF OTHERS

http://research.google.com/archive/spanner.html

http://research.google.com/archive/bigtable.html

BIG TABLE

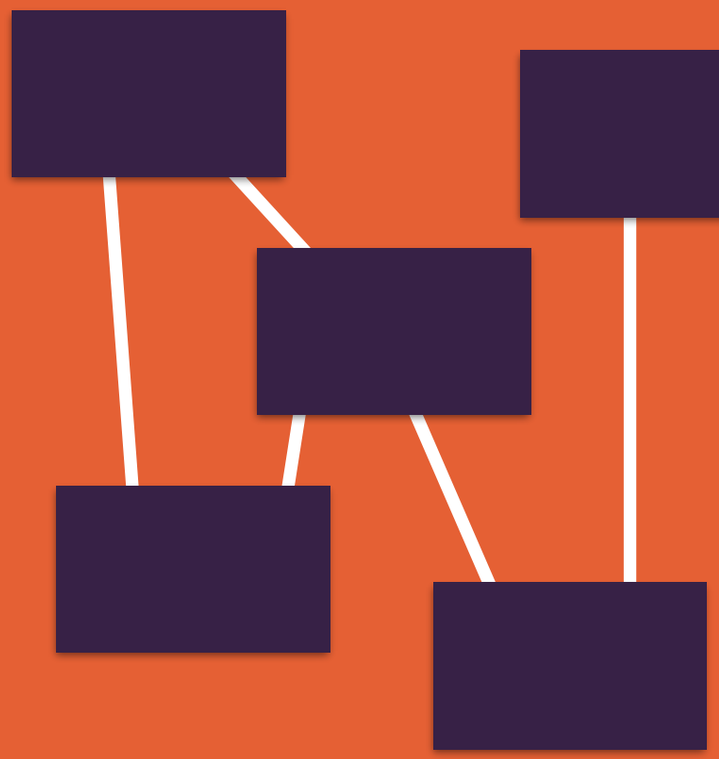

THIS IS A STANDARD RELATIONAL DATABASE

THIS IS A BIG TABLE DATABASE

(AND NOW THE NAME MAKES SENCE!)

BIG TABLE

"A Bigtable is a sparse, distributed, persistent multidimensional sorted map. The map is indexed by a row key, column key, and a timestamp; each value in the map is an uninterpreted array of bytes."

BIG TABLE

"A Bigtable is a sparse, distributed, persistent multidimensional sorted map. The map is indexed by a row key, column key, and a timestamp; each value in the map is an uninterpreted array of bytes."

BIG TABLE

"A Bigtable is a sparse, distributed, persistent multidimensional sorted map. The map is indexed by a row key, column key, and a timestamp; each value in the map is an uninterpreted array of bytes."

GRAPH DATABASE

FOCUS HERE IS ON MODELLING THE STRUCTURE OF THE DATA

INSPIRED BY GRAPH THEORY (GO MATHS!)

SCALES REALLY WELL TO THE STRUCTURE OF THE DATA

GRAPH
DATABASE

GRAPH
DATABASE

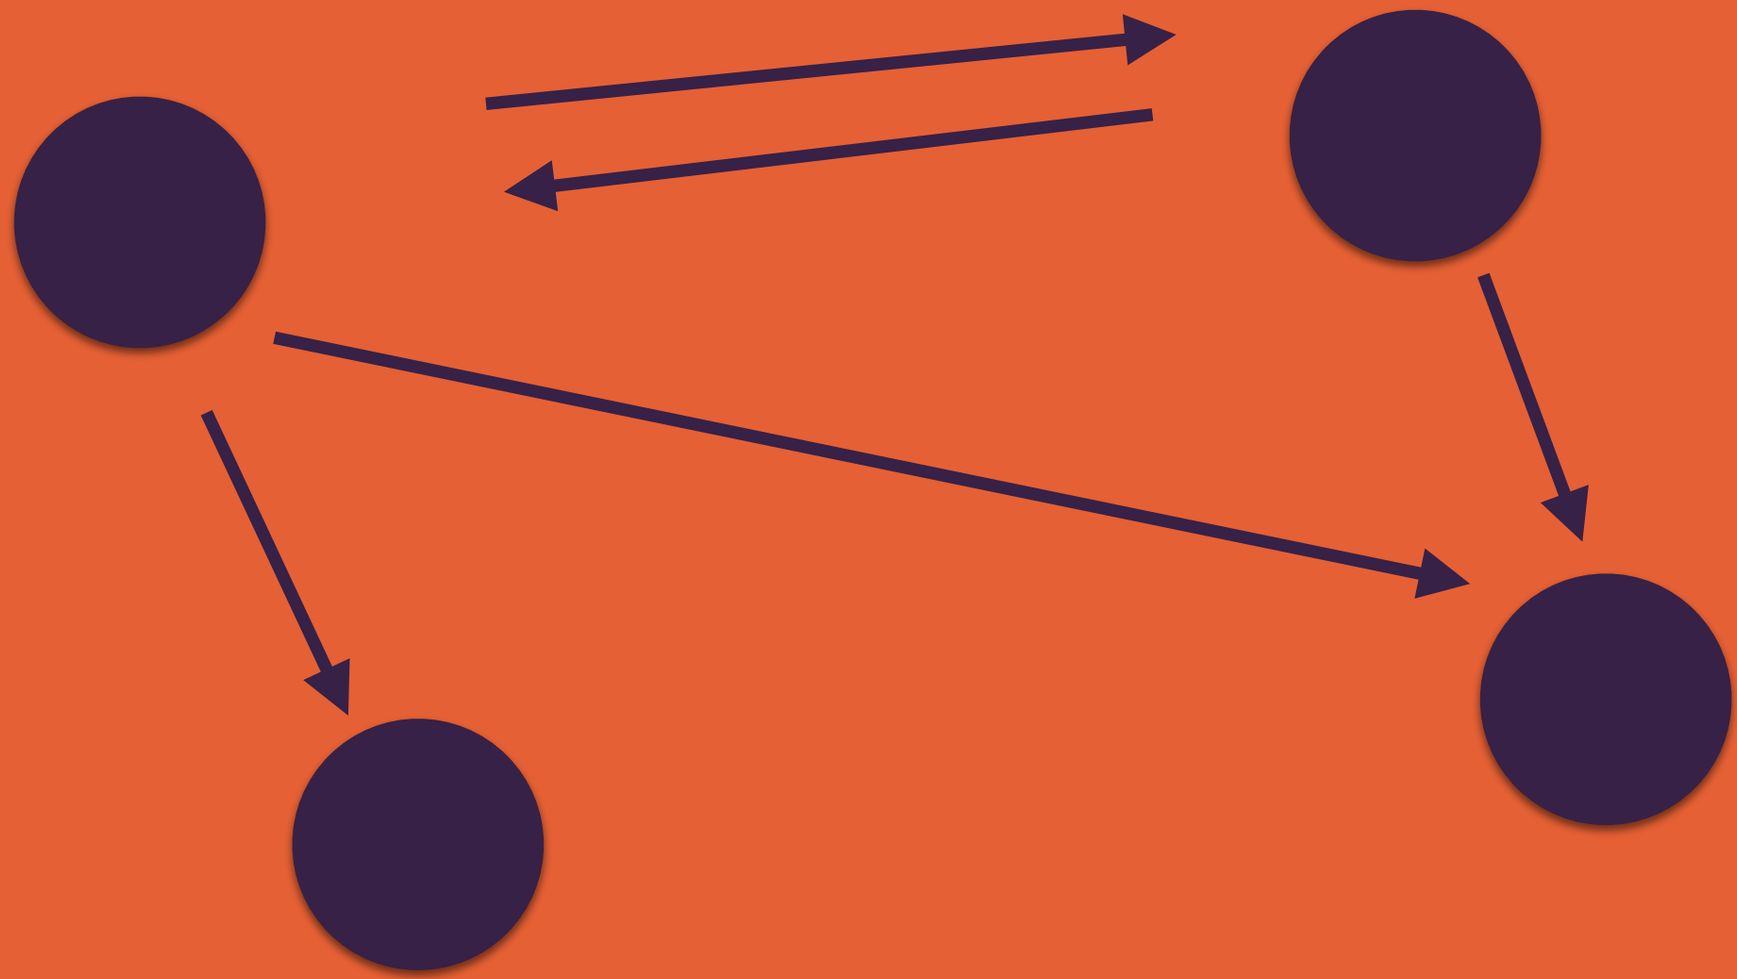

GRAPH
DATABASE

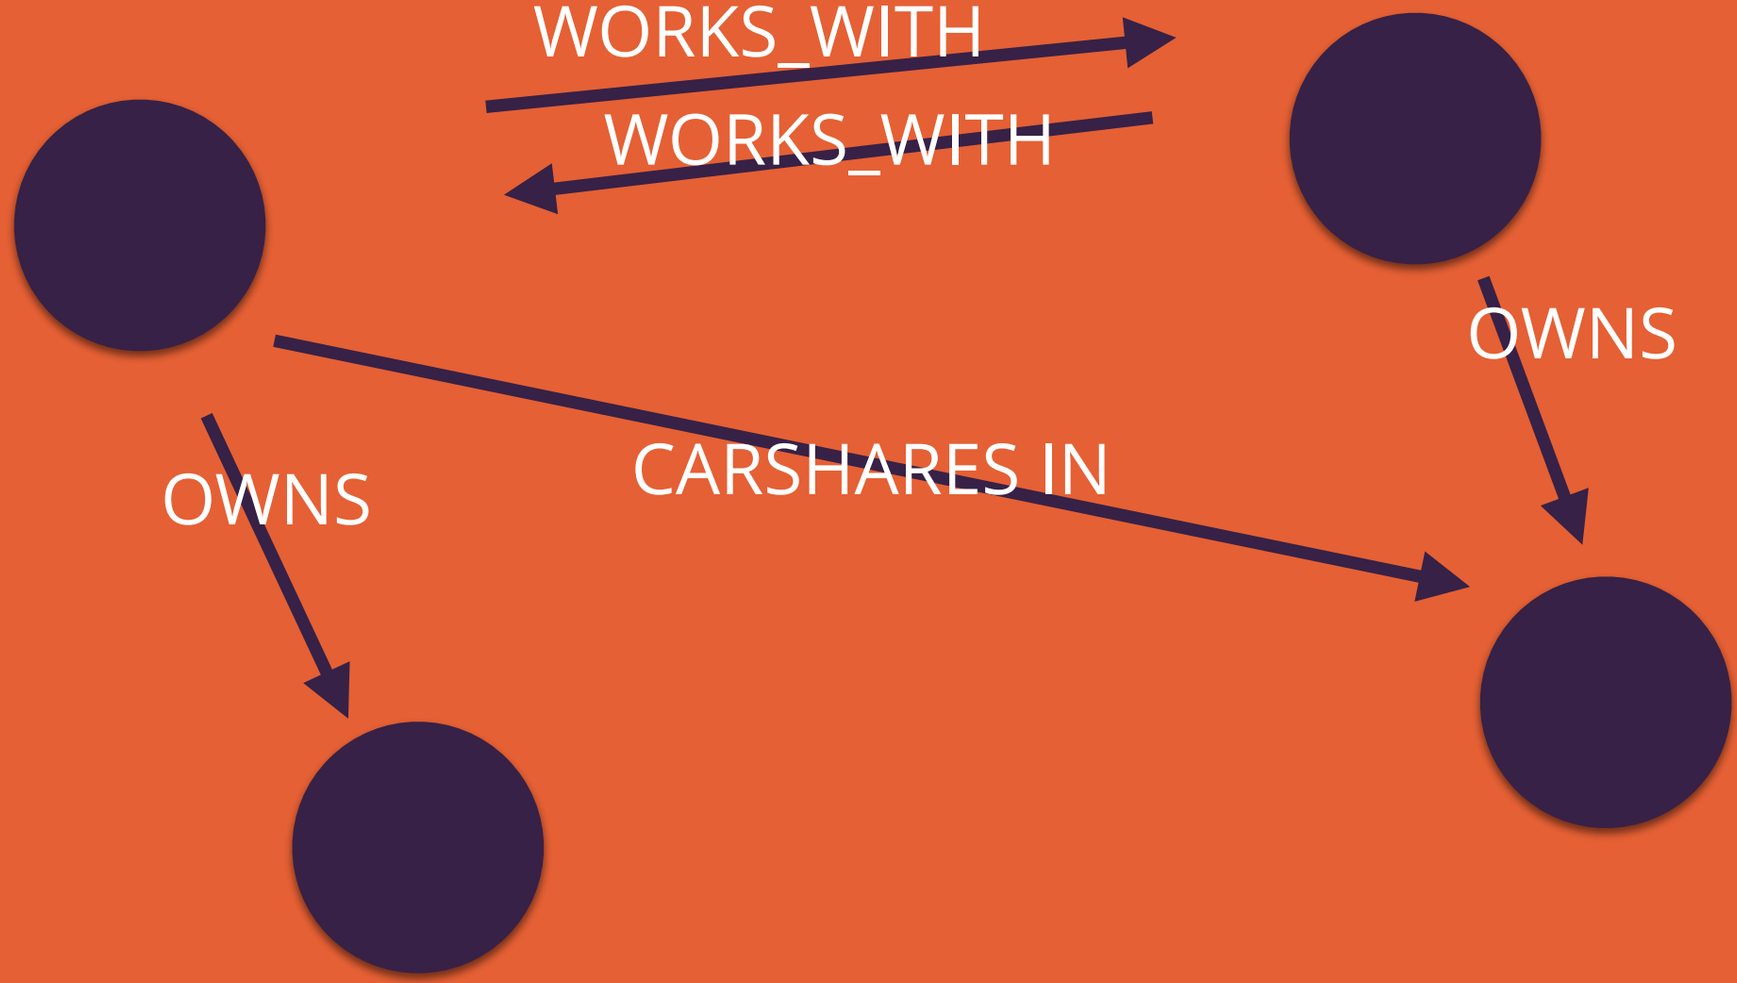

GRAPH DATABASE

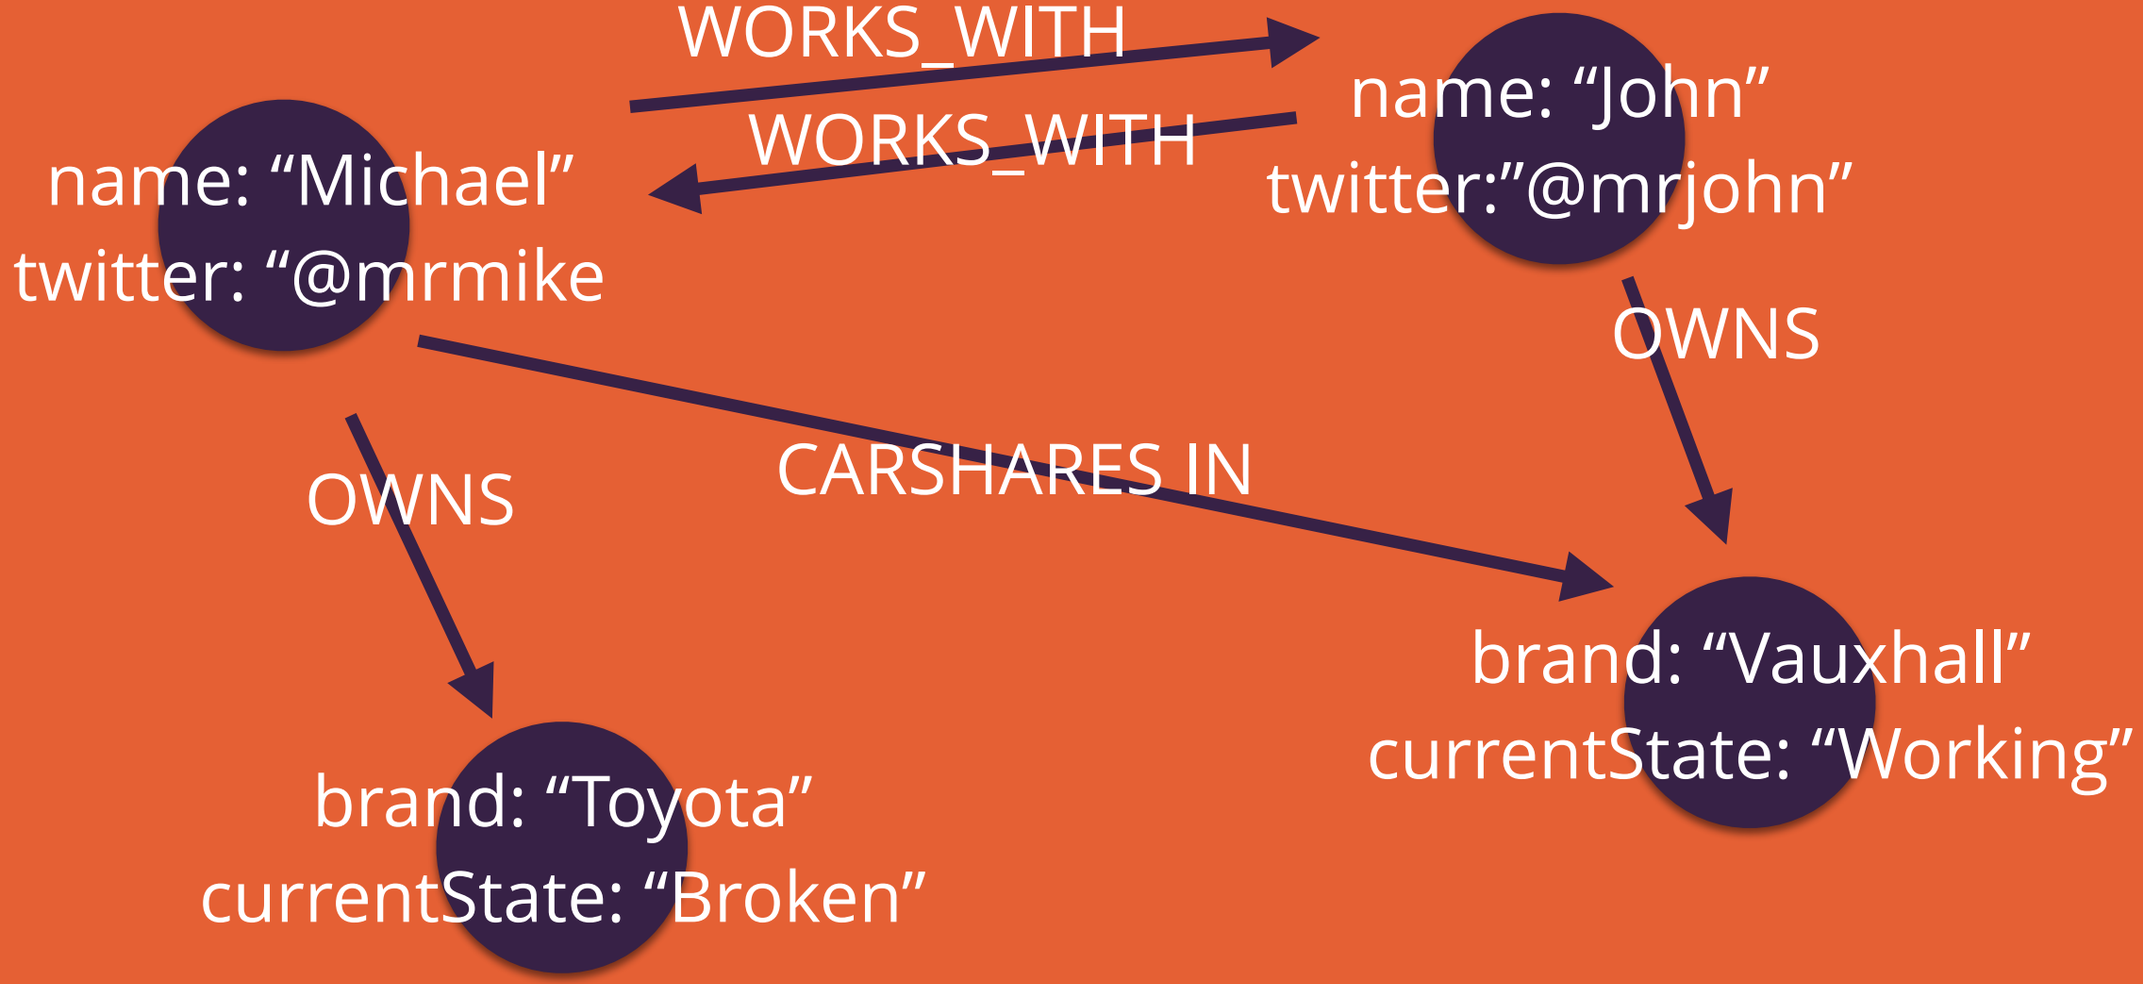

GRAPH DATABASE

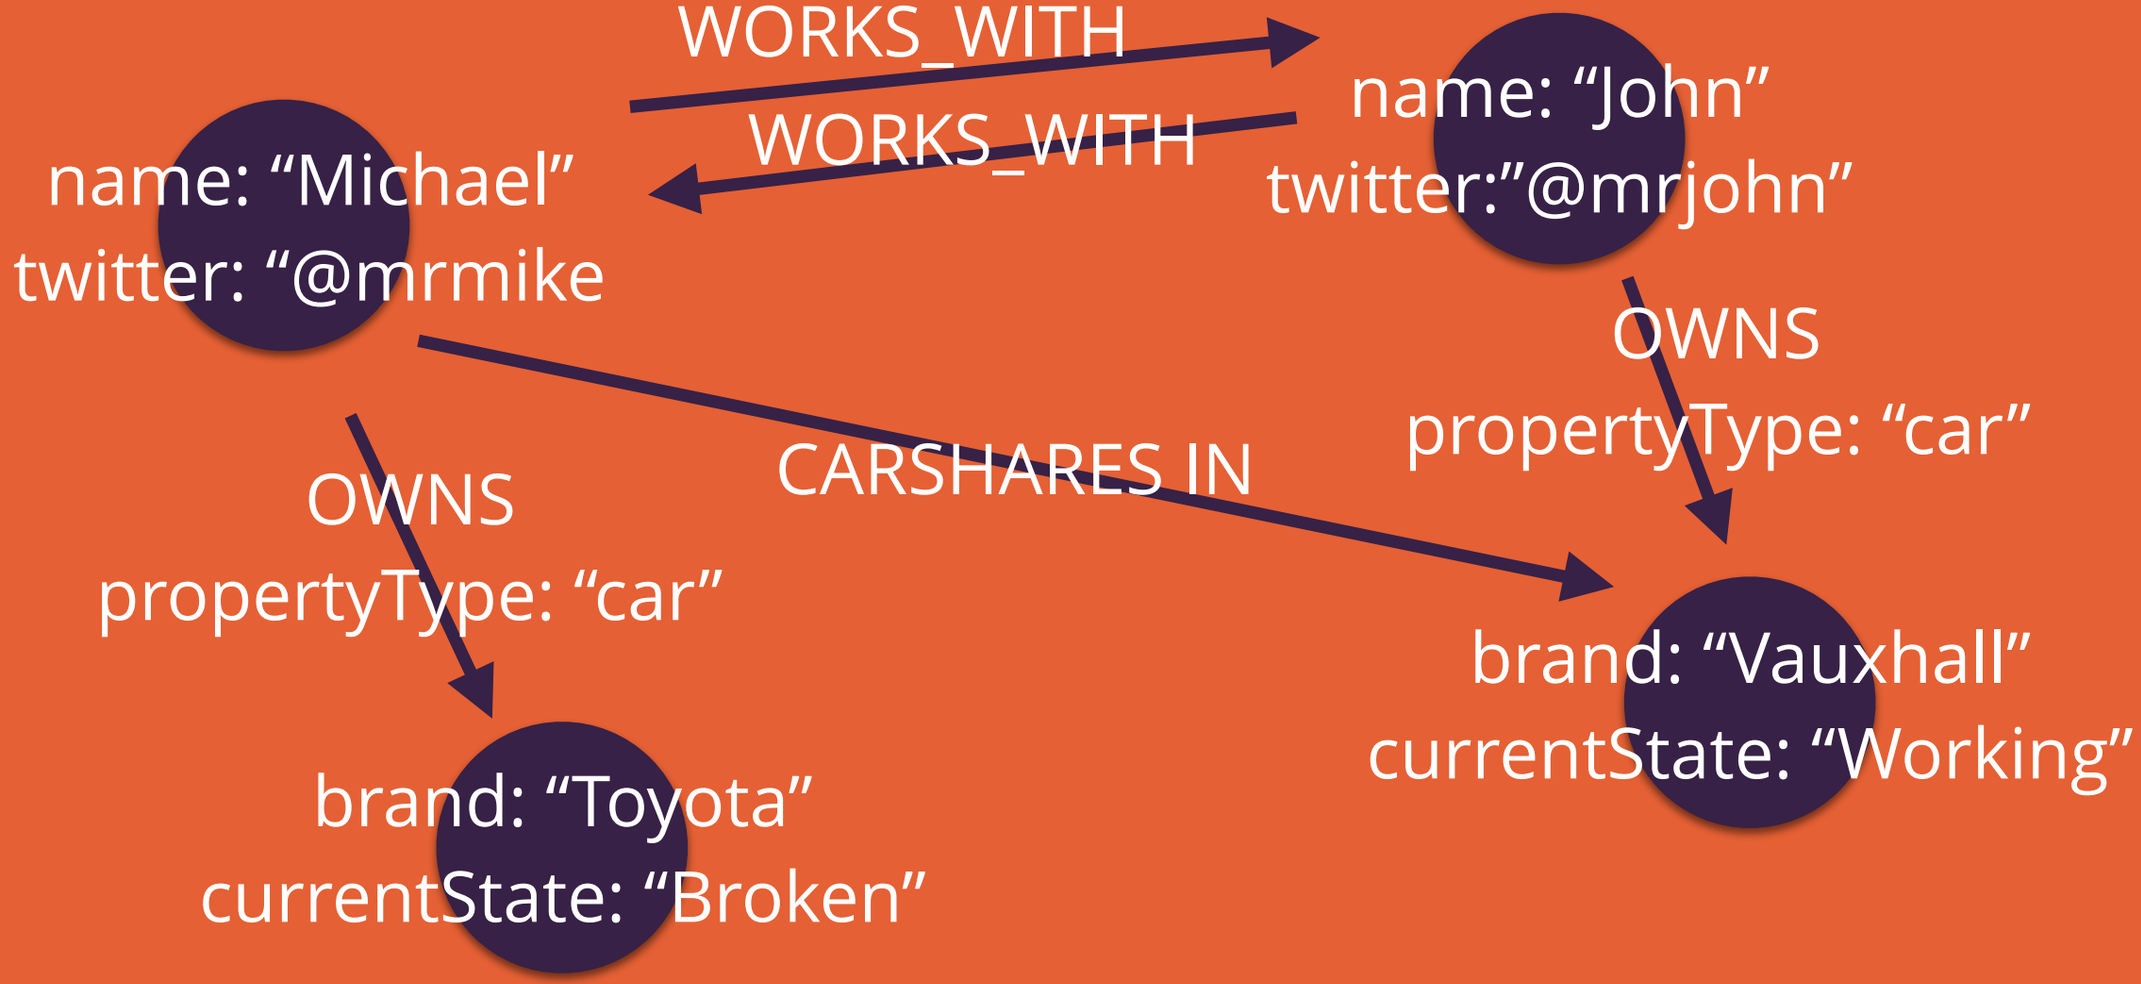

GRAPH
DATABASE

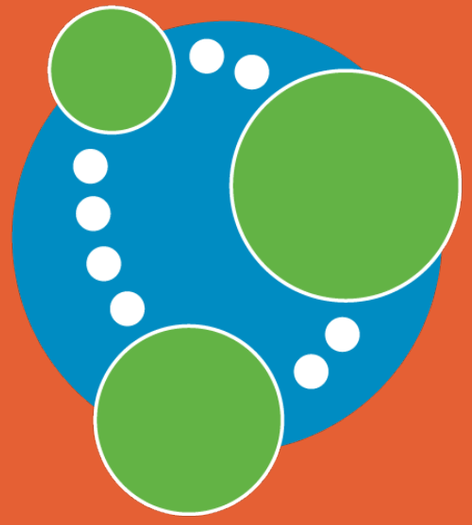

neo4j

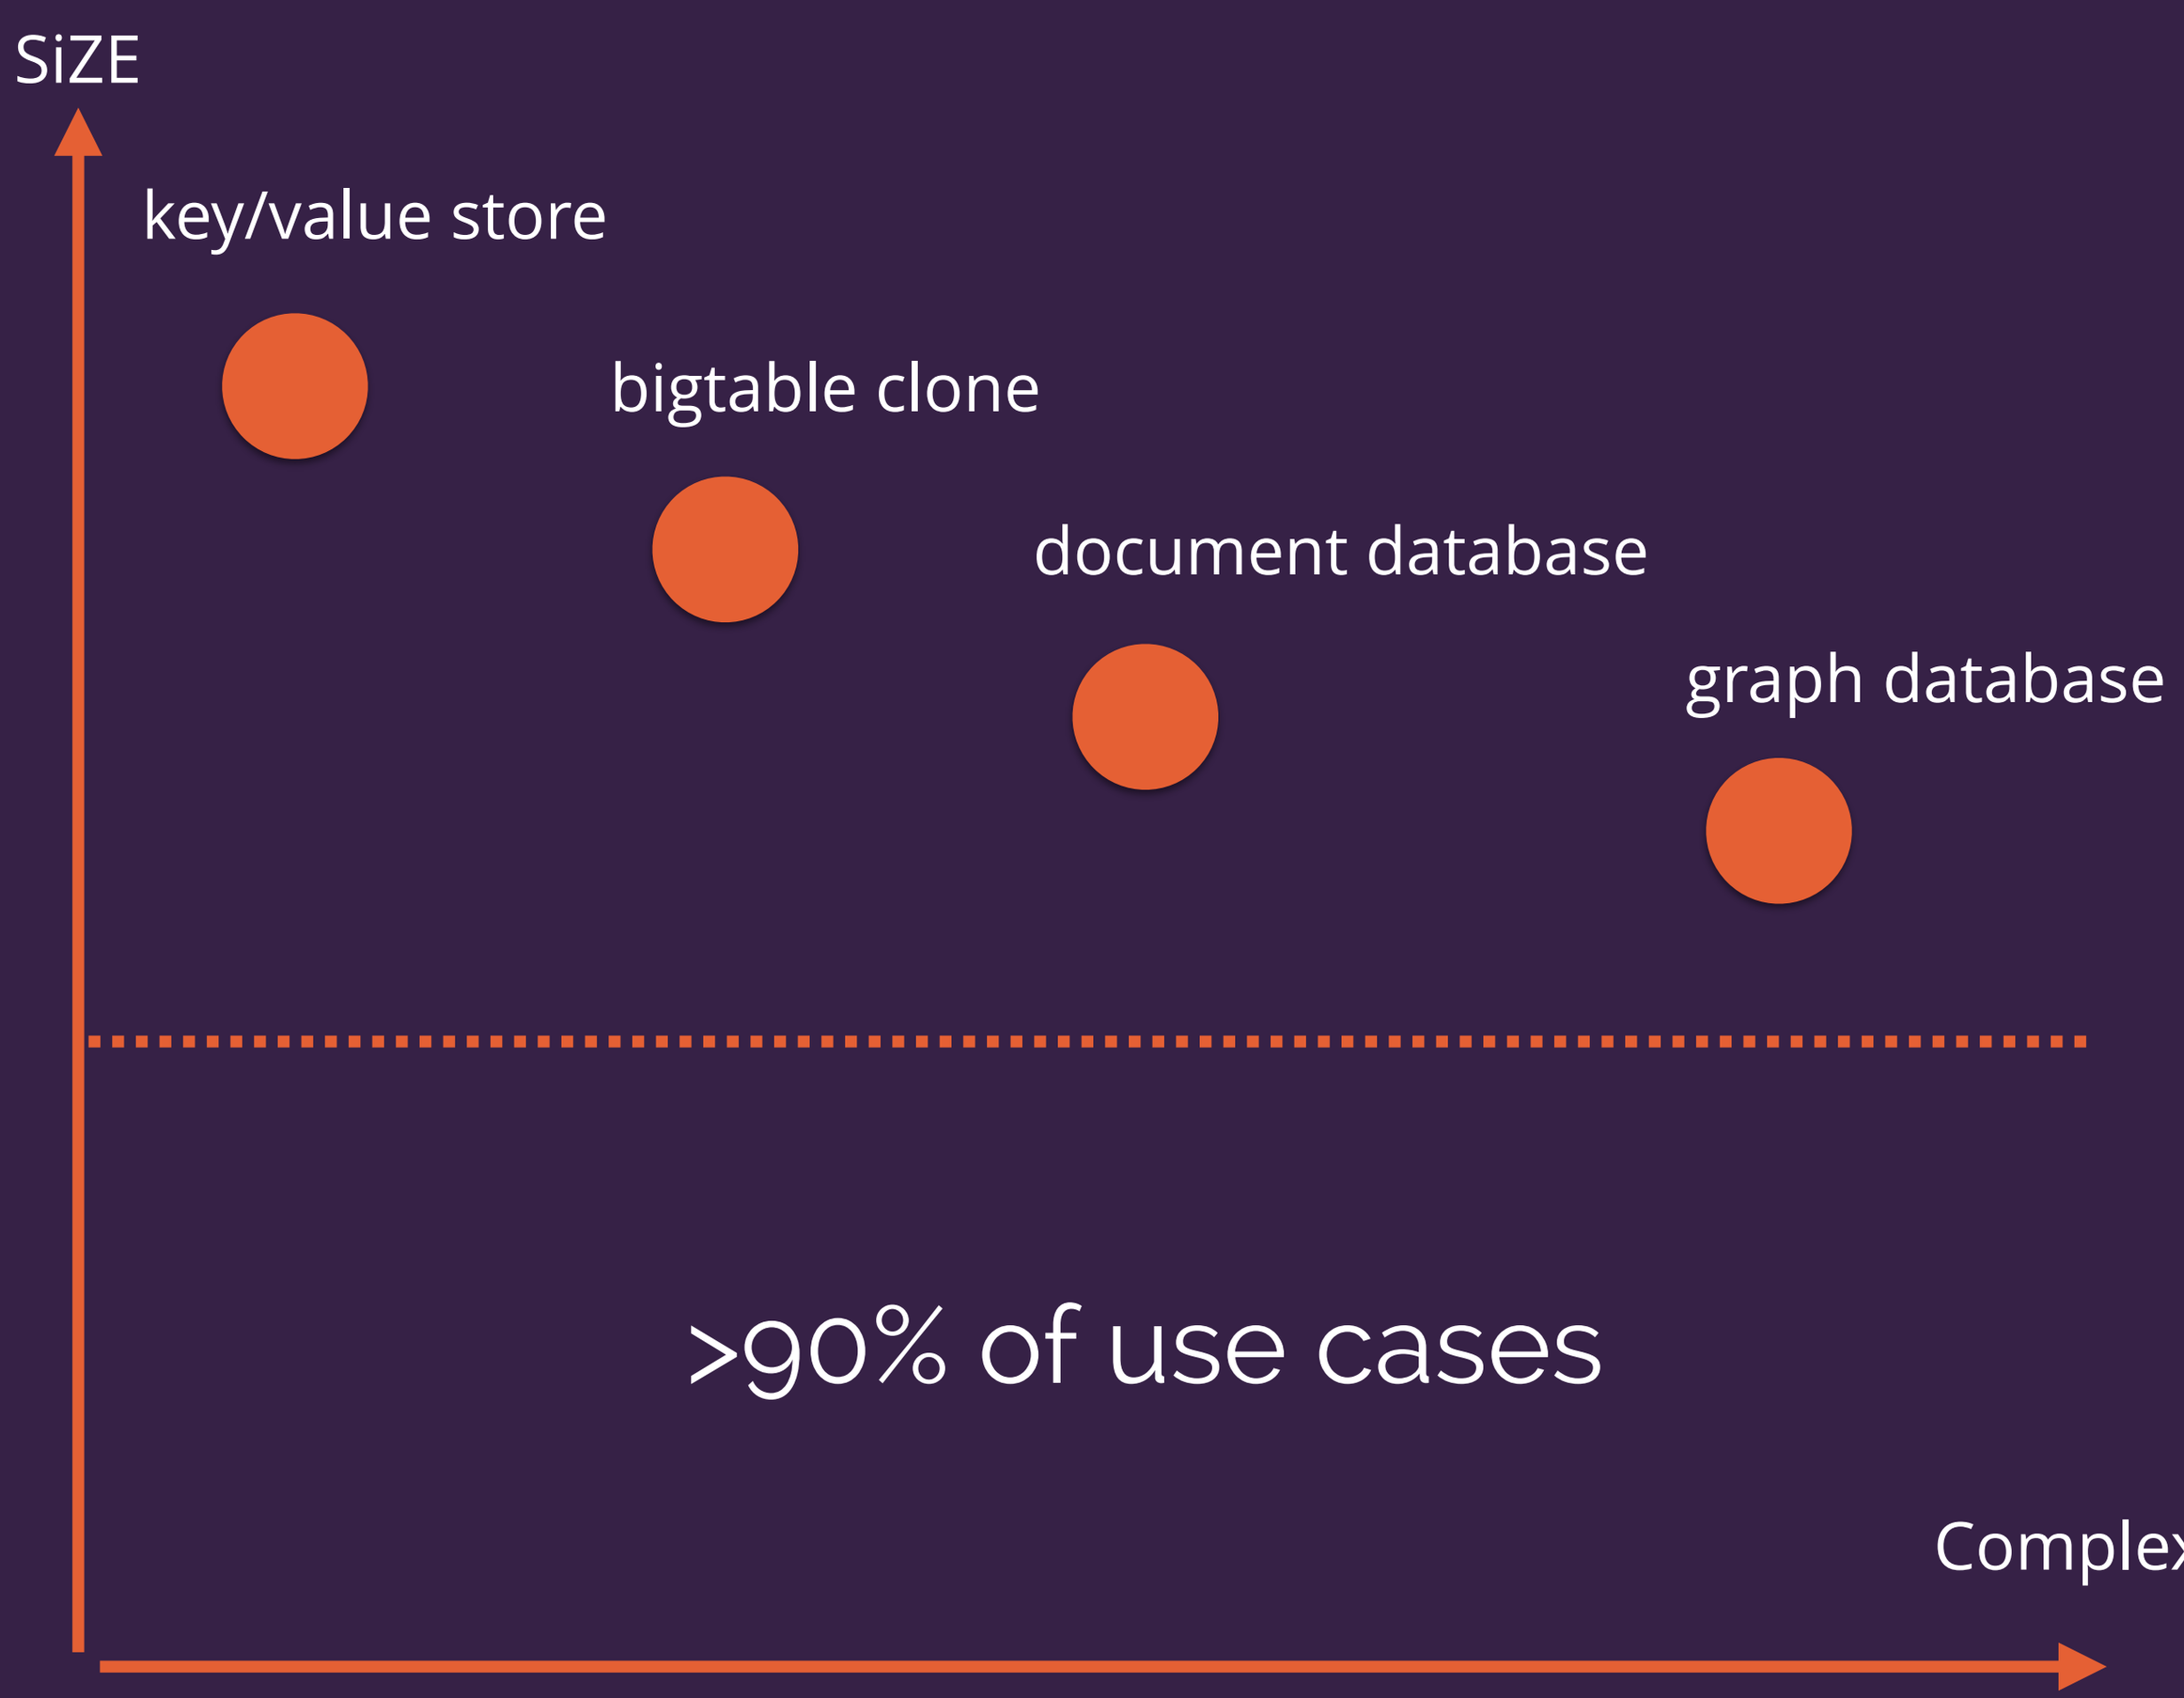
SiZE
key/value store
bigtable clone
document database
graph database
>90% of use cases
Complexity

WHEN TO USE

NOSQL

AND WHEN TO USE

SQL

THE BASICS

High availability and disaster recovery are a must

Understand the pros and cons of each design model

Don’t pick something just because it is new

Do you remember the zune?

Don’t pick something based JUST on performance

SQL

THE GOOD

High performance for transactions. Think ACID

Highly structured, very portable

Small amounts of data

SMALL IS LESS THAN 500GB

Supports many tables with different types of data

Can fetch ordered data

Compatible with lots of tools

SQL

Atomicity

Consistency

Isolation

Durability

SQL

THE BAD

Complex queries take a long time

The relational model takes a long time to learn

Not really scalable

Not suited for rapid development

noSQL

THE GOOD

Fits well for volatile data

High read and write throughput

In general it’s faster than SQL

Scales really well

Rapid development is possible

noSQL

Basically Available

Soft state Eventually consistent

noSQL

THE GOOD

Key/Value pairs need to be packed/unpacked all the time

Still working on getting security for these working as well as SQL

Lack of relations from one key to another

tl;dr

SQL

works great, can’t scale for large data

noSQL

works great, doesn't fit all situations

so use both, but think about when you want to use them!