



# Finding Similar Sets



Vassilis Christophides

christop@csd.uoc.gr

<http://www.csd.uoc.gr/~hy562>

University of Crete, Fall 2019



## Motivation

- Many Web-mining problems can be expressed as finding “similar” sets:
  - ◆ Pages with similar words, e.g., for classification by topic
  - ◆ NetFlix users with similar tastes in movies for recommendation systems
    - Dual: movies with similar sets of fans
  - ◆ Images of related things
- The best techniques depend on whether you are looking for items that are very similar or only somewhat similar
  - ◆ Special cases are easy, e.g., identical documents, or one document contained character-by-character in another
  - ◆ General case, where many small pieces of one document appear out of order in another, is very hard



## Comparing Documents for Near Duplicates

- **Applications:** Given a body of documents, find pairs of documents with a lot of text in common, e.g.:
    - ◆ **Mirror Web sites**, or approximate mirrors
      - Application: Don't want to show both in a search
    - ◆ **Plagiarism**, including large quotations
    - ◆ **Similar news articles** at many news sites
      - Application: Cluster articles by "same story"
  - Simple IR approaches are not suited:
    - ◆ Document = set of words appearing in document
    - ◆ Document = set of "important" words
- Why? **we need to account for ordering of words!**



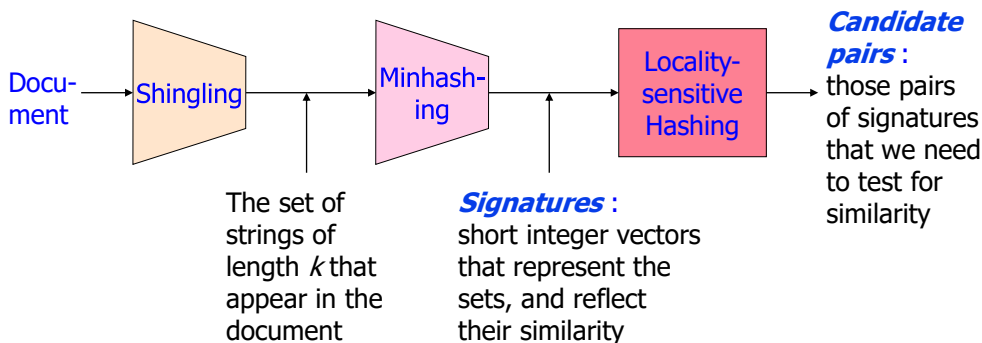
## Main Issues

- What is the **right representation** of the document when we check for similarity?
  - ◆ E.g., representing a document as a set of characters will not do (why?)
- When we have billions of documents, keeping the full text in memory is not an option
  - ◆ We need to find a **shorter representation**
- How do we do **pairwise comparisons** of billions of documents?
  - ◆ If exact match was the issue it would be ok, can we replicate this idea?



## Three Essential Techniques for Detecting Similar Documents

Fall 2019



- **Shingling**: convert documents, emails, etc., to *sets*
- **Minhashing**: convert *large sets* to *short signatures*, while preserving similarity
- **Locality-sensitive hashing**: focus on *pairs of signatures likely to be similar*

Unra

5



## Shingles

Fall 2019

- A  **$k$ -shingle** (or  **$k$ -gram**) for a document is a sequence of  $k$  characters that appears in the document
  - ◆ Represent a document by its set of  **$k$ -shingles**
- **Example**:  $k=2$ ; doc= abcab. Set of 2-shingles = {ab, bc, ca}
  - ◆ Option: regard shingles as a bag (multiset), and count ab twice
- **Working Assumption**: Documents that have lots of shingles in common have similar text, even if the text appears in different order
  - ◆ What if two documents differ by a word?
    - Affects only  **$k$ -shingles** within distance  $k$  from the word
  - ◆ What if we reorder paragraphs?
    - Affects only the  **$2k$  shingles** that cross paragraph boundaries

Unra

6



## Shingle Size

Fall 2019

- Is  $k=2$  a good choice for size?
- Example:  $k=2$ ;
  - ♦  $\text{doc1} = \text{abcab}$ . 2-shingles =  $\{\text{ab}, \text{bc}, \text{ca}\}$
  - ♦  $\text{doc2} = \text{cab c}$ . 2-shingles =  $\{\text{ab}, \text{bc}, \text{ca}\}$
- Careful decision: you must pick  $k$  large enough, or most documents will have most shingles
  - ♦  $k = 5$  is OK for short documents
  - ♦  $k = 10$  is better for long documents

Unia

7



## Shingles: Compression Option

Fall 2019

- How about space overhead?
    - ♦ Each character can be represented as a byte
    - ♦  $k$ -shingle requires  $k$  bytes
  - If  $k=9$ , to compare shingles we need to compare 9 bytes
  - To improve efficiency, we can compress long shingles:
    - ♦ hash them to (say) 4 bytes, and
    - ♦ represent a document by the set of hash values of its  $k$ -shingles
- $(\text{aaabbbccc})(\text{abcabcabc}) \rightarrow h(\text{aaabbbccc})h(\text{abcabcabc})$
- 18 bytes                       $\rightarrow$                       8 bytes
- Working Assumption: Two documents could (rarely) appear to have shingles in common, when in fact only the hash-values were shared

Unia

8



## Thought Question

- Why is it better to hash 9-shingles (say) to 4 bytes than to use 4-shingles?
  - ♦ There are many more possible shingles, this **reduces the likelihood that documents that share many shingles are not similar**
- **Hint:** How random are the 32-bit sequences that result from 4-shingling?
  - ♦ Assuming 20 characters are common in English, there are  $(20)^4 = 160000$  4-shingles  $< 2^{32}$
  - ♦ Using 9-shingles there are  $(20)^9 >> 2^{32}$



## MinHashing



## Basic Data Model: Sets


- Many similarity problems can be couched as finding subsets of some universal set that have significant intersection
- Examples include:
  - ♦ Documents represented by their sets of shingles (or hashes of those shingles):  $C_i = S(D_i)$
  - ♦ Similar customers or products
- Equivalently, each document is a 0/1 vector in the space of k-shingles
  - ♦ Each unique shingle is a dimension
  - ♦ Vectors are very sparse
- Interpret set intersection as bitwise AND, and set union as bitwise OR



## Jaccard Similarity of Sets

- The *Jaccard similarity* of two sets is the size of their intersection divided by the size of their union
  - ♦  $\text{sim}(C_1, C_2) = |C_1 \cap C_2| / |C_1 \cup C_2|$

- Example:



**3 in intersection**  
**8 in union**  
**Jaccard similarity**  
**= 3/8**





## Motivation for Min-Hash

Fall 2019

- Suppose we need to find near-duplicate documents among  $N=1$  million ( $10^6$ ) documents
- Naïvely, we would have to compute pairwise Jaccard similarities for every pair of docs
  - ◆  $N(N-1)/2 \approx 5 \cdot 10^{11}$  comparisons
  - ◆ At  $\approx 10^5$  secs/day and  $10^6$  comparisons/sec, it would take 5 days
- For  $N=10$  million ( $10^7$ ), it takes more than a year...

Inria

13



## From Sets to Boolean Matrices

Fall 2019

- Rows = elements (shingles) of the universal set
- Columns = sets (documents)
  - ◆ 1 in row  $e$  and column  $S$  if and only if  $e$  is a member of  $S$
  - ◆ Column similarity is the Jaccard similarity of the sets of their rows with 1
- Typical matrix is sparse (most rows are of type d, see later)
  - ◆ Sparse matrices are usually better represented by the list of places where there is a non-zero value
  - ◆ But the boolean matrix picture is conceptually useful

Shingles	Documents			
	1	1	1	0
	1	1	0	1
	0	1	0	1
	0	0	0	1
	1	0	0	1
	1	1	1	0
	1	0	1	0

Inria

14



## Example: Jaccard Similarity of Columns

	$C_1$	$C_2$		
a	0	1	*	
b	1	0	*	
c	1	1	*	*
d	0	0		
e	1	1	*	*
f	0	1	*	

$\text{Sim}(C_1, C_2) = \mathbf{2/5 = 0.4}$   
 $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = \mathbf{0.6}$



## Outline: Finding Similar Columns

- Naïve approach:
  - 1 **Compute** signatures of columns = small summaries of columns
  - 2 **Examine** pairs of signatures to find similar columns
    - **Requirement**: similarities of signatures and columns are related
  - 3 **Optional**: check that columns with similar signatures are really similar
- This scheme works but ...
  - ♦ What if the set of signatures (or k-shingles) is too large to fit in the memory?
  - ♦ Or the number of documents are too large?
- **Idea**: Find a way to **hash a document** (column) to a **single** (small size) **value**! and **similar documents** to the **same value**!
  - ♦ **Warning**: These methods can produce *false negatives*, and even *false positives* (if the above optional check is not made)





# Signatures

Fall 2019

- Key idea: “hash”  $h(\cdot)$  each column  $C$  to a small **signature**, such that:
  - ①  $h(C)$  is small enough that we can fit a signature in main memory for each column
  - ②  $Sim(C_1, C_2)$  is the same as the “similarity” of  $h(C_1)$  and  $h(C_2)$
- By hashing columns into buckets we expect that “most” pairs of near duplicate documents hash into the same bucket!
- Goal: Find a hash function  $h(\cdot)$  such that:
  - ◆ If  $sim(C_1, C_2)$  is high, then with high probability  $h(C_1) = h(C_2)$
  - ◆ If  $sim(C_1, C_2)$  is low, then with high probability  $h(C_1) \neq h(C_2)$
- Clearly, the **hash function depends on the similarity metric**:
  - ◆ Not all similarity metrics have a suitable hash function!
  - ◆ There is a suitable hash function for the Jaccard similarity:
    - It is called **Min-Hashing**!

Inria

17



# Minhashing

Fall 2019

- History: invented by Andrei Broder in 1997 (AltaVista) to detect near duplicate web pages
- Imagine the rows of the Boolean matrix permuted under **random permutation  $\pi$**
- Define a “hash” function  $h_{\pi}(C)$  = the index of the **first** (in the permuted order  $\pi$ ) row in which column  $C$  has value 1:
  - ◆  $h_{\pi}(C) = \min_{\pi} \pi(C)$




Inria

18



## Min-hashing Example

2nd element of the permutation  
is the first to map to a 1 in col 1



19



## Surprising Property

- The probability (over all permutations of the rows) that  $h(C_1)=h(C_2)$  is the same as  $\text{Sim}(C_1, C_2)$ :
  - ♦  $\Pr[h_\pi(C_1) = h_\pi(C_2)] = \text{sim}(C_1, C_2)$
- With multiple signatures we get a **good approximation**
- Use several **independent hash functions** to create a signature of a column
  - ♦ The **similarity of signatures** is the **fraction of the hash functions** in which they agree
  - ♦ Because of the Min-Hash property, the similarity of columns is the same as the expected similarity of their signatures

20



## Why?

Fall 2019

0	0
0	0
1	1
0	0
0	1
1	0

- Let  $X$  be a column (set of shingles),  $y \in X$  is a shingle
- Then:  $\Pr[\pi(y) = \min(\pi(X))] = 1/|X|$ 
  - It is equally likely that any shingle  $y \in X$  is mapped to the *min* element
- Let  $y$  be s.t.  $\pi(y) = \min(\pi(C_1 \cup C_2))$
- Then either:  $\pi(y) = \min(\pi(C_1))$  if  $y \in C_1$ , or  $\pi(y) = \min(\pi(C_2))$  if  $y \in C_2$ 
  - One of the two cols had to have 1 at position  $y$
- So the prob. that **both** are true is the prob.  $y \in C_1 \cap C_2$
- $\Pr[\min(\pi(C_1)) = \min(\pi(C_2))] = |C_1 \cap C_2| / |C_1 \cup C_2| = \text{sim}(C_1, C_2)$

Inria

21



## Four Types of Rows

Fall 2019

- Given columns  $C_1$  and  $C_2$ , rows may be classified as:

	$C_1$	$C_2$	
$a$	1	1	1 in both columns
$b$	1	0	columns are different
$c$	0	1	
$d$	0	0	0 in both columns


- Also,  $a$  = # rows of type  $a$ , etc.
- The ratio of type  $a$ ,  $b$ , and  $c$  that determine the similarity and the probability that  $h(C_1) = h(C_2)$ 
  - Note  $\text{sim}(C_1, C_2) = a / (a + b + c)$
  - Then:  $\Pr[h(C_1) = h(C_2)] = \text{sim}(C_1, C_2)$
- Look down the permuted columns  $C_1$  and  $C_2$  until we see a 1
  - If it's a type- $a$  row, then  $h(C_1) = h(C_2)$
  - If a type- $b$  or type- $c$  row, then not

Inria

22



## Min Hashing – Example



## MinHash – False Positive/Negative

- Instead of comparing sets, we now compare only 1 bit!
- **False positive?**
  - ♦ False positive can be easily dealt with by doing an additional layer of checking (treat minhash as a filtering mechanism)
- **False negative?**
  - ♦ Requiring full match of signature is strict, some similar sets will be lost
- High error rate! Can we do better?



## Minhash Signatures

Fall 2019

- Pick (say) 100 random permutations of the rows
- Think of  $Sig(C)$  as a column vector
- Let  $sig(C)[i] = \min(\pi_i(C))$   
according to the  $i$ th permutation, the number of the first row that has a 1 in column  $C$
- **Note:** The sketch (signature) of column  $C$  is small **~400** bytes!
  - ◆ We achieved our goal! We “compressed” long bit vectors into short signatures

Inria

25



## Implementation Trick

Fall 2019

- **Permuting rows even once is prohibitive**
  - ◆ Suppose 1 billion rows
  - ◆ Hard to pick a random permutation from 1...billion
    - Sorting would take a long time
    - Representing a random permutation requires 1 billion entries
- **A good approximation to permuting rows:** pick 100 (?) hash functions  $h_i$ 
  - ◆ Simulate the effect of a random permutation by a random hash function that **maps row numbers to as many buckets as there are rows**
  - ◆ **Row hashing:** ordering under  $h_i$  gives a **random row permutation!**
- **One-pass implementation**
  - ◆ For each column  $C$  and each hash function  $h_i$ , keep a “slot”  $M(i, C)$  for the min-hash value
  - ◆ **Intent:**  $M(i, C)$  will become the smallest value of  $h_i(r)$  for which column  $C$  has 1 in row  $r$ 
    - i.e.,  $h_i(r)$  gives order of rows for  $i$ th permutation

Inria

26



## Implementation

Fall 2019

```

M(i,C) = ∞
for each row r
  for each column C
    if C has 1 in row r // Scan rows looking for 1s
      for each hash function hi do
        // Suppose row r has 1 in column C
        if hi(r) is a smaller value than M(i,C) then
          M(i,C) := hi(r);
  
```

How to pick a random hash function h(x)?

Universal hashing:

$h_{a,b}(x) = ((a \cdot x + b) \bmod p) \bmod N$  where:

a, b ... random integers

p ... prime number (p > N)

Inria

27



## Example

Fall 2019

Row	C1	C2
1	1	0
2	0	1
3	1	1
4	1	0
5	0	1

Jaccard=1/5

$$h(x) = x \bmod 5$$

$$g(x) = 2x + 1 \bmod 5$$

$$h(1) = 1$$

$$g(1) = 3$$

$$h(2) = 2$$

$$g(2) = 0$$

$$h(3) = 3$$

$$g(3) = 2$$

$$h(4) = 4$$

$$g(4) = 4$$

$$h(5) = 0$$

$$g(5) = 1$$

Sig1 Sig2

$$1 \quad \infty$$

$$3 \quad \infty$$

$$1 \quad 2$$

$$3 \quad 0$$

$$1 \quad 2$$

$$2 \quad 0$$

$$1 \quad 2$$

$$2 \quad 0$$

$$1 \quad 0$$

$$2 \quad 0$$

M(i,C)

Inria

28



## So far ...

Fall 2019

- Represent a document as a set of hash values (of its k-shingles)
- Transform set of k-shingles to a set of minhash signatures
- Use Jaccard to compare two documents by comparing their signatures
- Is this method (i.e., transforming sets to signature) necessarily “better”??


Inria

30



Fall 2019

## Locality-Sensitive Hashing



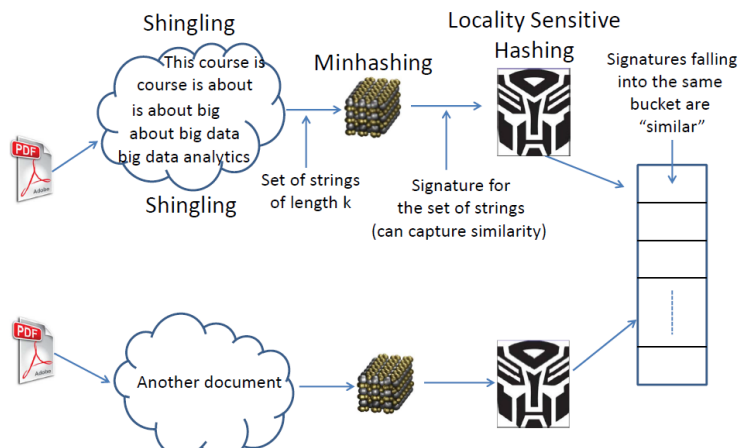
Inria

31



## The BIG Picture (All-pair comparison)

Fall 2019



- Suppose, in **main memory**, a **representation** of a large number of objects
  - ◆ May be signatures of documents as in minhashing
- We want to **pair-wise compare each** for finding those pairs that are sufficiently similar

Inria

32



## Finding Similar Pairs

Fall 2019

- While the signatures of all columns may fit in main memory, **comparing the signatures of all pairs of columns is quadratic in the number of columns**
- **Naïve solution**
  - ◆ For each document, compare with the other **N-1** documents
    - Takes N-1 comparisons
    - Can be optimized using *filter-and-refine* mechanisms
  - ◆ Requires  $N(N-1)/2$  comparisons
- **Example:**
  - ◆  $10^7$  columns implies  $\sim 10^{14}$  column-comparisons
  - ◆ At 1  $\mu$ s/comparison  $10^8$  ( $\sim 3$  years!)

Inria

33





## Locality-Sensitive Hashing

Fall 2019

- Use a function  $f(x, y)$  that tells whether or not  $x$  and  $y$  is a **candidate pair**: a pair of elements whose similarity must be evaluated
- With only one hash function on one entire column of signature, likely to have many *false negatives*
- **Key idea**: Apply the **hash** function on the **columns** of signature matrix  $M$  **multiple times**, each on a partition of the column
  - ◆ Arrange that (only) **similar columns are likely to hash** (i.e., with high probability) **to the same bucket**
  - ◆ Each pair of columns that hashes **at least once** into the same bucket is a **candidate pair**

Unlabeled


34



Fall 2019

## Candidate Generation from Minhash Signatures

- Pick a similarity threshold  $s$ , a fraction  $0 < s < 1$
- A pair of columns  $x$  and  $y$  is a candidate pair if their signatures agree in at least fraction  $s$  of the rows
  - ◆ i.e.,  $M(i, x) = M(i, y)$  for at least fraction  $s$  values of  $i$
  - ◆ we expect documents  $x$  and  $y$  to have the same (Jaccard) similarity as their signatures



Unlabeled

35




## Partition Into Bands

- Divide matrix  $M$  into  $b$  bands of  $r$  rows

- ◆ For each document, compute  $b$  sets of  $r$  minhash values
- ◆ Each set is a **mini-signature** with  $r$  minhash functions (or a concatenation of the  $r$  minhash values together)

$b$  mini-signatures




Unria

36



## Partition into Bands

- For **each band**, hash its portion of each column (i.e., the concatenated values) to a hash table with  $k$  buckets
  - ◆ this has the “same” effect as ensuring all columns have the same values
  - ◆ make  $k$  as large as possible to minimize collision
- **Candidate** column pairs are those that hash to the same bucket for  $\geq 1$  band
- Tune  $b$  and  $r$  to catch *most similar pairs*, but few *non-similar pairs*



Unria

37



## Simplifying Assumption

- There are **enough buckets** that **columns** are unlikely to hash to the same bucket unless they are *identical* in a particular band
  - ◆ Hereafter, we assume that “**same bucket**” means “**identical in that band**”
  - ◆ Assumption needed only to simplify analysis, not for correctness of algorithm
- **Finding all pairs within a bucket become computationally cheaper!**
  - ◆ Declare all pairs within a bucket to be “matching” OR
  - ◆ Perform pair-wise comparisons for those documents that fall into the same bucket
    - Much smaller than pair-wise over all documents

Blue	White	White	White
White	Blue	White	White
White	White	Blue	White
White	White	White	Blue



## Example: Effect of Bands

- Suppose  $10^5$  columns of  $M$  (100k docs)
- Signatures of  $10^2$  integers (rows)
- If each signature is represented as a 4 byte integer value, we need only  $10^2 * 4 * 10^5 = 40\text{Mb}$  of memory!
- $5 * 10^9$  pairs of signatures can take a while to compare
- Choose 20 bands of 5 integers/band
- Goal: Find pairs of documents that are *at least*  $s = 0.8$  similar



## Suppose $C_1, C_2$ are 80% Similar

- Find pairs of  $\geq s=0.8$  similarity, set  $b=20, r=5$
- Assume:  $\text{sim}(C_1, C_2) = 0.8$ 
  - ◆ Since  $\text{sim}(C_1, C_2) \geq s$ , we want  $C_1, C_2$  to be a **candidate pair**
  - ◆ We want them to hash to at **least 1 common bucket** (at least one band is identical)
- Probability  $C_1, C_2$  identical in one particular band:  $(0.8)^5 = 0.328$
- Probability  $C_1, C_2$  are **not** similar in all of the 20 bands:  $(1-0.328)^{20} = 0.00035$ 
  - ◆ i.e., about 1/3000th of the 80%-similar column pairs are **false negatives** (we miss them)
- We would find 99.965% pairs of truly similar documents

40



## Suppose $C_1, C_2$ are 30% Similar

- Find pairs of  $\geq s=0.8$  similarity, set  $b=20, r=5$
- Assume:  $\text{sim}(C_1, C_2) = 0.3$ 
  - ◆ Since  $\text{sim}(C_1, C_2) < s$  we want  $C_1, C_2$  to hash to **NO common buckets** (all bands should be different)
- Probability  $C_1, C_2$  identical in one particular band:  $(0.3)^5 = 0.00243$ 
  - ◆ Probability  $C_1, C_2$  identical in at least 1 of 20 bands:  $1 - (1-0.00243)^{20} = 0.0474$
  - ◆ In other words, approximately 4.74% pairs of docs with similarity 0.3 end up becoming **candidate pairs**
  - ◆ They are **false positives** since we will have to examine them (they are candidate pairs) but then it will turn out their similarity is below threshold  $s$

41




# LSH Involves a Tradeoff

- How to get a step-function?
- Pick:
  - ◆ The number of Min-Hashes (rows of  $M$ )
  - ◆ The number of bands  $b$ , and
  - ◆ The number of rows  $r$  per band
 to balance false positives/negatives
- Example: if we had only 20 bands of 5 rows, the number of false negatives would go down, but the number of false positives would go up

Probability  
of sharing  
a bucket

## Analysis of LSH – What We Want



Similarity  $t = \text{sim}(C_1, C_2)$  of two sets


Unria

42



# What One Band Gives You

Probability  
of sharing  
a bucket



Single hash  
signature

Similarity  $t = \text{sim}(C_1, C_2)$  of two sets

- This is what 1 hash-code gives you  
 $\Pr[h_\pi(C_1) = h_\pi(C_2)] = \text{sim}(C_1, C_2)$

Unria


43



# What $b$ Bands of $r$ Rows Gives You

Fall 2019

- The S-curve is where the “magic” happens



Inria

44



## Example: $b = 20$ ; $r = 5$

Fall 2019

$t$	$1 - (1 - t^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996




Figure 3.7: The S-curve

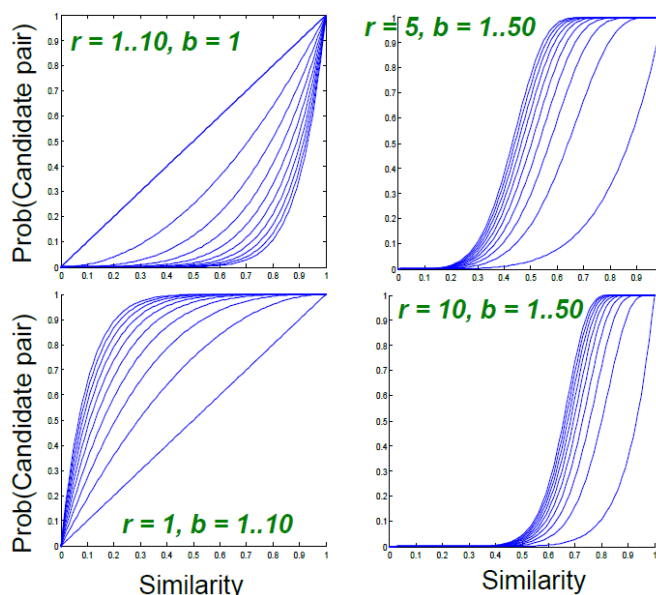
Inria

45



## S-curves as a Function of $b$ and $r$

- Given a fixed threshold  $s$
- We want choose  $r$  and  $b$  such that the  $\text{Pr}(\text{Candidate pair})$  has a “step” right around  $s$




Unria



## Picking $r$ and $b$ : The S-Curve

- Picking  $r$  and  $b$  to get the best S-curve



**Blue area:** False Negative rate  
These are pairs with  $\text{sim} > s$  but the X fraction won't share a band and then will never become candidates. This means we will never consider these pairs for (slow/exact) similarity calculation!

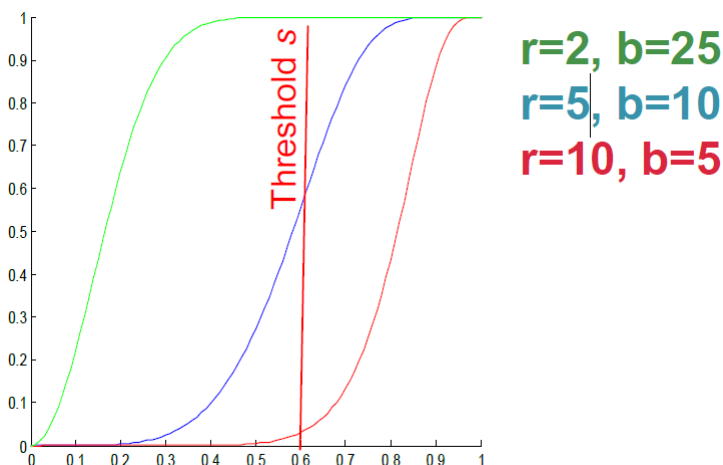
**Green area:** False Positive rate  
These are pairs with  $\text{sim} < s$  but we will consider them as candidates. This is not too bad, we will consider them for (slow/exact) similarity computation and discard them.

Unria



## Picking $r$ and $b$ to Get Desired Performance

- 50 hash-functions ( $r * b = 50$ )



## Limitations of Minhash

- Minhash is great for near-duplicate detection
  - ◆ Set high threshold for Jaccard similarity
- Limitations:
  - ◆ Jaccard similarity only
  - ◆ Set-based representation, no way to assign weights to features
- Random projections:
  - ◆ Works with arbitrary vectors using cosine similarity
  - ◆ Same basic idea, but details differ
  - ◆ Slower but more accurate: no free lunch!





## LSH Generalizations



## Multiple Hash Functions

- For Min-Hashing signatures, we got a Min-Hash function for each permutation of rows
- So far, we have **assumed only one hash function** (even applied multiple times)
  - ◆ Shorthand:  $h(x)=h(y)$  implies “**h says x and y are equal**”
- We could have used a **family of hash functions**
  - ◆ A (large) set of related hash functions generated by some mechanism
  - ◆ We should be able to efficiently **pick a hash function at random** from such a family




## Locality-Sensitive (LS) Families

Fall 2019

- Consider a space  $S$  of points with a distance measure  $d$
- A family  $H$  of hash functions is said to be  $(d_1, d_2, p_1, p_2)$ -sensitive if for any  $x$  and  $y$  in  $S$ :
  - If  $d(x, y) \leq d_1$ , then prob over all  $h$  in  $H$  that  $h(x)=h(y)$  is at least  $p_1$
  - If  $d(x, y) \geq d_2$ , then prob over all  $h$  in  $H$  that  $h(x)=h(y)$  is at most  $p_2$

Small distance,  
high probability  
of hashing to  
the same value



Large distance,  
low probability  
of hashing to  
the same value

Inria

53



## Example of LS Family: MinHash

Fall 2019

- Let
  - $S$  = space of all sets,
  - $d$  = Jaccard distance,
  - $H$  is family of Min-Hash functions for all permutations of rows
- Minhashing gives a  $(d_1, d_2, p_1, p_2)$ -sensitive family for any  $d_1 < d_2$ 
  - E.g.,  $H$  is a  $(1/3, 2/3, 2/3, 1/3)$ -sensitive family for  $S$  and  $d$
  - If distance  $\leq 1/3$  (i.e., similarity  $\geq 2/3$ ), then probability that minhash values agree is  $\geq 2/3$
  - This is because for any hash function  $h \in H$ 

$$\Pr(h(x)=h(y))=1-d(x,y)$$
- Simply restates theorem about Min-Hashing in terms of distances rather than similarities!

Inria

54



## Example of LS Family: MinHash

Fall 2019

- **Claim:** Min-hash  $H$  is a  $(\boxed{1/3}, \boxed{2/3}, \boxed{2/3}, 1/3)$ -sensitive family for  $S$  and  $d$

If distance  $< 1/3$   
(so similarity  $\geq 2/3$ )

Then probability that Min-Hash values agree  $\geq 2/3$

- For Jaccard similarity, Min-Hashing gives a  $(d_1, d_2, (1-d_1), (1-d_2))$ -sensitive family for any  $d_1 < d_2$
- Theory leaves unknown what happens to pairs that are at distance between  $d_1$  and  $d_2$ 
  - ♦ **Consequence:** No guarantees about fraction of **false positives** in that range

Inria

55



## Amplifying a LS-family

Fall 2019

- Can we reproduce the “S-curve” effect we saw before for any LS family?
- The “bands” technique we learned for signature matrices carries over to this more general setting
  - ♦ So we can do LSH with any  $(d_1, d_2, p_1, p_2)$ -sensitive family
- Two constructions:
  - ♦ **AND** construction like “rows in a band”
  - ♦ **OR** construction like “many bands”

Inria

56



## AND Construction of Hash Functions

- Given family  $H$ , construct family  $H'$  consisting of  $r$  functions from  $H$
- For  $h=[h_1, \dots, h_r]$  in  $H'$ ,  $h(x)=h(y)$  if and only if  $h_i(x)=h_i(y)$  for all  $i$ ,  $1 \leq i \leq r$
- Note this has the same effect as “ $r$  signatures”
  - $x$  and  $y$  are considered a candidate pair if every one of the  $r$  rows say that  $x$  and  $y$  are equal
- Theorem: If  $H$  is  $(d_1, d_2, p_1, p_2)$ -sensitive, then  $H'$  is  $(d_1, d_2, p_1^r, p_2^r)$ -sensitive
  - That is, for any  $p$ , if  $p$  is the probability that a member of  $H$  will declare  $(x, y)$  to be a candidate pair, then the probability that a member of  $H'$  will so declare is  $p^r$
  - Proof: Use the fact that  $h_i$ 's are independent

57



## OR Construction of Hash Functions



- Given family  $H$ , construct family  $H'$  consisting of  $b$  functions from  $H$
- For  $h=[h_1, \dots, h_b]$  in  $H'$ ,  $h(x)=h(y)$  if and only if  $h_i(x)=h_i(y)$  for at least one  $i$ ,  $1 \leq i \leq b$
- Mirrors the effect of combining “ $b$  bands”:
  - $x$  and  $y$  become a candidate pair if any set makes them a candidate pair
- Theorem: If  $H$  is  $(d_1, d_2, p_1, p_2)$ -sensitive, then  $H'$  is  $(d_1, d_2, 1-(1-p_1)^b, 1-(1-p_2)^b)$ -sensitive
  - That is, for any  $p$ , if  $p$  is the probability that a member of  $H$  will declare  $(x, y)$  to be a candidate pair, then  $(1-p)$  is the probability that it will not declare so
  - $(1-p)^b$  is the probability that none of the family  $h_1, h_b$  will declare  $(x, y)$  a candidate pair
  - $1-(1-p)^b$  is the probability that at least one  $h_i$  will declare  $(x, y)$  a candidate pair, and therefore that  $H'$  will declare  $(x, y)$  to be a candidate pair

59



## Effect of AND & OR Constructions

- **AND** makes all probabilities **shrink**, but by choosing **r** correctly, we can make the *lower probability approach 0* while the higher does not
- **OR** makes all probabilities **grow**, but by choosing **b** correctly, we can make the *upper probability approach 1* while the lower does not



## Composing Constructions: AND-OR Composition

- **r**-way **AND** construction followed by **b**-way **OR** construction
  - ◆ Exactly what we did with minhashing
    - If **b** bands match in all **r** values hash to same bucket
    - Columns that are hashed into  $\geq 1$  common bucket  $\rightarrow$  candidate
- Take points **x** and **y** s.t.  $\Pr[h(x)=h(y)] = p$ 
  - ◆ **H** will make **(x, y)** a **candidate** pair with probability **p**
- Construction makes **(x, y)** a **candidate** pair with probability  $1-(1-pr)^b$ 
  - ◆ The S-Curve!



## Example

- **Example:** Take **H** and construct **H'** by the **AND** construction with  $r = 4$ . Then, from **H'**, construct **H''** by the **OR** construction with  $b = 4$
- E.g., transform a (0.2, 0.8, 0.8, 0.2)-sensitive family into a (0.2, 0.8, 0.8785, 0.0064)-sensitive family

$p$	$1-(1-p^4)^4$
.2	.0064
.3	.0320
.4	.0985
.5	.2275
.6	.4260
.7	.6666
.8	.8785
.9	.9860



## Composing Constructions: OR-AND Composition


- $b$ -way **OR** construction followed by  $r$ -way **AND** construction
- Transforms probability  $p$  into  $(1-(1-p)^b)^r$ 
  - ◆ The same S-curve, mirrored horizontally and vertically



## Example

- **Example:** Take **H** and construct **H'** by the **OR** construction with  $b = 4$ . Then, from **H'**, construct **H''** by the **AND** construction with  $r = 4$
- E.g., transform a (0.2, 0.8, 0.8, 0.2)-sensitive family into a (0.2, 0.8, 0.9936, 0.1215)-sensitive family

p	$(1-(1-p)^4)^4$
.1	.0140
.2	.1215
.3	.3334
.4	.5740
.5	.7725
.6	.9015
.7	.9680
.8	.9936



64



## Cascading Constructions

- **Example:** Apply the (4,4) **OR-AND** construction followed by the (4,4) **AND-OR** construction
- Transforms a (.2,.8,.8,.2)-sensitive family into a (.2,.8,.9999996,.0008715)-sensitive family
  - ◆ Note this family uses  $256 (= 4*4*4*4)$  of the original hash functions

65




## Applications of LSH



### Application 2: A LHS Family for Fingerprint Matching

- Fingerprint can be uniquely defined by its minutiae
- By overlaying a grid on the fingerprint image, we can extract the grid squares where the minutiae are located
- Two fingerprints are similar if the set of grid squares significantly overlap
  - ♦ Jaccard distance and minhash can be used, but ...
- Let  $F$  be a family of functions
  - ♦  $f \in F$  is defined by, say 3, grid squares such that  $f$  returns the same bucket whenever the fingerprint has minutiae in all three grid squares
  - ♦  $f$  sends all fingerprints that have minutiae in all three of  $f$ 's grid points to the same bucket
  - ♦ Two fingerprints match if they are in the same bucket







## LSH for Fingerprint Matching

Fall 2019

- Suppose probability of finding a minutiae in a random grid square of a random finger is 0.2
- And probability of finding one in the same grid square of the same finger (different fingerprint) is 0.8
- Prob two fingerprints from different fingers match  $= (0.2)^3 \times (0.2)^3 = 0.000064$
- Prob two fingerprints from the same finger match  $= (0.2)^3 \times (0.8)^3 = 0.004096$
- Use more functions from F!
- Take 1024 functions and do a OR construction
  - ◆ Prob putting the fingerprints from the same finger in at least one bucket is  $1 - (1 - 0.004096)^{1024} = 0.985$
  - ◆ Prob two fingerprints from different fingers falling into the same bucket is  $1 - (1 - 0.000064)^{1024} = 0.063$
  - ◆ We have 1.5% false negatives and 6.3% false positives
- Using AND construction will
  - ◆ Greatly reduce the prob of a false positive
  - ◆ Small increase in false-negative rate

Inria

71



## References

Fall 2019

- CS9223 – Massive Data Analysis J. Freire & J. Simeon New York University Course 2013
- CS246: Mining Massive Datasets Jure Leskovec, Stanford University, 2014
- CS5344: Big Data Analytics Technology, TAN Kian-Lee, National University of Singapore 2014

Inria

72