




Relational Data Processing on MapReduce

Fall 2019



Vassilis Christophides

christop@csd.uoc.gr

<http://www.csd.uoc.gr/~hy562>

University of Crete, Fall 2019


Inria

1



Peta-scale Data Analysis

Fall 2019



Inria


© 2014 IBM Corporation 2



Big Data Analysis

Fall 2019

- A lot of these datasets have some structure
 - ◆ Query logs
 - ◆ Point-of-sale records
 - ◆ User data (e.g., demographics)
 - ◆ ...
- How do we perform data analysis at scale?
 - ◆ Relational databases and SQL
 - ◆ MapReduce (Hadoop)



Inria

3



Relational Databases vs. MapReduce

Fall 2019

- Relational databases:
 - ◆ Multipurpose: analysis and transactions; batch and interactive
 - ◆ Data integrity via ACID transactions
 - ◆ Lots of tools in software ecosystem (for ingesting, reporting, etc.)
 - ◆ Supports SQL (and SQL integration, e.g., JDBC)
 - ◆ Automatic SQL query optimization
- MapReduce (Hadoop):
 - ◆ Designed for large clusters, fault tolerant
 - ◆ Data is accessed in “native format”
 - ◆ Supports many query languages
 - ◆ Programmers retain control over performance

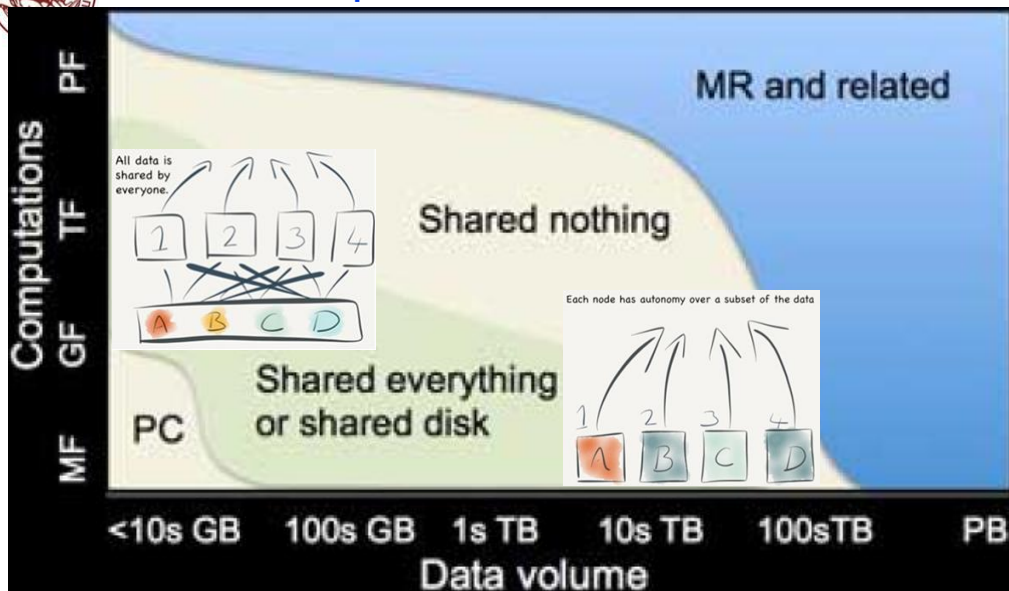
Inria

4



Parallel Computation & Data Size Matters!

Fall 2019



Inria

5



Parallel Relational Databases vs. MapReduce

Fall 2019

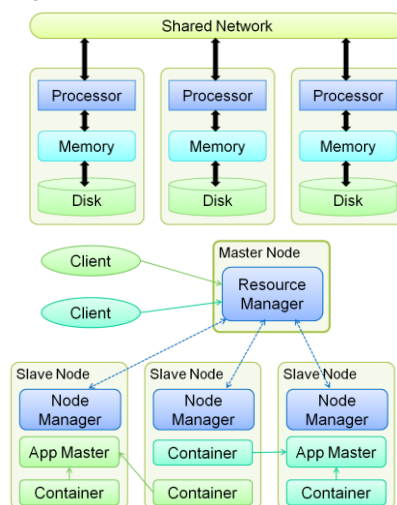
Parallel relational databases

- ◆ Schema on “write”
- ◆ Failures are relatively infrequent
- ◆ “Possessive” of data
- ◆ Mostly proprietary

MapReduce

- ◆ Schema on “read”
- ◆ Failures are relatively common
- ◆ In situ data processing
- ◆ Open source

Shared-nothing architecture for parallel processing



Inria

Hadoop NextGen (YARN) architecture

6



MapReduce: A Major Step Backwards?

- MapReduce is a step backward in database access
 - ◆ Separation of the schema from the application is good
 - Sharing across multiple MR programs is difficult
 - ◆ Declarative access languages are good
 - Does not requires highly-skilled programmers
- MapReduce is poor implementation
 - ◆ Brute force and only brute force
 - no indexes: Wasteful access to unnecessary data
 - ◆ Don't need 1000 nodes to process petabytes
 - Parallel DBs do it in fewer than 100 nodes
- MapReduce is missing features
 - ◆ Bulk loader, indexing, updates, transactions...
 - ◆ No support for JOINS:
 - Requires multiple MR phases for the analysis



Map Reduce vs Parallel DBMS

	Parallel DBMS	MapReduce
Schema Support	✓	Not out of the box
Indexing	✓	Not out of the box
Programming Model	Declarative (SQL)	Imperative (C/C++, Java, ...) Extensions through Pig and Hive
Optimizations (Compression, Query Optimization)	✓	Not out of the box
Flexibility	Not out of the box	✓
Fault Tolerance	Coarse grained techniques	✓



Database Workloads

Fall 2019

- OLTP (online transaction processing)

- ◆ Typical applications: e-commerce, banking, airline reservations
- ◆ User facing: *real-time*, *low latency*, *highly-concurrent*
- ◆ Tasks: relatively small set of “standard” *transactional queries*
- ◆ Data access pattern: *random reads*, *updates*, *writes* (involving relatively small amounts of data)



- OLAP (online analytical processing)

- ◆ Typical applications: business intelligence, data mining
- ◆ Back-end processing: *batch workloads*, *less concurrency*
- ◆ Tasks: complex *analytical queries*, often ad hoc
- ◆ Data access pattern: *table scans*, large amounts of data involved per query

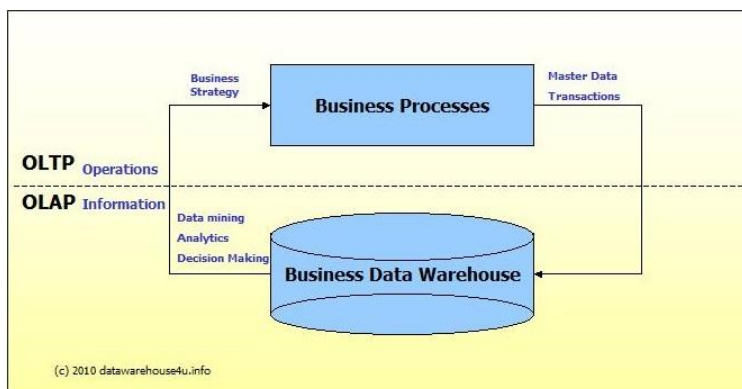
Inria

9



One Database or Two?

Fall 2019



- Downsides of co-existing OLTP and OLAP workloads

- ◆ Poor memory management
- ◆ Conflicting data access patterns
- ◆ Variable latency

- Solution: separate databases

- ◆ User-facing OLTP database for high-volume transactions
- ◆ Data warehouse for OLAP workloads
- ◆ How do we connect the two?

Inria

10



OLTP/OLAP Integration

Fall 2019



- OLTP database for user-facing transactions
 - ◆ Retain records of all activity
 - ◆ Periodic ETL (e.g., nightly)
- Extract-Transform-Load (ETL)
 - ◆ Extract records from source
 - ◆ Transform: clean data, check integrity, aggregate, etc.
 - ◆ Load into OLAP database
- OLAP database for data warehousing
 - ◆ Business intelligence: reporting, ad hoc queries, data mining, etc.
 - ◆ Feedback to improve OLTP services

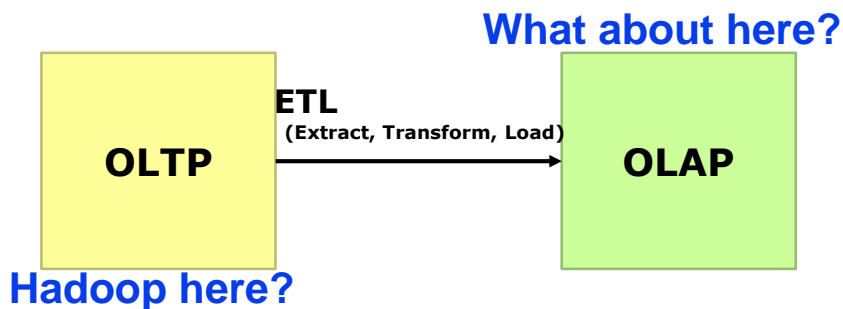
Inria

11



OLTP/OLAP Architecture: Hadoop?

Fall 2019




Inria

12



OLTP/OLAP/Hadoop Architecture

Fall 2019



- Why does this make sense?

Inria

13



ETL Bottleneck

Fall 2019

- Reporting is often a nightly task:
 - ◆ ETL is *often slow* (see next picture!)
 - What happens if processing 24 h of data takes longer than 24 h?
- Often, with noisy datasets, ETL is the analysis!
 - ◆ ETL necessarily involves brute force data scans: L, then E and T?
- Hadoop is perfect:
 - ◆ Most likely, you already have some data warehousing solution
 - ◆ *Ingest is limited by speed of HDFS*
 - ◆ *Scales out* with more nodes
 - ◆ *Massively parallel* and much cheaper than parallel databases
 - ◆ Ability to use *any processing tool*
 - ◆ ETL is a *batch process* anyway!

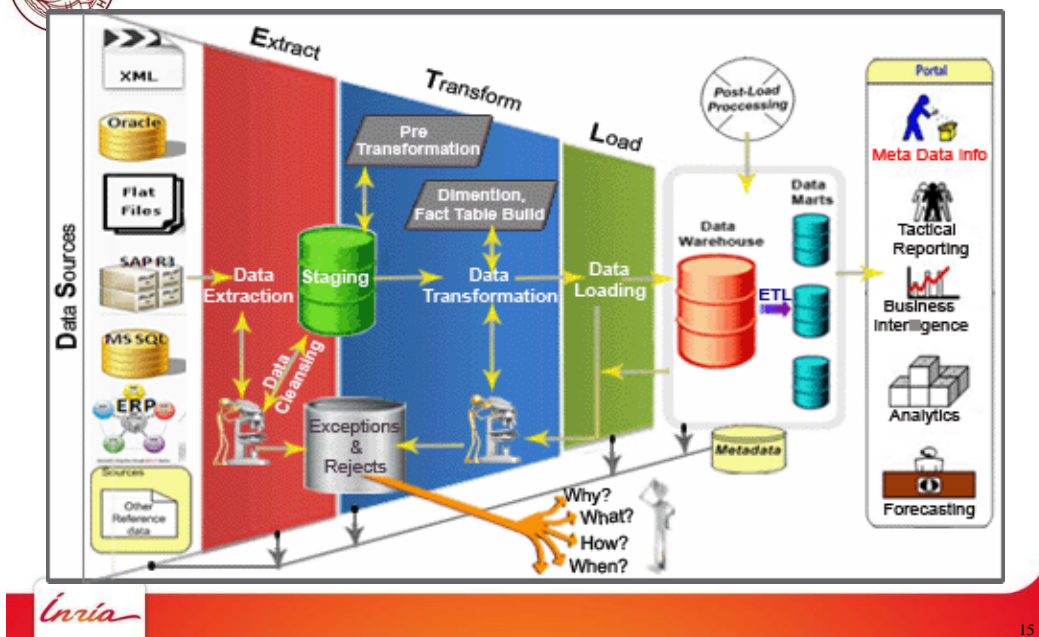
Inria

14



A Closer Look at ETL

Fall 2019



Fall 2019

MapReduce Algorithms for Processing Relational Data



Secondary Sorting

Fall 2019

- MapReduce sorts input to reducers by key

- ◆ Values are arbitrarily ordered

- What if want to sort value also?


- ◆ E.g., $k \rightarrow (v1, R), (v3, R), (v4, R), (v8, R)...$

- Solution 1:

- ◆ Buffer values in memory, then sort
 - ◆ Why is this a bad idea?

- Solution 2:

- ◆ “Value-to-key conversion”: extends the key with part of the value
 - ◆ Let execution framework do the sorting
 - ◆ Preserve state across multiple key-value pairs to handle processing
 - ◆ Anything else we need to do?



Inria

18



Value-to-Key Conversion

Fall 2019

Before

$k \rightarrow (v1, R), (v4, R), (v8, R), (v3, R)...$

Values arrive in arbitrary order...

After

$(k, v1) \rightarrow (v1, R)$ Values arrive in sorted order...

$(k, v3) \rightarrow (v3, R)$ Process by preserving state across multiple keys!

$(k, v4) \rightarrow (v4, R)$

$(k, v8) \rightarrow (v8, R)$

...

- Default comparator, group comparator, and Partitioner has to be tuned to use the appropriate part of the key

Inria

19



Working Scenario

Fall 2019

- Two tables:
 - ◆ User demographics (gender, age, income, etc.)
 - ◆ User page visits (URL, time spent, etc.)
- **Analyses** we might want to perform:
 - ◆ Statistics on demographic characteristics
 - ◆ Statistics on page visits
 - ◆ Statistics on page visits by URL
 - ◆ Statistics on page visits by demographic characteristic
 - ◆ ...

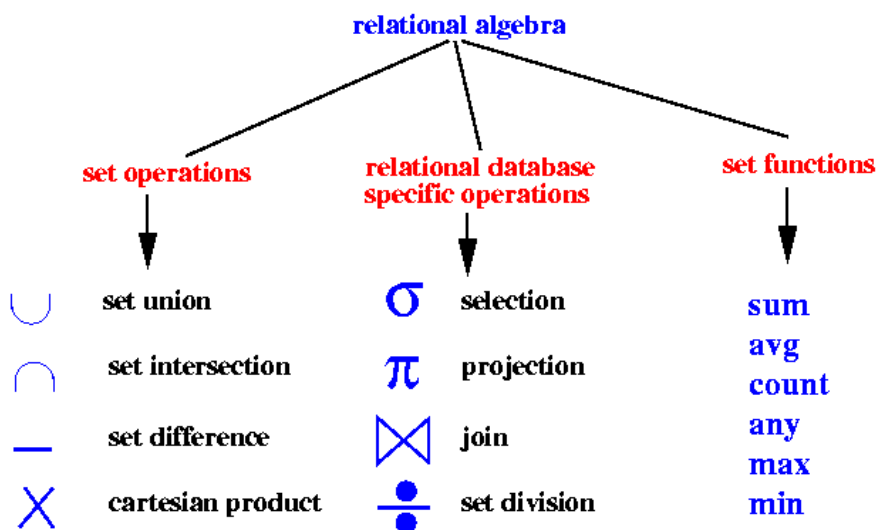
Inria

20



Relational Algebra

Fall 2019




Inria

www.mathcs.emory.edu/~cheung/Courses/377/Syllabus/4-RelAlg/intro.html



Projection

Fall 2019



Inria

22



Projection in MapReduce

Fall 2019

- Easy!
 - ◆ Map over tuples, **emit new tuples with the projected attributes**
 - For each tuple t in R , construct a tuple t' by eliminating those components whose attributes are not in S , emit a key/value pair (t', t')
 - ◆ **No reducers** (reducers are the *identity* function), unless for regrouping or resorting tuples
 - the Reduce operation performs **duplicate elimination**
 - ◆ Alternatively: perform in reducer, after some other processing
- Basically **limited by HDFS streaming speeds**
 - ◆ Speed of *encoding/decoding* tuples becomes important
 - ◆ Relational databases take advantage of *compression*
 - ◆ Semi-structured data? No problem!


Inria

23



Selection

Fall 2019



Inria

24



Selection in MapReduce

Fall 2019

- Easy!
 - ♦ Map over tuples, **emit only tuples that meet selection criteria**
 - For each tuple t in R , check if t satisfies C and If so, emit a key/value pair (t, t)
 - ♦ **No reducers** (reducers are the *identity* function), unless for regrouping or resorting tuples
 - ♦ Alternatively: perform in reducer, after some other processing
- Basically **limited by HDFS streaming speeds**:
 - ♦ Speed of *encoding/decoding tuples* becomes important
 - ♦ Relational databases take advantage of *compression*
 - ♦ Semistructured data? No problem!

Inria

25



Set Operations in Map Reduce

Fall 2019

- $R(X,Y) \cup S(Y,Z)$
 - ◆ **Map**: for each tuple t either in R or in S , emit (t,t)
 - ◆ **Reduce**: either receive $(t,[t,t])$ or $(t,[t])$
 - Always emit (t,t)
 - We perform **duplicate elimination**
- $R(X,Y) \cap S(Y,Z)$
 - ◆ **Map**: for each tuple t either in R or in S , emit (t,t)
 - ◆ **Reduce**: either receive $(t,[t,t])$ or $(t,[t])$
 - Emit (t,t) in the former case and nothing (t, NULL) in the latter
- $R(X,Y) - S(Y,Z)$
 - ◆ **Map**: for each tuple t either in R or in S , emit $(t, R \text{ or } S)$
 - ◆ **Reduce**: receive $(t,[R])$ or $(t,[S])$ or $(t,[R,S])$
 - Emit (t,t) only when received $(t,[R])$ otherwise nothing (t, NULL)

Unria

26



Group by... Aggregation

Fall 2019


- Example: What is the average time spent per URL?
- In SQL:
 - ◆ `SELECT url, AVG(time) FROM visits GROUP BY url`
- In MapReduce: Let $R(A, B, C)$ be a relation to which we apply $\gamma_{A, \theta(B)}(R)$
 - ◆ The **map** operation prepares the **grouping** (e.g., emit time, keyed by url)
 - ◆ The grouping is done by the framework
 - ◆ The **reducer** computes the **aggregation** (e.g. average)
 - ◆ Eventually, **optimize with combiners**
 - ◆ Simplifying assumptions: *one grouping attribute and one aggregation function*

Unria


27



Relational Joins



Types of Relationships





Join Algorithms in MapReduce

Fall 2019

- “Join” usually just means **equi-join**, but we also want to support *other* join predicates
- Hadoop has some built-in join support, but our goal is to understand **important algorithm design principles**
- **Algorithms**
 - ◆ Reduce-side join
 - ◆ Map-side join
 - ◆ In-memory join
 - Striped variant
 - Memcached variant

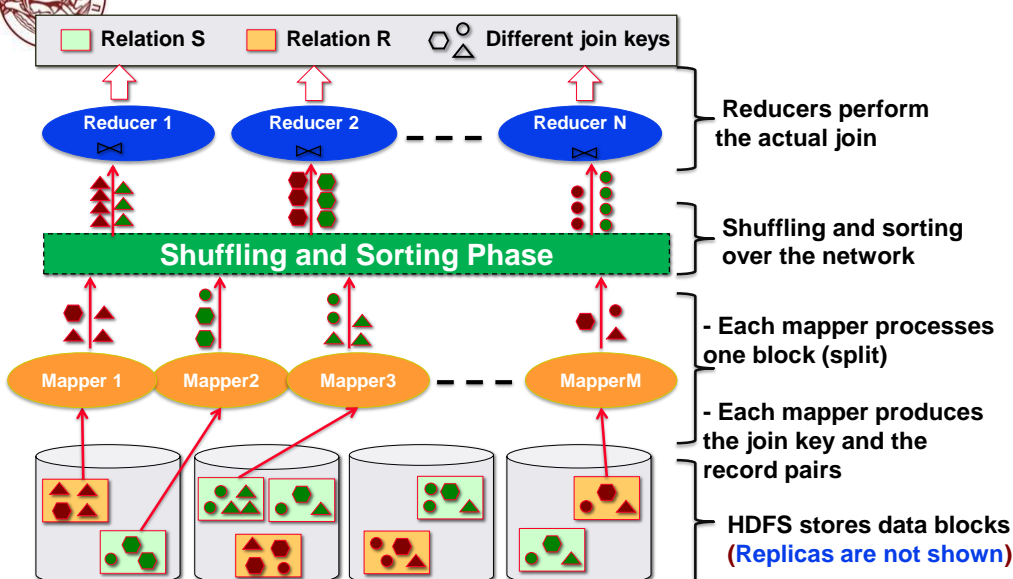
Inria

30



Re-Partition Join

Fall 2019



Inria

31



Reduce-side Join

Fall 2019

- Basic idea: **group by join key**
 - ◆ Execution framework brings together tuples sharing the same key
 - ◆ Similar to a “**sort-merge join**” in the database terminology
- A **map** function
 - ◆ Receives a record in R and S
 - ◆ Emits its **join attribute value as a key and the record as a value**
- A **reduce** function
 - ◆ Receives each join attribute value with its records from R and S
 - ◆ **Perform actual join between the records in R and S**
- Two variants
 - ◆ 1-to-1 joins
 - ◆ 1-to-many and many-to-many joins

Unria


32




Reduce-side Join: 1-to-1

Fall 2019

Map



Reduce



Note: **no guarantee if R is going to come first or S!**

Unria


33



Reduce-side Join: 1-to-Many

Fall 2019

Map



Reduce



- What's the problem?
- ♦ R is the one side, S is the many

inkling

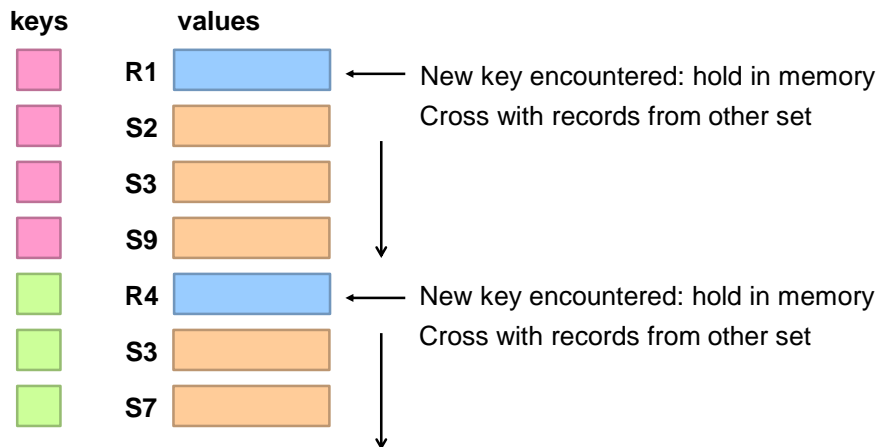
34



Reduce-side Join: Value-to-Key Conversion

Fall 2019

In reducer... Buffer all values in memory, pick out the tuple from R, and then cross it with every tuple from S to perform the join



inkling

www.inkling.com/read/hadoop-definitive-guide-tom-white-3rd/chapter-8/example-8-9

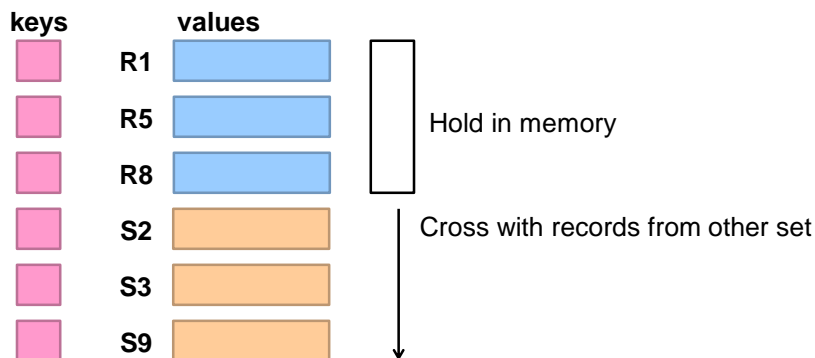
35



Reduce-side Join: Many-to-Many

Fall 2019

In reducer...



- What's the problem?
- ♦ R is the smaller dataset

Inria


36



Map-side Join: Basic Idea

Fall 2019

- What are the limitations of reduce-side joins?
- ♦ Both relations are transferred over the network
- Assume two datasets are sorted by the join key:



A sequential scan through both relations to join:
called a "sort-merge join" in database terminology

Inria

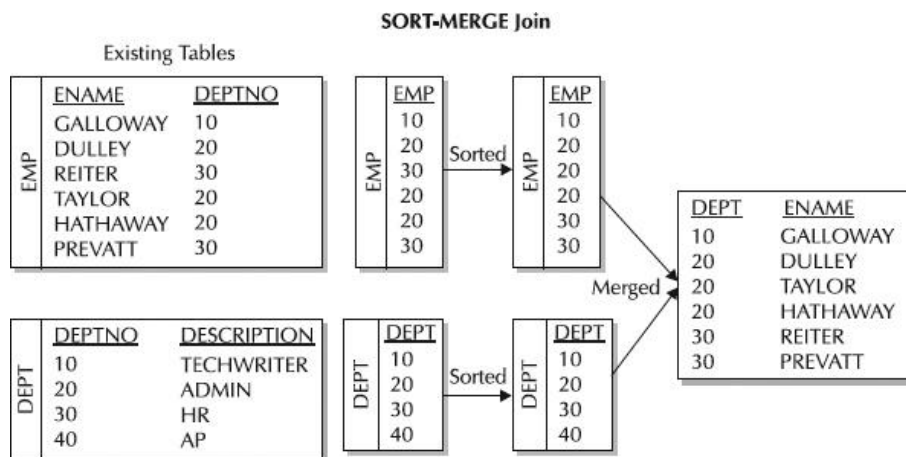
37



Map-side Join: Parallel Scans

Fall 2019

- If datasets are sorted by join key, join can be accomplished by a scan over both relations



<https://logicalread.com/oracle-11g-sort-merge-joins-mc02/>

38



Map-side Join: Parallel Scans

Fall 2019

- How can we accomplish this in parallel?
 - ◆ Partition and sort both relations in the same manner
- In MapReduce:
 - ◆ Map over one relation, read from other corresponding partition
 - ◆ No reducers necessary (unless to repartition or resort)
- Consistently partitioned relations: realistic to expect?
 - ◆ Depends on the workflow
 - ◆ For ad hoc data analysis, reduce-side are more general, although less efficient



39

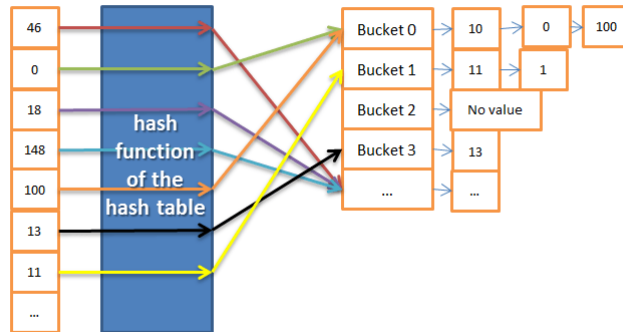


In-Memory Join: Variants

Fall 2019

- Basic idea: load one dataset into memory, stream over other dataset
 - Works if $R \ll S$ and R fits into memory
 - Called a “hash join” in database terminology

Hash Join



Outer relation

Inner relation (in-memory hash table)

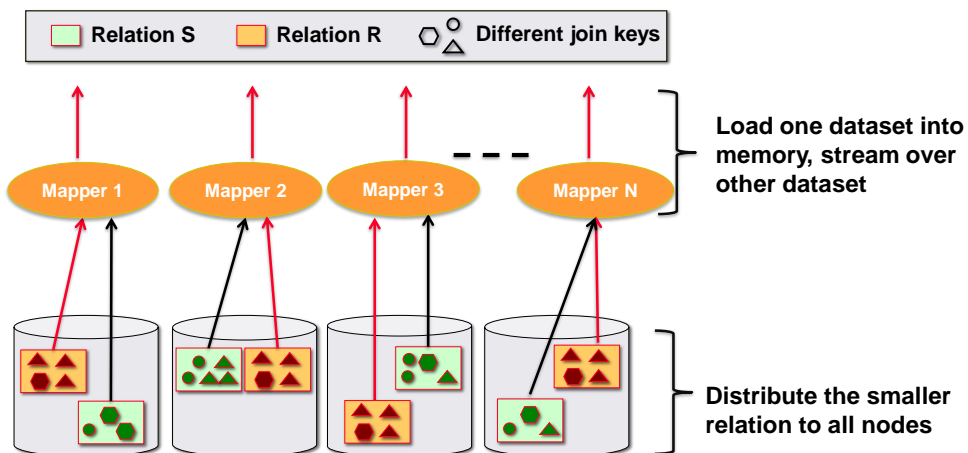
Unria

40



Broadcast/Replication Join

Fall 2019



Unria

41



In-Memory Join: Variants

Fall 2019

- MapReduce implementation
 - ◆ Distribute R to all nodes
 - ◆ Map over S, each mapper loads R in memory, hashed by join key
 - ◆ For every tuple in S, look up join key in R
 - ◆ No reducers, unless for regrouping or resorting tuples
- Downside: need to copy R to all mappers
 - ◆ Not so bad, since R is small

Inria

42



Which Join to Use?

Fall 2019

- In-memory join > map-side join > reduce-side join
 - ◆ Why?
- Limitations of each?
 - ◆ In-memory join: memory
 - ◆ Map-side join: sort order and partitioning
 - ◆ Reduce-side join: general purpose algorithm but sensible to data skewness?
- What about non-equi joins?
 - ◆ Inequality ($S.A < R.A$): map just forwards R-tuples, but replicates S-tuples for all larger R.A values as keys

Inria

46



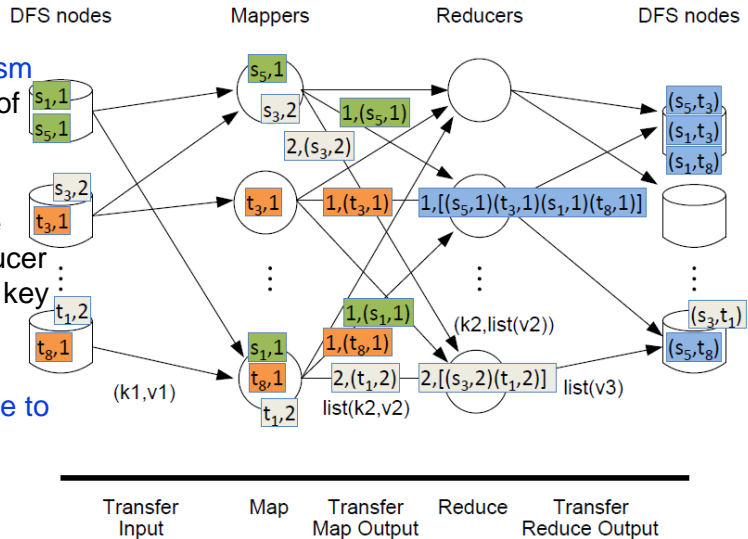
Problems With Standard Repartition Equi-Joins

- Degree of parallelism limited by number of distinct join values

- Data skew

◆ If one join value dominates, reducer processing that key will become bottleneck

- Does not generalize to other joins



Standard Repartition Equi-Join Algorithm

- Consider only the pairs with the same join attribute values

		s ₁	s ₂	s ₃	s ₄	s ₅	s ₆
		a ₁	a ₁	a ₂	a ₂	a ₃	a ₄
r ₁	a ₁	○	○				
r ₂	a ₁	○	○				
r ₃	a ₂			○	○		
r ₄	a ₃					○	

Naïve join algorithm

enumerated pairs

Standard repartition join algorithm

		s ₁	s ₂	s ₃	s ₄	s ₅	s ₆
		a ₁	a ₁	a ₂	a ₂	a ₃	a ₄
r ₁	a ₁	○	○				
r ₂	a ₁	○	○				
r ₃	a ₂			○	○		
r ₄	a ₃					○	

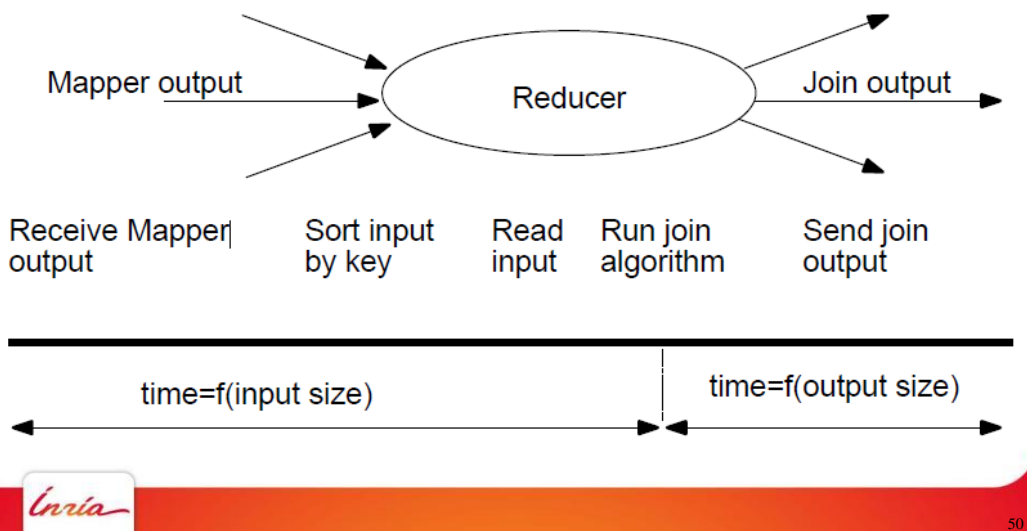




Reducer-Centric Cost Model

Fall 2019

- Difference between join implementations starts with Map output



Optimization Goal: Minimal Job Completion time

Fall 2019

- Job completion time depends on the slowest map and reduce functions
- Balancing the workloads of map functions is easy and thus we ignore them
- Balance the workloads of reduce functions as evenly as possible
 - ◆ Assume all reducers are similarly capable
- Processing time at reducer is approximately **monotonic** in input and output size
- Hence need to minimize **max-reducer-input** or **max-reducer-output**
- Join problem classification
 - ◆ **Input-size dominated**: minimize max-reducer-input
 - ◆ **Output-size dominated**: minimize max-reducer-output
 - ◆ **Input-output balanced**: minimize combination of both






Join Model

Fall 2019

- Join-matrix M : $M(i, j) = \text{true}$, if and only if (s_i, t_j) in join result
- Cover each true-valued cell by exactly one reducer



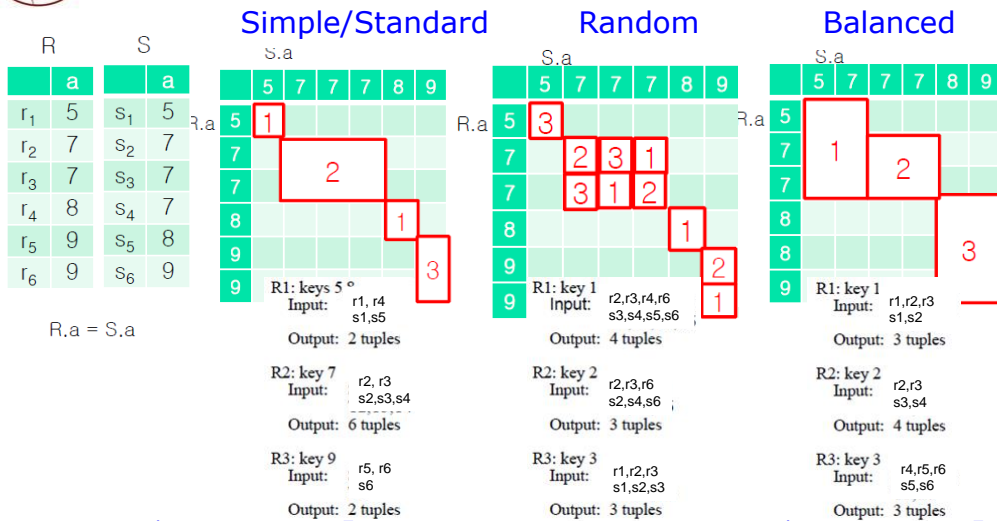
© Kyuseok Shim (VLDB 2012 TUTORIAL)

52



Fall 2019

Reduce Allocations for Repartition Equi-joins



Max reduce input size = 5 Max reduce input size = 8 Max reduce input size = 5
 Max reduce output size = 6 Max reduce output size = 4 Max reduce output size = 4



©Kyuseok Shim (VLDB 2012 TUTORIAL)

53



Comparison of Reduce Allocation Methods

- Simple allocation
 - ◆ Minimize the maximum input size of reduce functions
 - ◆ Output size may be skewed
- Random allocation
 - ◆ Minimize the maximum output size of reduce functions
 - ◆ Input size may be increased due to duplication
- Balanced allocation
 - ◆ Minimize both maximum input and output sizes



How to Balance Reduce Allocation

- Assume r is desired number of reduce functions
- Partition join-matrix M into r regions
- A map function sends each record in R and S to mapped regions
- A reduce function outputs all possible (r,s) pairs satisfying the join predicates in its value-list
- Propose M-Bucket-I algorithm [Okcan Riedewald: SIGMOD 2011]

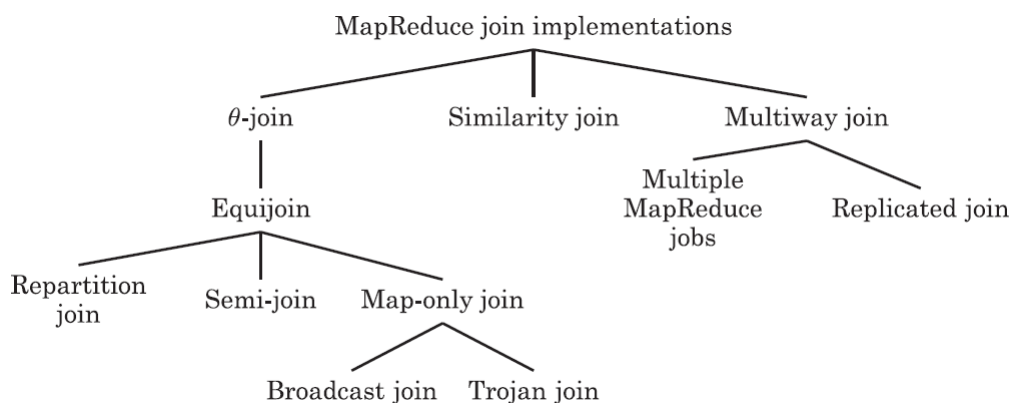


Processing Relational Data: Summary

- MapReduce algorithms for processing relational data:
 - ◆ Group by, sorting, partitioning are handled automatically by shuffle/sort in MapReduce
 - ◆ Selection, projection, and other computations (e.g., aggregation), are performed either in mapper or reducer
- Complex operations require multiple MapReduce jobs
 - ◆ Example: top ten URLs in terms of average time spent
 - ◆ Opportunities for automatic optimization
- Multiple strategies for relational joins



Join Implementations on MapReduce






Evolving Roles for Relational Database and MapReduce

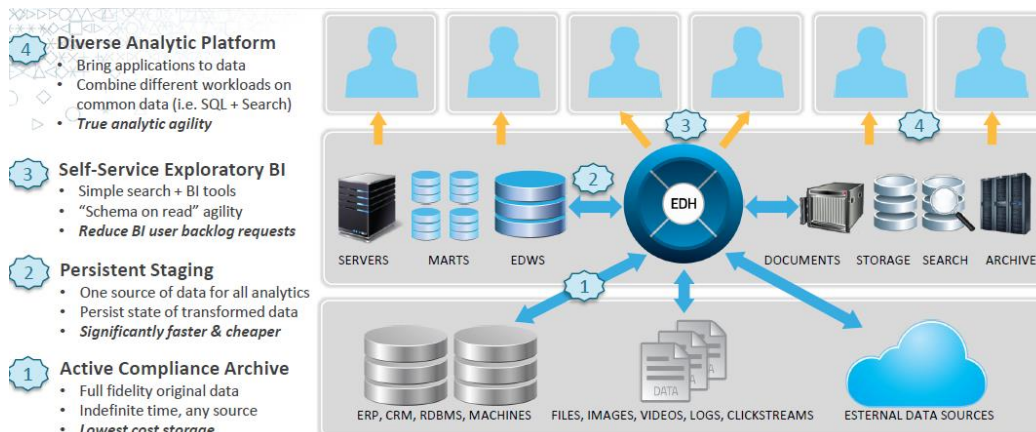


The Traditional Way: Bringing Data to Compute





The New Way: Bringing Compute to Data



Evolution from Apache Hadoop to the Enterprise Data Hub A. Awadallah Co-Founder & CTO of Cloudera SMDB 2014

60



Need for High-Level Languages

- Hadoop is great for large-data processing!
 - ◆ But writing Java programs for everything is **verbose** and **slow**
 - ◆ Analysts don't want to (or can't) write Java
- **Solution:** develop higher-level data processing languages
 - ◆ Hive: HQL is like SQL
 - ◆ Pig: Pig Latin is a bit like Perl



61



Hive and Pig

- **Hive: data warehousing application in Hadoop**
 - ◆ Query language is HQL, variant of SQL
 - ◆ Tables stored on HDFS as flat files
 - ◆ Developed by Facebook, now open source
- **Pig: large-scale data processing system**
 - ◆ Scripts are written in Pig Latin, a dataflow language
 - ◆ Developed by Yahoo!, now open source
 - ◆ Roughly 1/3 of all Yahoo! internal jobs
- **Common idea:**
 - ◆ Provide higher-level language to facilitate large-data processing
 - ◆ Higher-level language “compiles down” to Hadoop jobs



62



Hive: Example

- Hive looks similar to an SQL database
- Relational join on two tables:
 - ◆ Table of word counts from Shakespeare collection
 - ◆ Table of word counts from the bible

```
SELECT s.word, s.freq, k.freq FROM shakespeare s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```

the	25848	62394
I	23031	8854
and	19671	38985
to	18038	13526
of	16700	34654
a	14170	8057
you	12702	2720
my	11297	4135
in	10797	12445
is	8882	6884



Source: Material drawn from Cloudera training VM

63



Hive: Behind the Scenes

Fall 2019

```
SELECT s.word, s.freq, k.freq FROM shakespear s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```



(Abstract Syntax Tree)

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF shakespear s) (TOK_TABREF bible k) (= (
(TOK_TABLE_OR_COL s) word) (. (TOK_TABLE_OR_COL k) word)))) (TOK_INSERT
(TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT (TOK_SELEXPR (
(TOK_TABLE_OR_COL s) word)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL s) freq)) (TOK_SELEXPR
(. (TOK_TABLE_OR_COL k) freq))) (TOK_WHERE (AND (>= (. (TOK_TABLE_OR_COL s) freq) 1)
(>= (. (TOK_TABLE_OR_COL k) freq) 1))) (TOK_ORDERBY (TOK_TABSORTCOLNAMEDESC (.
(TOK_TABLE_OR_COL s) freq))) (TOK_LIMIT 10)))
```



(one or more of MapReduce jobs)



64



Hive: Behind the Scenes

Fall 2019

STAGE DEPENDENCIES:
Stage-1 is a root stage
Stage-2 depends on stages: Stage-1
Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1

Map Reduce

Alias -> Map Operator Tree:

s

TableScan

alias: s

Filter Operator

predicate:

expr: (freq >= 1)

type: boolean

Reduce Output Operator

key expressions:

expr: word

type: string

sort order: +

Map-reduce partition columns:

expr: word

type: string

tag: 0

value expressions:

expr: freq

type: int

expr: word

type: string

k

TableScan

alias: k

Filter Operator

predicate:

expr: (freq >= 1)

type: boolean

Reduce Output Operator

key expressions:

expr: word

type: string

sort order: +

Map-reduce partition columns:

expr: word

type: string

tag: 1

value expressions:

expr: freq

type: int

Reduce Operator Tree:

Join Operator

condition map:

Inner Join 0 to 1

condition expressions:

0 (VALUE_col0) (VALUE_col1)

1 (VALUE_col0)

outputColumnNames: _col0, _col1, _col2

Filter Operator

predicate:

expr: ((_col0 >= 1) and (_col2 >= 1))

type: boolean

Select Operator

expressions:

expr: _col1

type: string

expr: _col0

type: int

expr: _col2

type: int

outputColumnNames: _col0, _col1, _col2

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.SequenceFileInputFormat

output format: org.apache.hadoop.hive.q1.io.HiveSequenceFileOutputFormat

Stage: Stage-2

Map Reduce

Alias -> Map Operator Tree:

hdfs://localhost:8022/tmp/hive-training/364214370/10002

Reduce Output Operator

key expressions:

expr: _col1

type: int

sort order: -

tag: -1

value expressions:

expr: _col0

type: string

expr: _col1

type: int

expr: _col2

type: int

Reduce Operator Tree:

Extract

Limit

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.TextInputFormat

output format: org.apache.hadoop.hive.q1.io.HiveIgnoreKeyTextOutputFormat

Stage: Stage-0

Fetch Operator

limit: 10



65



Pig: Example

Fall 2019

- Task: Find the top 10 most visited pages in each category

Visits

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	flickr.com	10:05
Fred	cnn.com	12:00

⋮

Url Info

Url	Category	PageRank
cnn.com	News	0.9
bbc.com	News	0.8
flickr.com	Photos	0.7
espn.com	Sports	0.9

⋮

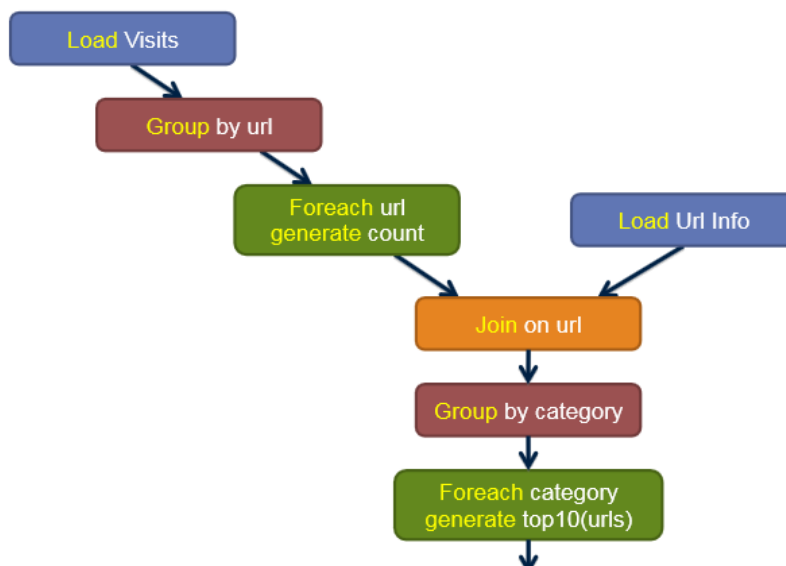
Unria

66



Pig Query Plan

Fall 2019



Unria

67



Pig Script

Fall 2019

```
visits = load '/data/visits' as (user, url, time);
gvisits = group visits by url;
visitCounts = foreach gvisits generate url, count(visits);
urlInfo = load '/data/urlInfo' as (url, category, pRank);
visitCounts = join visitCounts by url, urlInfo by url;
gCategories = group visitCounts by category;
topUrls = foreach gCategories generate top(visitCounts,10);

store topUrls into '/data/topUrls';
```

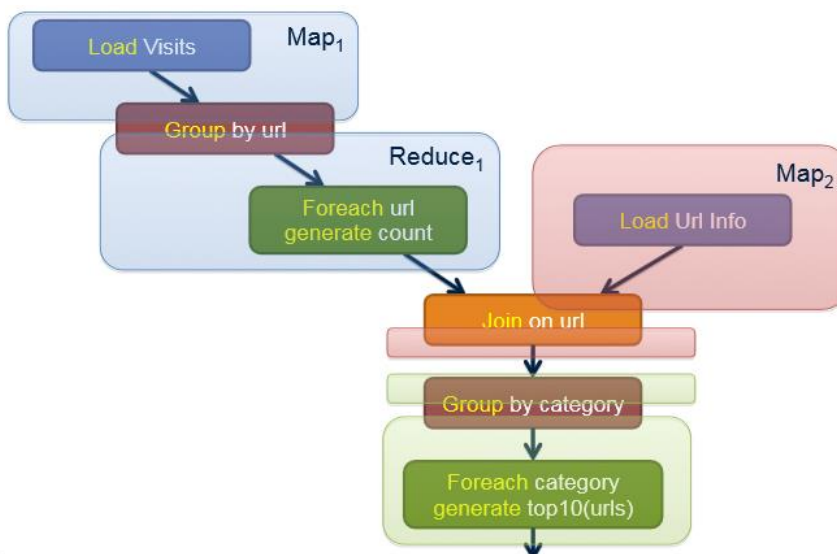
Unria

68



Pig Query Plan

Fall 2019



Unria

69



References

Fall 2019

- CS9223 – Massive Data Analysis J. Freire & J. Simeon New York University Course 2013
- INFM 718G / CMSC 828G Data-Intensive Computing with MapReduce J. Lin University of Maryland 2013
- CS 6240: Parallel Data Processing in MapReduce Mirek Riedewald Northeastern University 2014
- Extreme Computing Stratis D. Viglas University of Edinburg 2014
- MapReduce Algorithms for Big Data Analysis Kyuseok Shim VLDB 2012 TUTORIAL

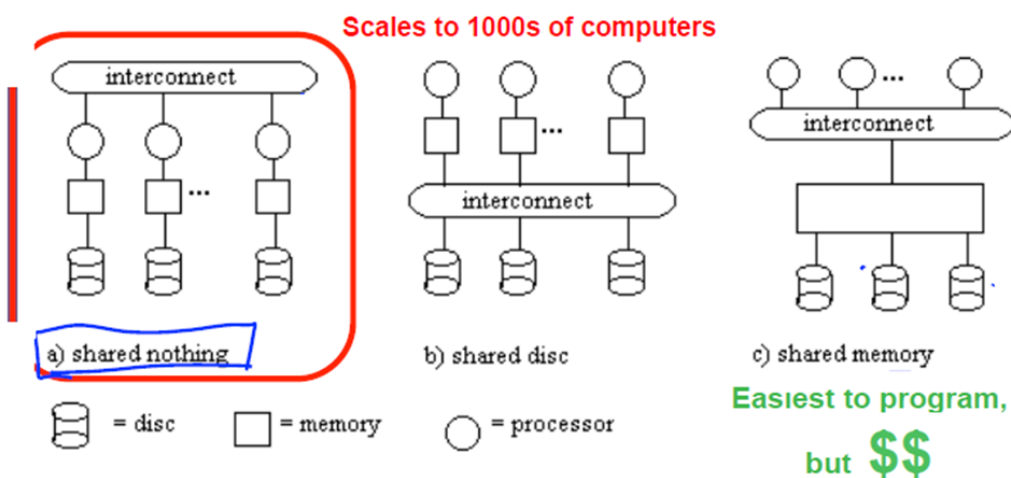
Inria

70



Taxonomy of Parallel Architectures

Fall 2019




Inria

71



Unicore vs Multi-core Architectures



Positioning Big Data

